

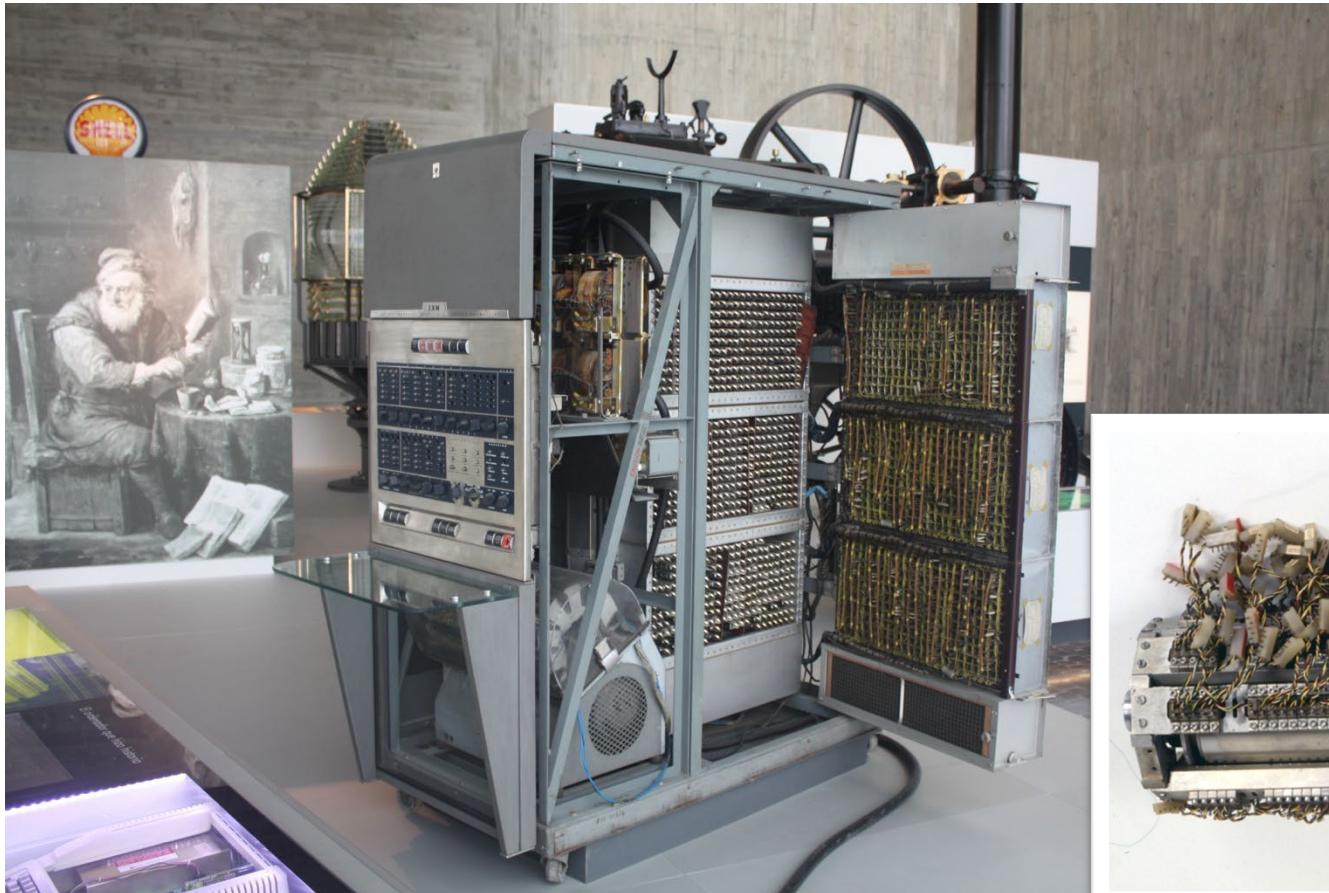
Instruction Set Architecture

Joel Emer

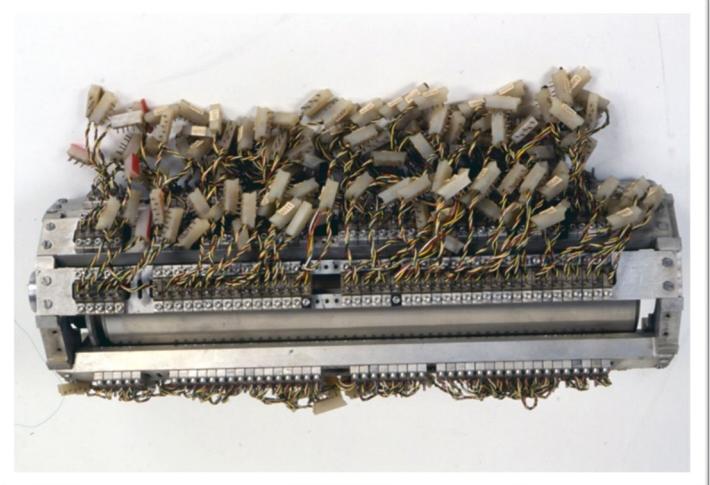
Computer Science & Artificial Intelligence Lab
M.I.T.

The IBM 650 (1953-4)

The first mass produced computer in the world.

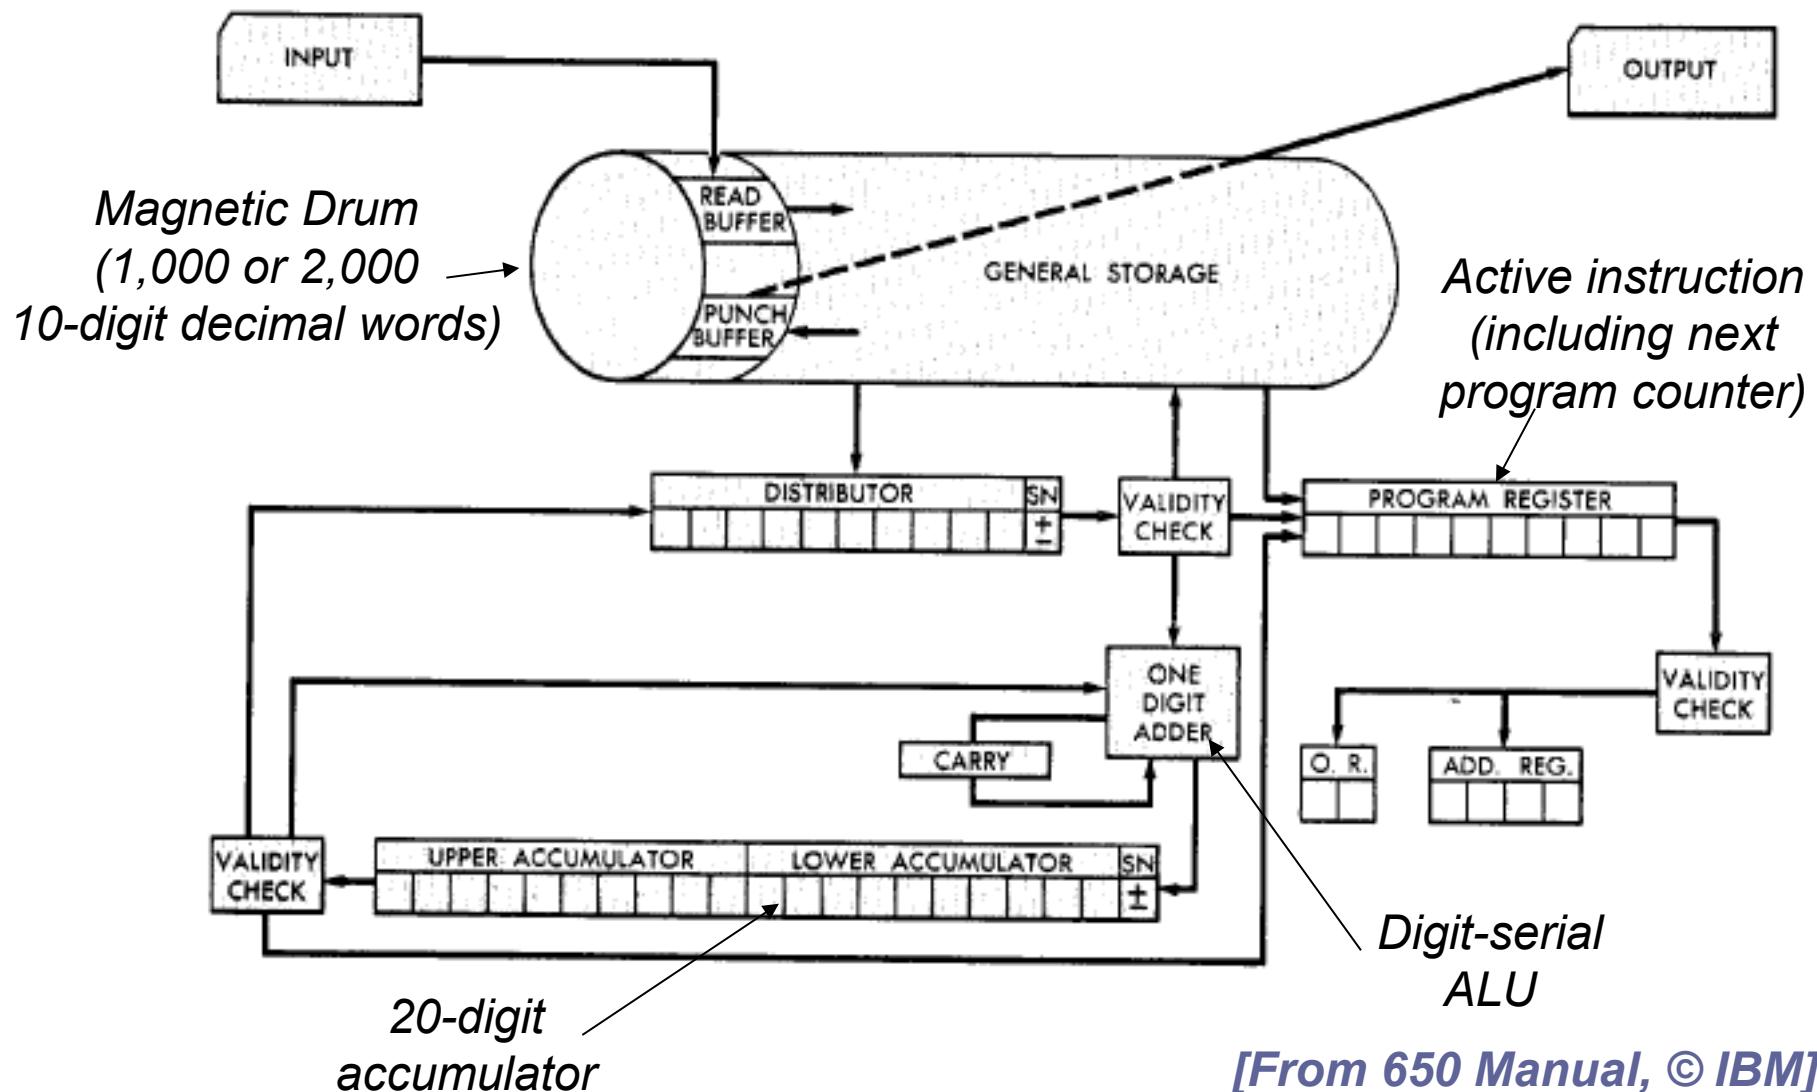


Memory drum



Side view of an IBM 650 Console Unit. First computer in Spain (1959) now at National Museum of Science and Technology in A Coruña. From Wikipedia.

The IBM 650 (1953-4)



[From 650 Manual, © IBM]

Programmer's view of a machine: IBM 650

A drum machine with 44 instructions

Instruction: 60 1234 1009

“Load the contents of location 1234 into the *distributor*;
put it also into the *upper accumulator*; set *lower
accumulator* to zero; and then go to location 1009 for the
next instruction.”

- Programmer's view of the machine was inseparable from the actual hardware implementation
- Good programmers optimized the placement of instructions on the drum to reduce latency!

Compatibility Problem at IBM

By early 60's, *IBM had 4 incompatible lines of computers!*

| | | |
|------|---|------|
| 701 | → | 7094 |
| 650 | → | 7074 |
| 702 | → | 7080 |
| 1401 | → | 7010 |

Each system had its own

- Instruction set
- I/O system and Secondary Storage:
magnetic tapes, drums and disks
- Assemblers, compilers, libraries,...
- Market niche
business, scientific, real time, ...

⇒ *IBM 360*

IBM 360: Design Premises

Amdahl, Blaauw, and Brooks, 1964

The design must lend itself to *growth and successor Machines*

- General method for connecting I/O devices
- Total performance - answers per month rather than bits per microsecond \Rightarrow *programming aids*
- Machine must be capable of *supervising itself* without manual intervention
- Built-in *hardware fault checking* and locating aids to reduce down time
- Simple to assemble systems with redundant I/O devices, memories, etc. for *fault tolerance*
- Some problems required floating point words larger than 36 bits

Processor State and Data Types

"The information held in the processor at the end of an instruction to provide the processing context for the next instruction."

Program Counter, Accumulator, ...

- The information held in the processor will be interpreted as having data types manipulated by the instructions.
- If the processing of an instruction can be **interrupted** then the *hardware* must save and restore the state in a transparent manner

Programmer's machine model is a **contract** between the hardware and software

Instruction Set

*The control for **changing** the information held in the processor are specified by the instructions available in the instruction set architecture or ISA.*

Some things an ISA must specify:

- *A way to reference registers and memory*
- *The computational operations available*
- *How to control the sequence of instructions*
- *A binary representation for all of the above*

*ISA must satisfy the needs of the software:
- assembler, compiler, OS, VM*

IBM 360: A General-Purpose Register (GPR) Machine

- Processor State
 - 16 General-Purpose 32-bit Registers
 - 4 Floating Point 64-bit Registers
 - A Program Status Word (PSW)
 - *PC, Condition codes, Control flags*
- Data Formats
 - 8-bit bytes, 16-bit half-words, 32-bit words, 64-bit double-words
 - 24-bit addresses
- A 32-bit machine with 24-bit addresses
 - *No instruction contains a 24-bit address!*
- Precise interrupts

IBM 360: Initial Implementations (1964)

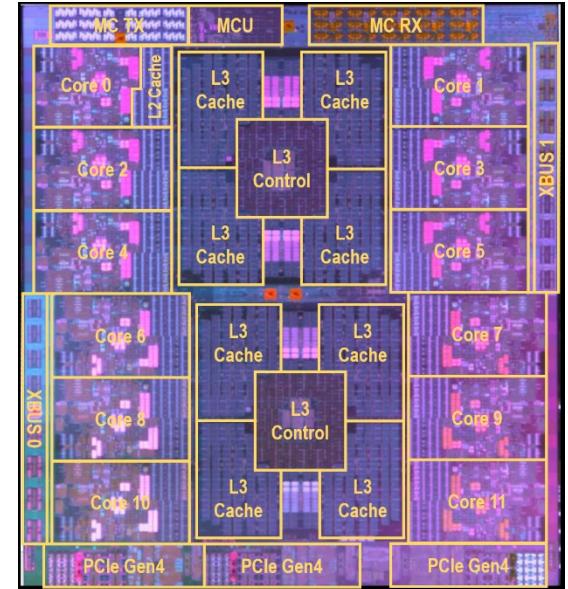
| | <i>Model 30</i> | ... | <i>Model 70</i> |
|------------------------|-----------------------|-----|--------------------|
| <i>Memory Capacity</i> | 8K - 64 KB | | 256K - 512 KB |
| <i>Memory Cycle</i> | 2.0 μ s | ... | 1.0 μ s |
| <i>Datapath</i> | 8-bit | | 64-bit |
| <i>Circuit Delay</i> | 30 nsec/level | | 5 nsec/level |
| <i>Registers</i> | in Main Store | | in Transistor |
| <i>Control Store</i> | Read only 1 μ sec | | Dedicated circuits |

- Six implementations (Models, 30, 40, 50, 60, 62, 70)
- 50x performance difference across models
- *ISA completely hid the underlying technological differences between various models*

With minor modifications, IBM 360 ISA is still in use

IBM 360: Fifty-five years later... z15 Microprocessor

- 9.2 billion transistors, 12-core design
- Up to 190 cores (2 spare) per system
- 5.2 GHz, 14nm CMOS technology
- 64-bit virtual addressing
 - Original 360 was 24-bit; 370 was a 31-bit extension
- Superscalar, out-of-order
 - 12-wide issue
 - Up to 180 instructions in flight
- 16K-entry Branch Target Buffer
 - Very large buffer to support commercial workloads
- Four Levels of caches
 - 128KB L1 I-cache, 128KB L1 D-cache
 - 4MB L2 cache per core
 - 256MB shared on-chip L3 cache
 - 960MB shared off-chip L4 cache
- Up to 40TB of main memory per system



September 2019
Image credit: IBM

Iron Law of Processor Performance

$$\frac{\text{Time}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} * \frac{\text{Cycles}}{\text{Instruction}} * \frac{\text{Time}}{\text{Cycle}}$$

- Instructions per program depends on source code, compiler technology and ISA
- Cycles per instructions (CPI) depends upon the ISA and the microarchitecture
- Time per cycle depends upon the microarchitecture and the base technology

| Microarchitecture | CPI | cycle time |
|--------------------------|-----|------------|
| Microcoded | >1 | short |
| Single-cycle unpipelined | 1 | long |
| Pipelined | 1 | short |

Summary: Instruction Set Architecture (ISA) versus Implementation

- ISA is the hardware/software interface
 - Examples: *MIPS, RISC-V, Alpha, x86, IBM 360, VAX, ARM, JVM*
- Many possible implementations of one ISA
 - 360 implementations: model 30 (c. 1964), z15 (c. 2019)
 - x86 implementations: *8086* (c. 1978), *80186*, *286*, *386*, *486*, *Pentium*, *Pentium Pro*, *Pentium-4*, *Core i7*, *AMD Athlon*, *AMD Opteron*, *Transmeta Crusoe*, *SoftPC*
 - MIPS implementations: *R2000*, *R4000*, *R10000*, ...
 - JVM: *HotSpot*, *PicoJava*, *ARM Jazelle*, ...

What potential negative consequences of having ISAs?

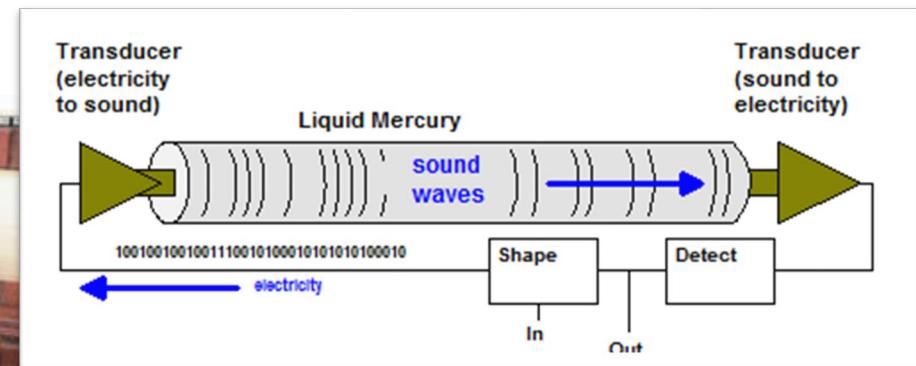
Memory and Caches

Joel Emer

Computer Science and Artificial Intelligence Laboratory
M.I.T.

Memory Technology

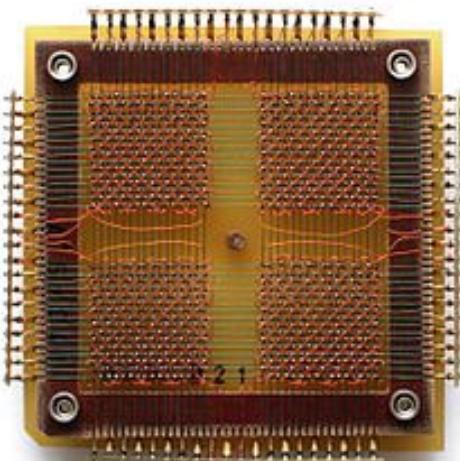
- Early machines used a variety of memory technologies
 - Manchester Mark I used CRT Memory Storage
 - EDVAC used a mercury delay line



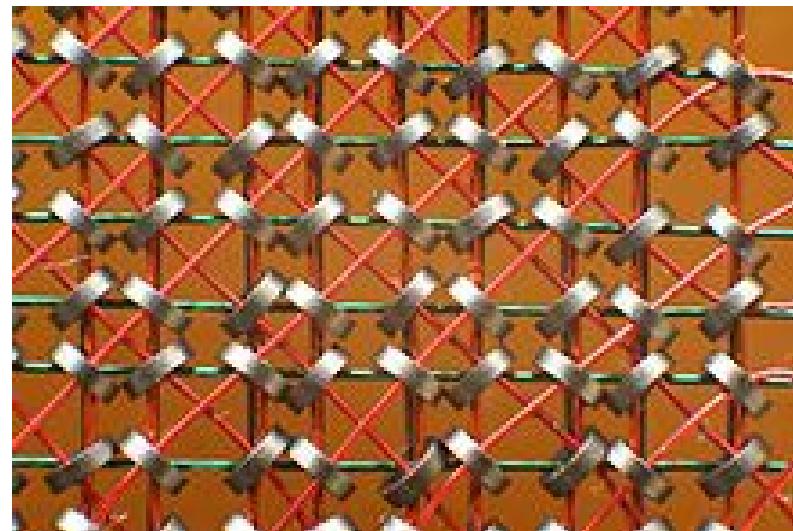
store acoustic signals in mercury

Memory Technology

- Early machines used a variety of memory technologies
 - Manchester Mark I used CRT Memory Storage
 - EDVAC used a mercury delay line
- Core memory was first large scale reliable main memory
 - Invented by Forrester in late 40s at MIT for Whirlwind project
 - Bits stored as magnetization polarity on small ferrite cores threaded onto 2 dimensional grid of wires



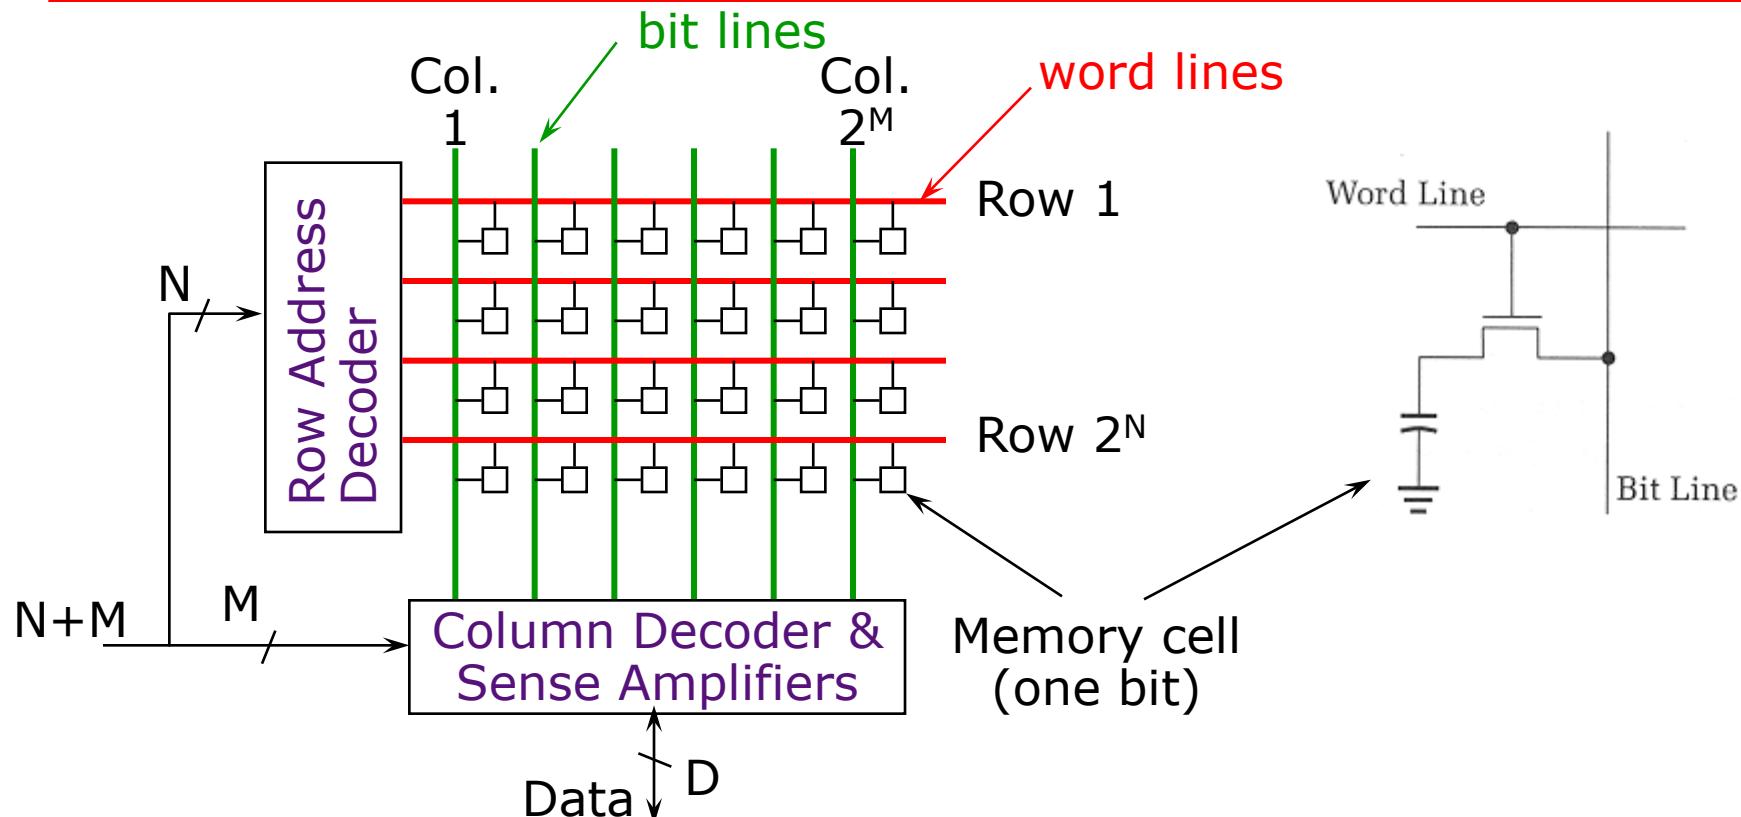
From Wikipedia



Memory Technology

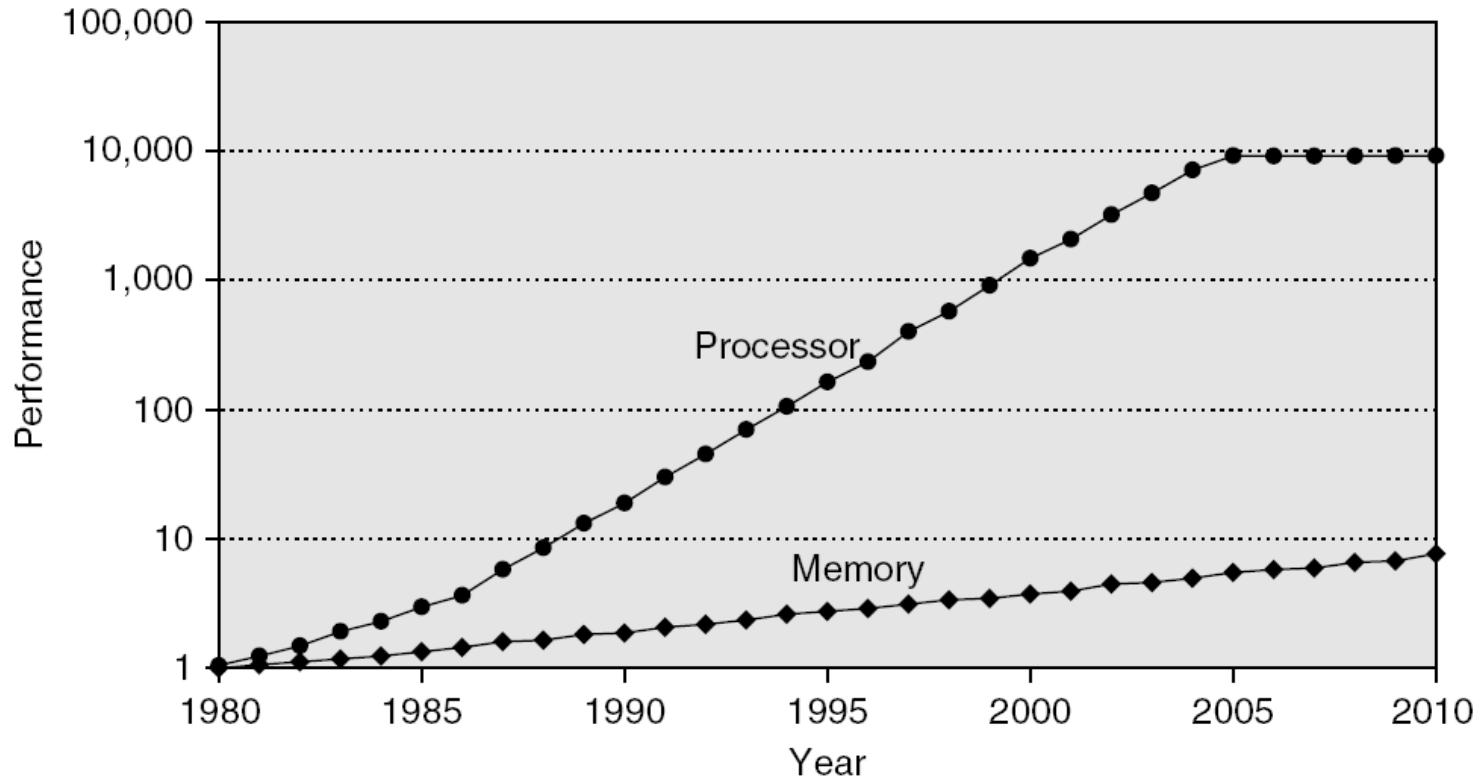
- Early machines used a variety of memory technologies
 - Manchester Mark I used CRT Memory Storage
 - EDVAC used a mercury delay line
- Core memory was first large scale reliable main memory
 - Invented by Forrester in late 40s at MIT for Whirlwind project
 - Bits stored as magnetization polarity on small ferrite cores threaded onto 2 dimensional grid of wires
- First commercial DRAM was Intel 1103
 - 1Kbit of storage on single chip
 - charge on a capacitor used to hold value
- Semiconductor memory quickly replaced core in 1970s
 - Intel formed to exploit market for semiconductor memory
- Flash memory
 - Slower, but denser than DRAM. Also non-volatile, but with wearout issues
- Phase change memory (PCM, 3D XPoint)
 - Slightly slower, but much denser than DRAM and non-volatile

DRAM Architecture



- Bits stored in 2-dimensional arrays on chip
- Modern chips have around 8 logical banks on each chip
 - Each logical bank physically implemented as many smaller arrays

Processor-DRAM Gap (latency)

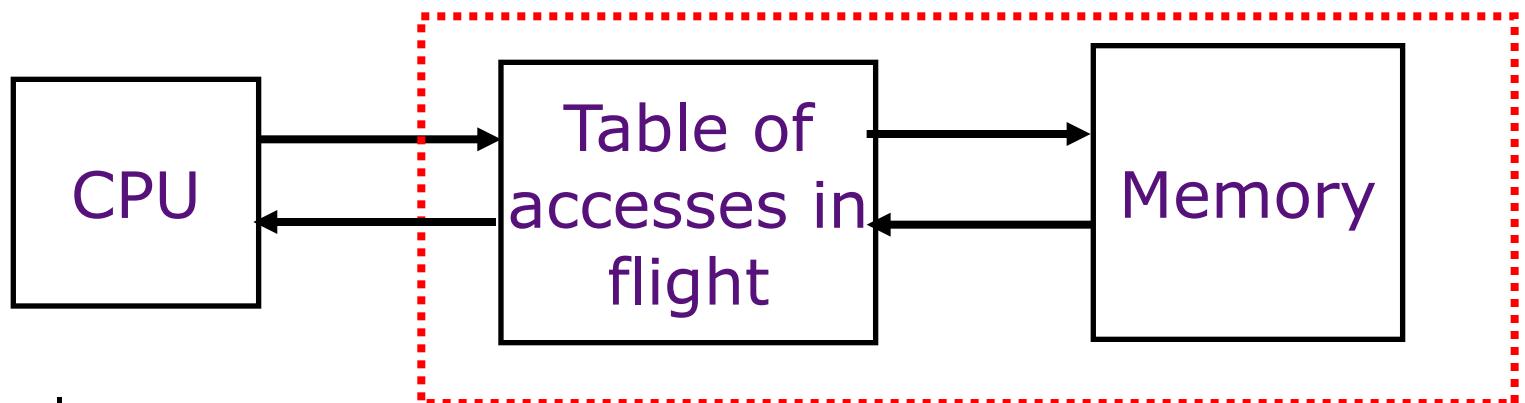


Four-issue 2GHz superscalar accessing 100ns DRAM could execute 800 instructions during time for one memory access!

How to bridge this gap?

Little's Law

$$\text{Throughput } (T) = \text{Number in Flight } (N) / \text{Latency } (L)$$



Example:

- Assume infinite-bandwidth memory
- 100 cycles / memory reference
- 1 + 0.2 memory references / instruction

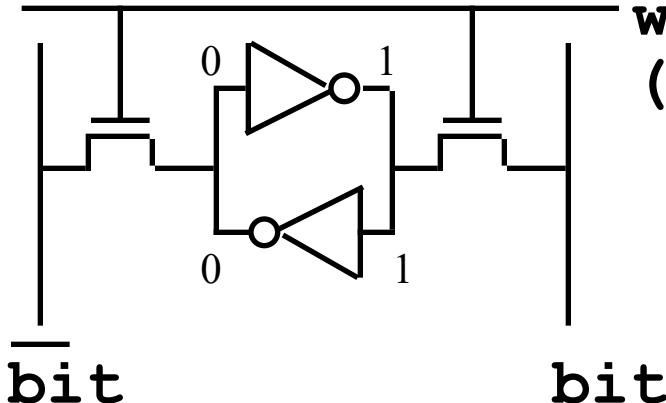
How many inflight accesses do we need to hold to keep the CPU busy?

$$\Rightarrow \text{Table size} = 1.2 * 100 = 120 \text{ entries}$$

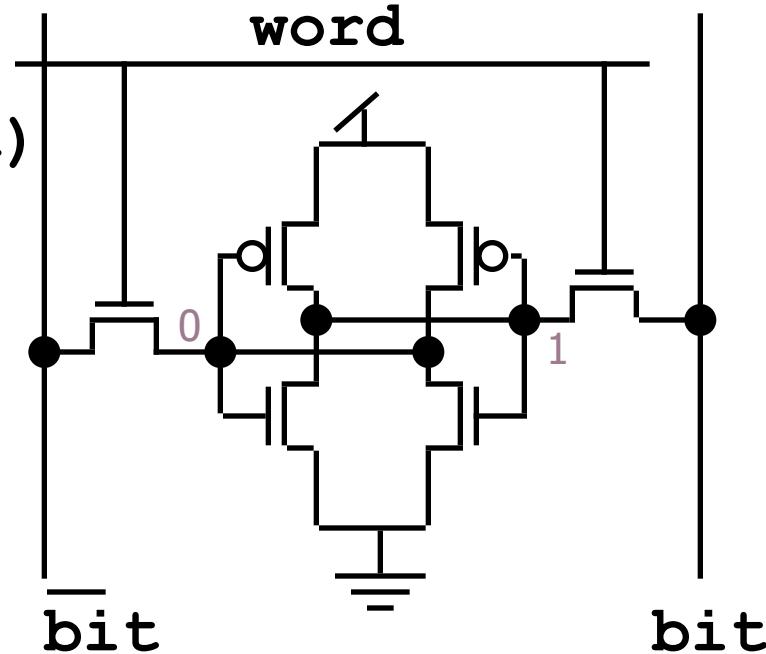
120 independent memory operations in flight!

Basic Static RAM Cell

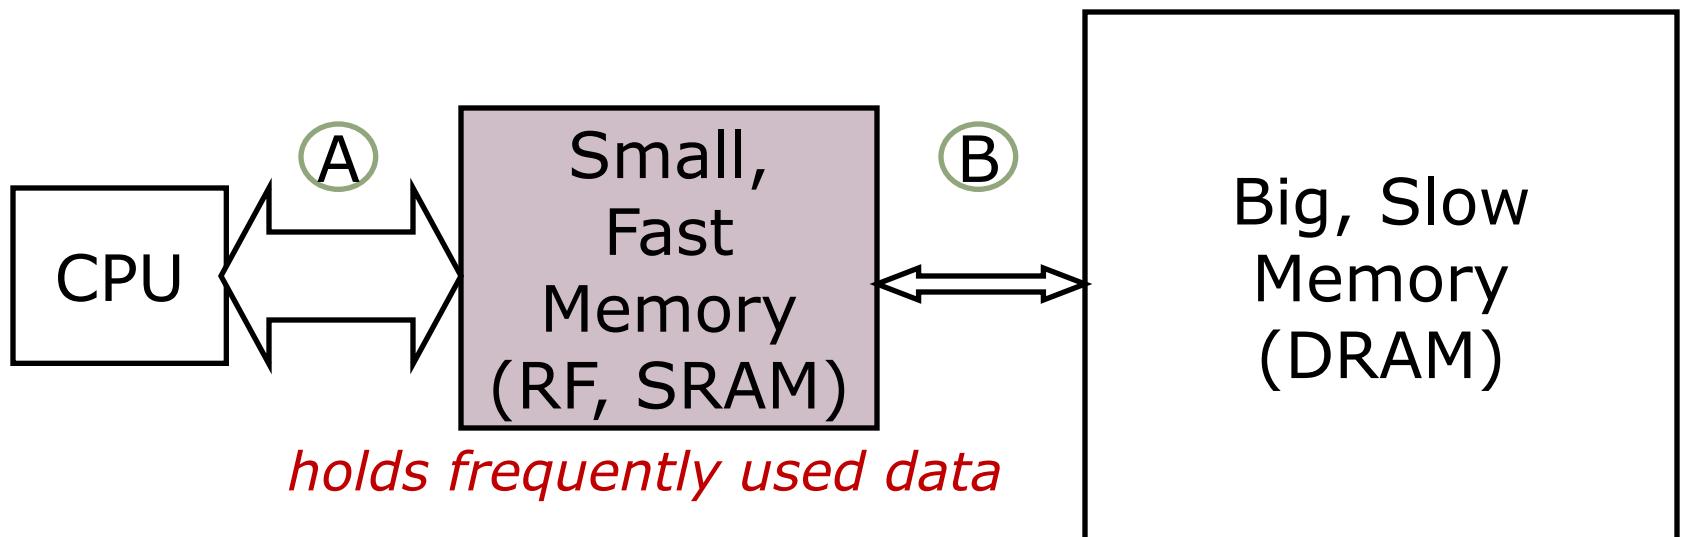
6-Transistor SRAM Cell



- Write:
 1. Drive bit lines ($\text{bit}=1$, $\overline{\text{bit}}=0$)
 2. Select word line
- Read:
 1. Precharge bit and $\overline{\text{bit}}$ to Vdd
 2. Select word line
 3. Cell pulls one bit line low
 4. Column sense amp detects difference between bit & $\overline{\text{bit}}$



The Concept of Caching



On a data access:

$\text{data} \in \text{fast memory}$

\Rightarrow low latency access

$\text{data} \notin \text{fast memory}$

\Rightarrow long latency access (*DRAM*)

- *size*: Register << SRAM << DRAM *why?*
- *latency*: Register << SRAM << DRAM
- *bandwidth*: on-chip >> off-chip

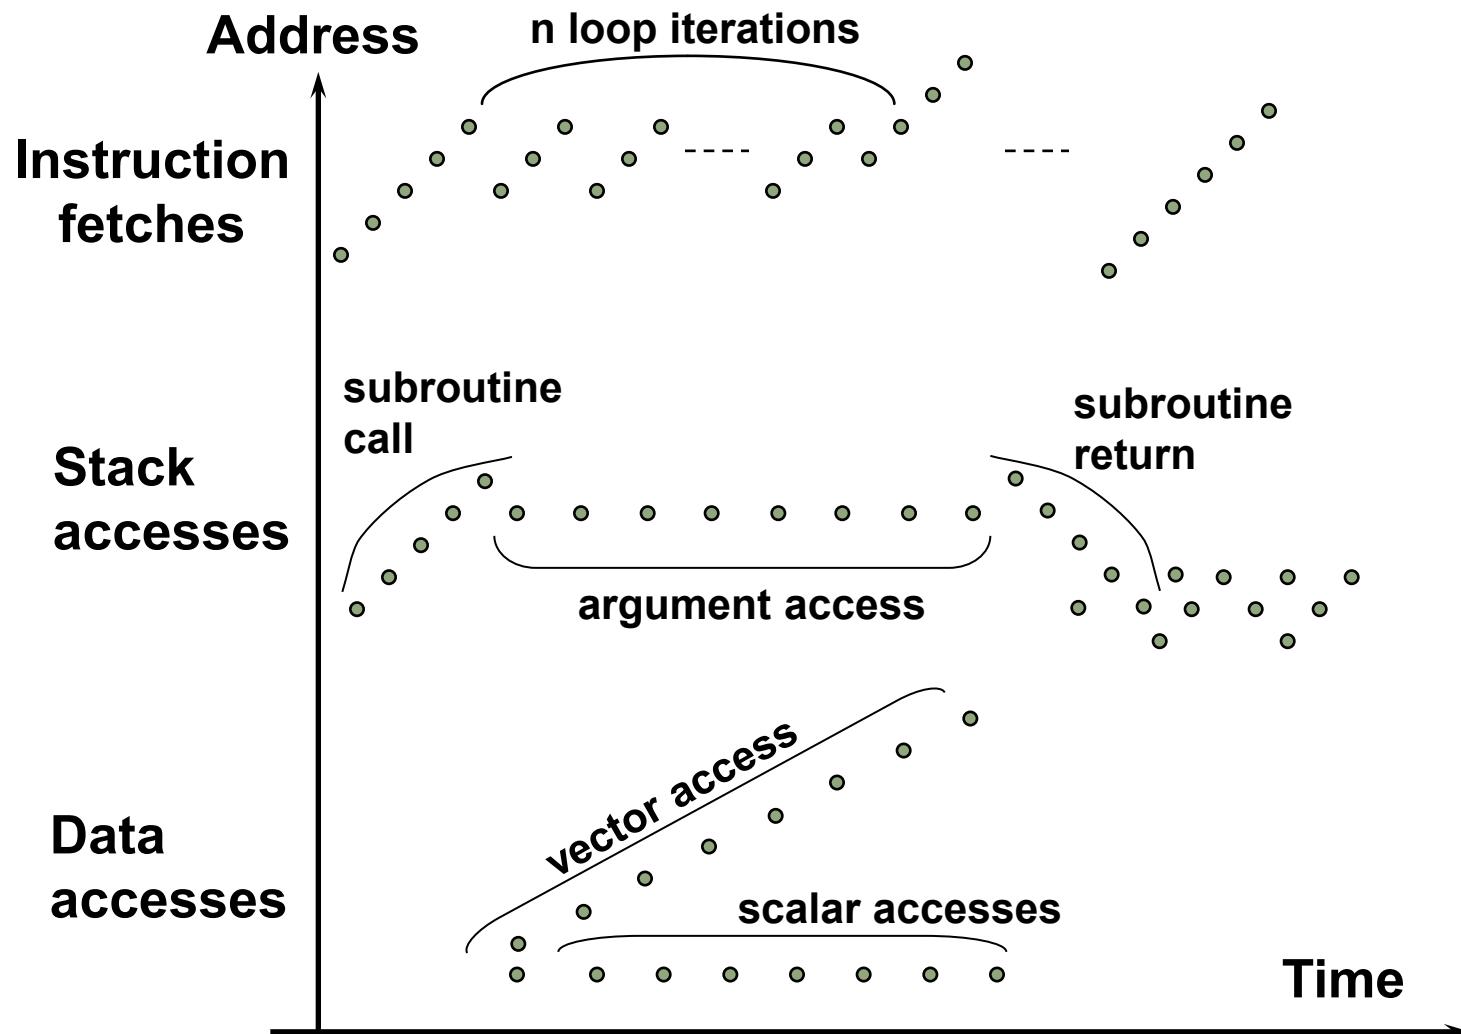
Multilevel Memory

Strategy: Reduce average latency using small, fast memories called caches.

Caches are a mechanism to reduce memory latency based on the **empirical** observation that the patterns of memory references made by a processor are often highly predictable:

| | <i>PC</i> |
|-------------------------|------------|
| ... | 96 |
| <i>Loop:</i> | |
| <i>add x2, x1, x1</i> | <i>100</i> |
| <i>addi x3, x3, -1</i> | <i>104</i> |
| <i>bne x3, x0, Loop</i> | <i>108</i> |
| ... | 112 |

Typical Memory Reference Patterns

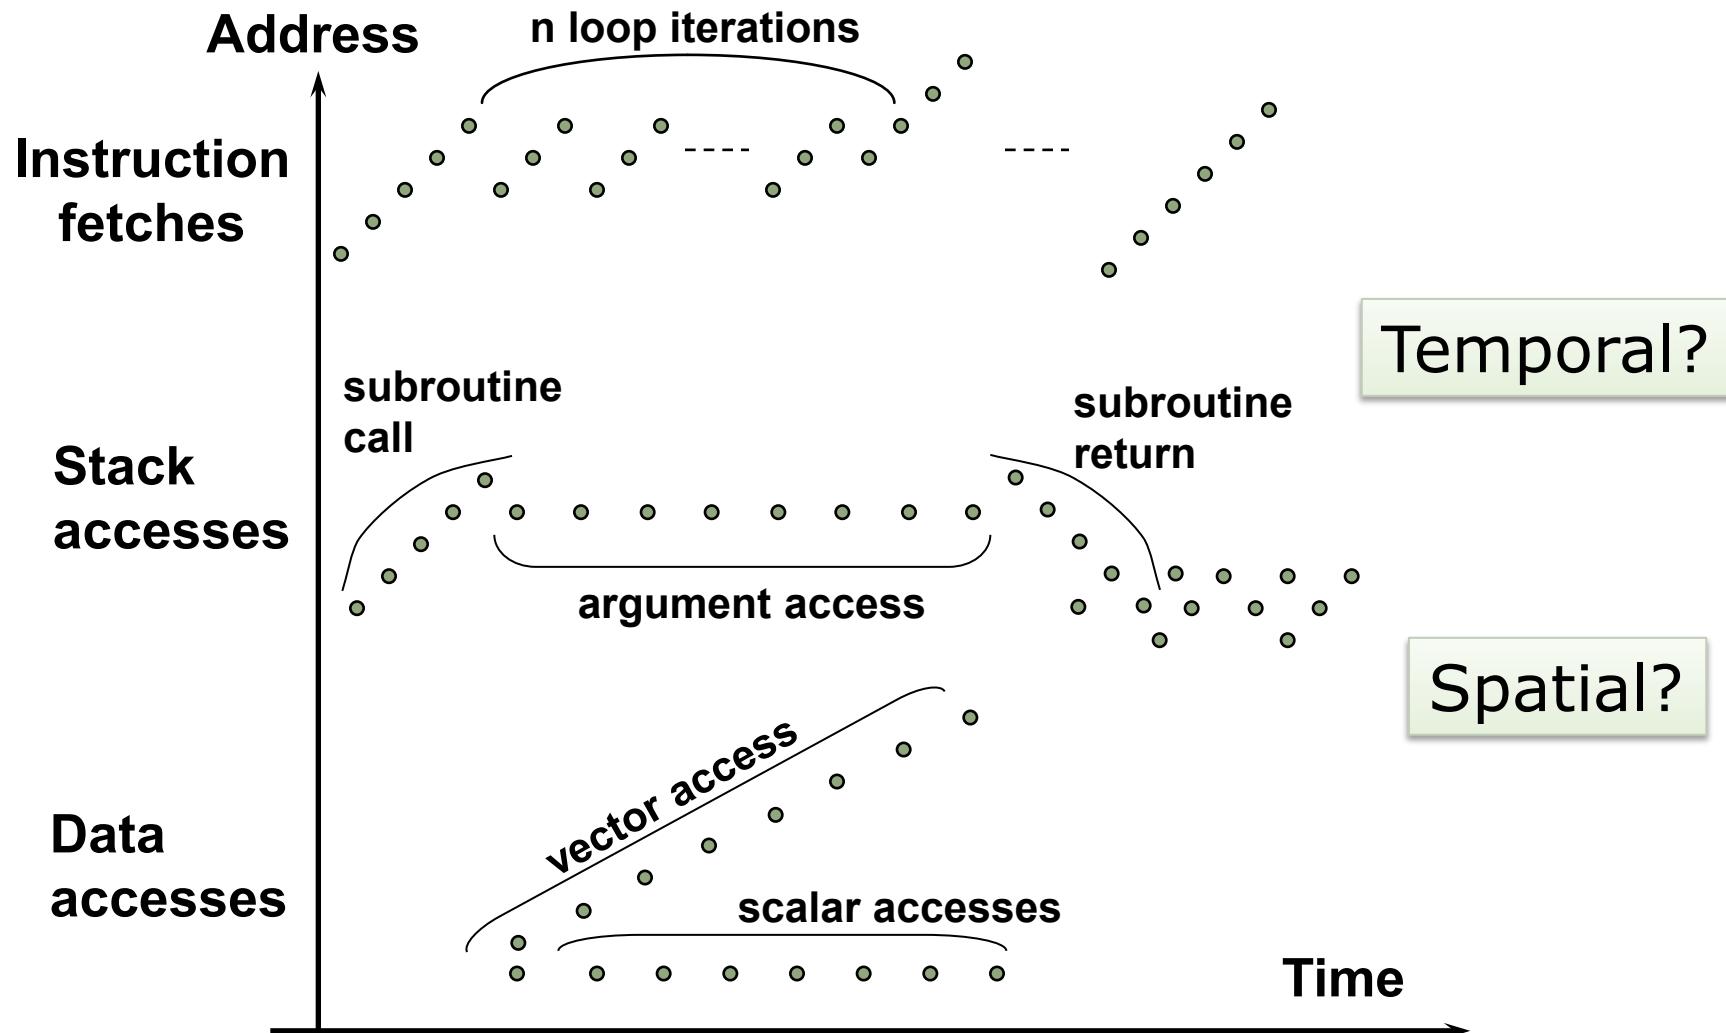


Common Predictable Patterns

Two predictable properties of memory references:

- *Temporal Locality*: If a location is referenced, it is likely to be referenced again in the near future
- *Spatial Locality*: If a location is referenced, it is likely that locations near it will be referenced in the near future

Temporal vs. Spatial Reuse?



Data Orchestration Techniques

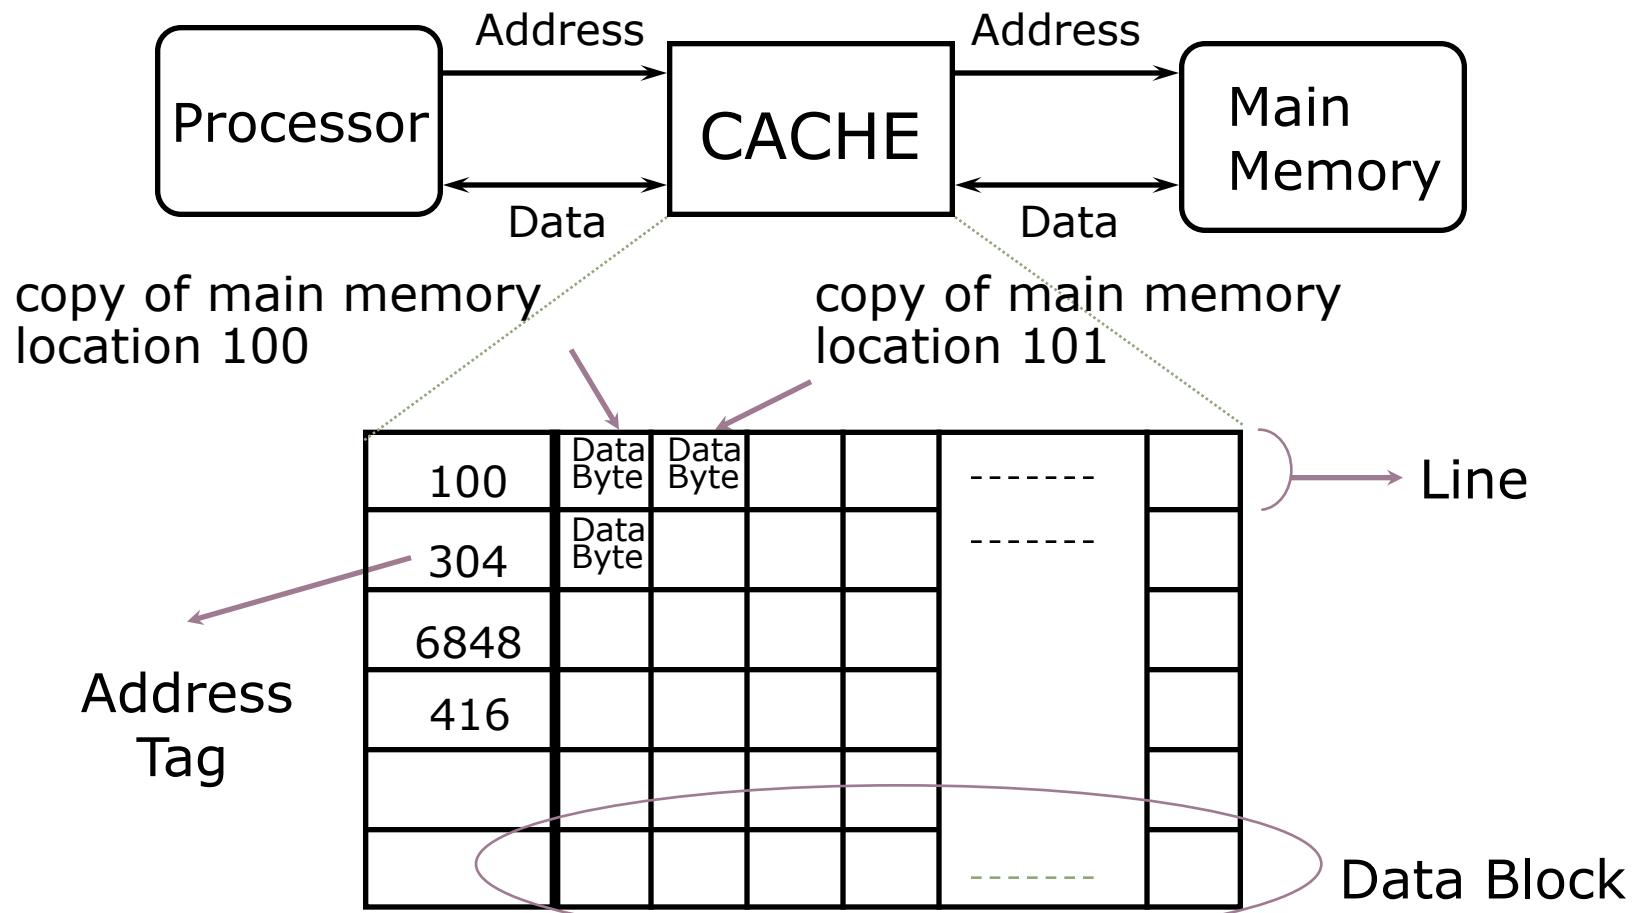
Two approaches to controlling data movement in the memory hierarchy:

- *Explicit*: Manually at the direction of the programmer using instructions
- *Implicit*: Automatically by the hardware in response to a request by an instruction, but transparent to the programmer.

Registers

Cache hierarchy

Inside a Cache



Q: How many bits needed in tag? Enough to uniquely identify block

Cache Algorithm (Read)

Look at Processor Address, search cache tags to find match.
Then either

Found in cache
a.k.a. HIT

Return copy
of data from
cache

Not in cache
a.k.a. MISS

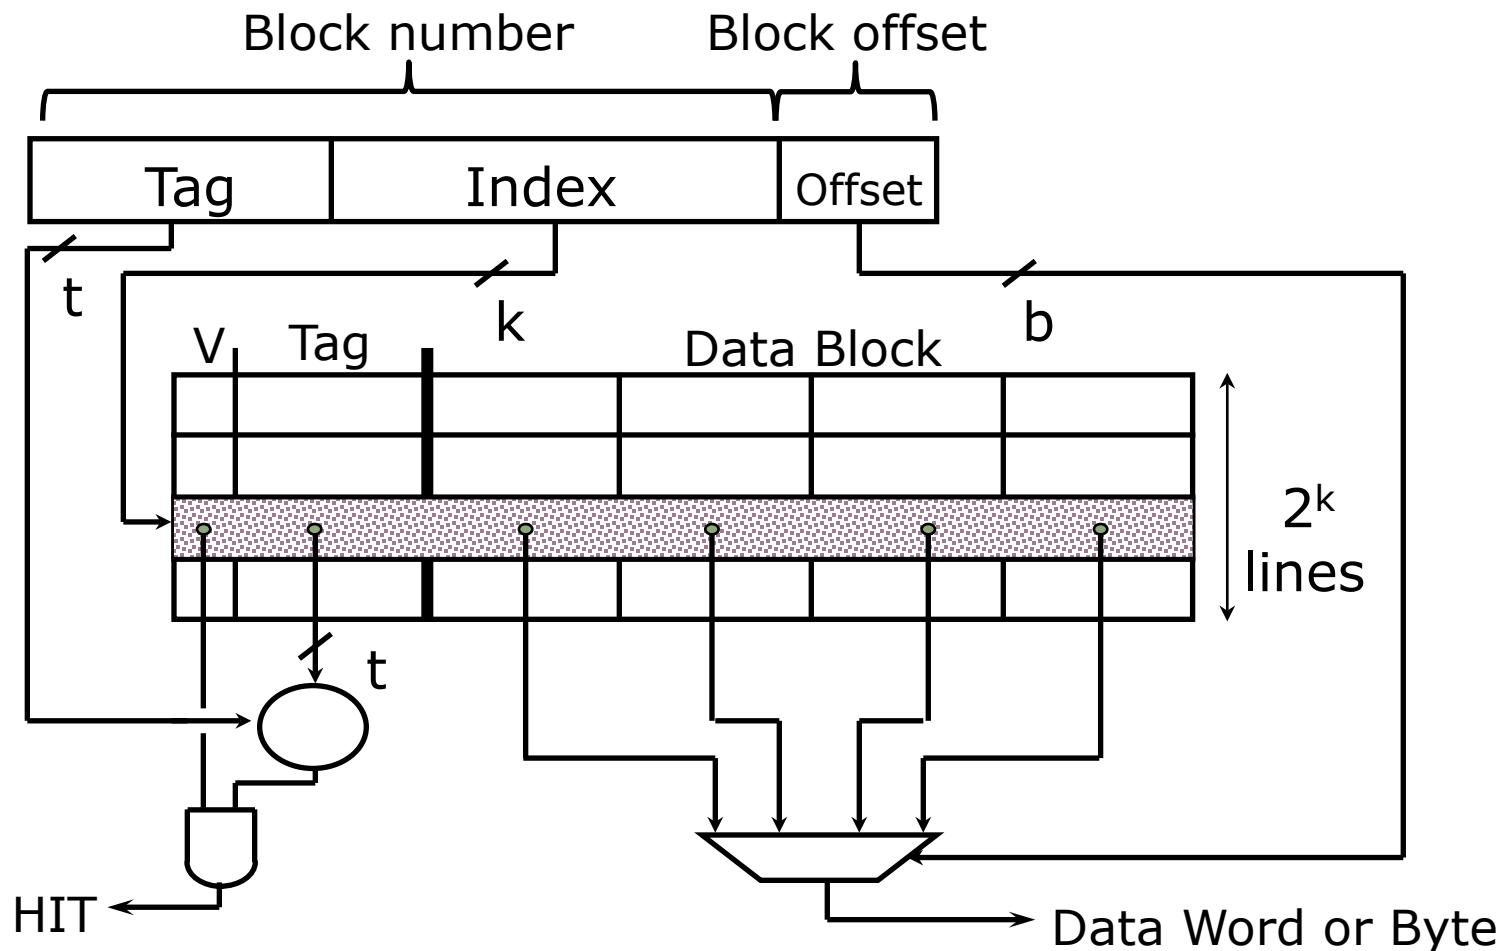
Read block of data from
Main Memory

Wait ...

Return data to processor
and update cache

Which line do we replace?

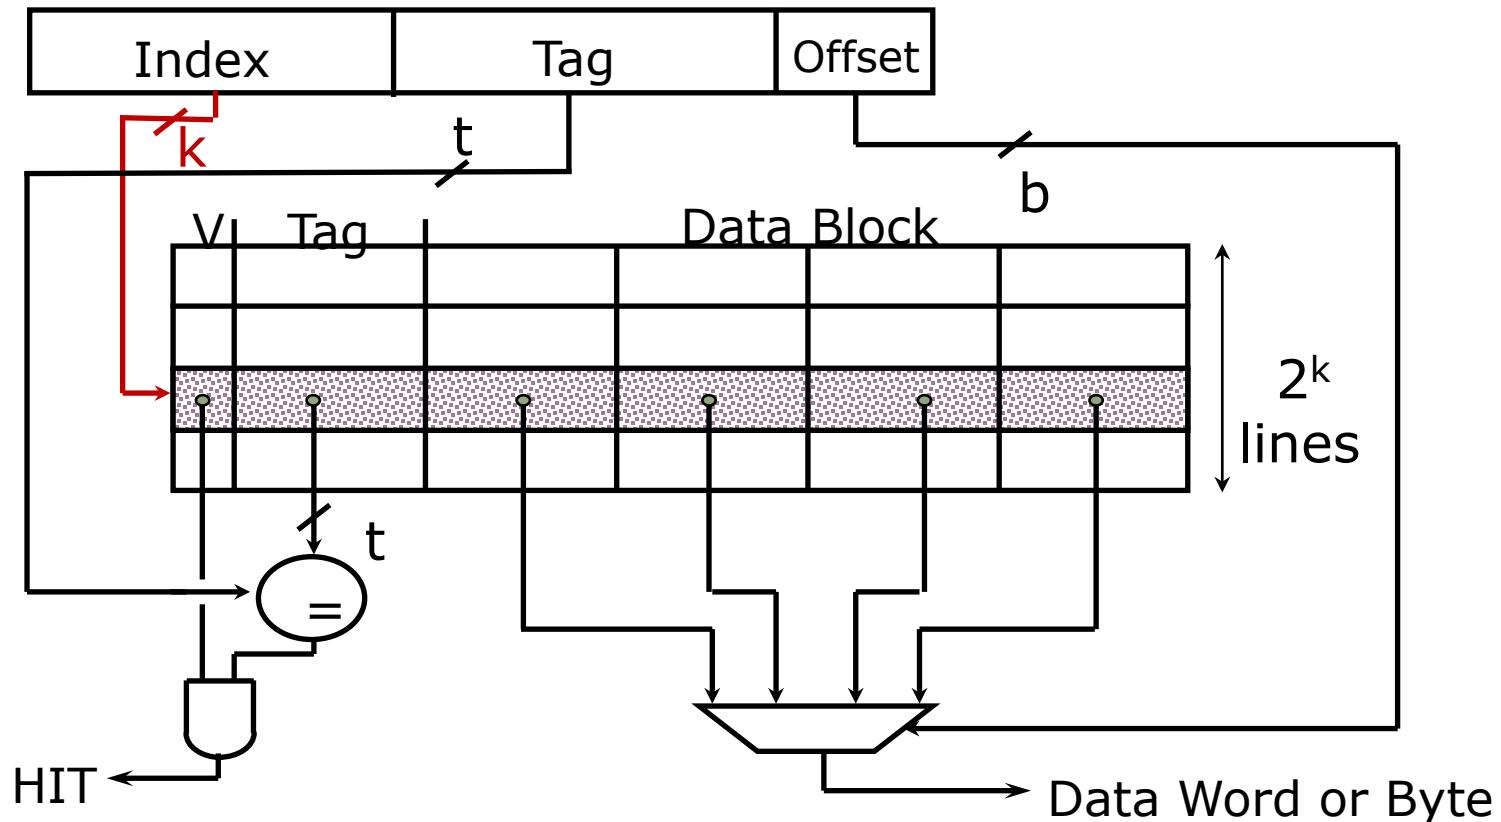
Direct-Mapped Cache



Q: What is a bad reference pattern? Strided at size of cache

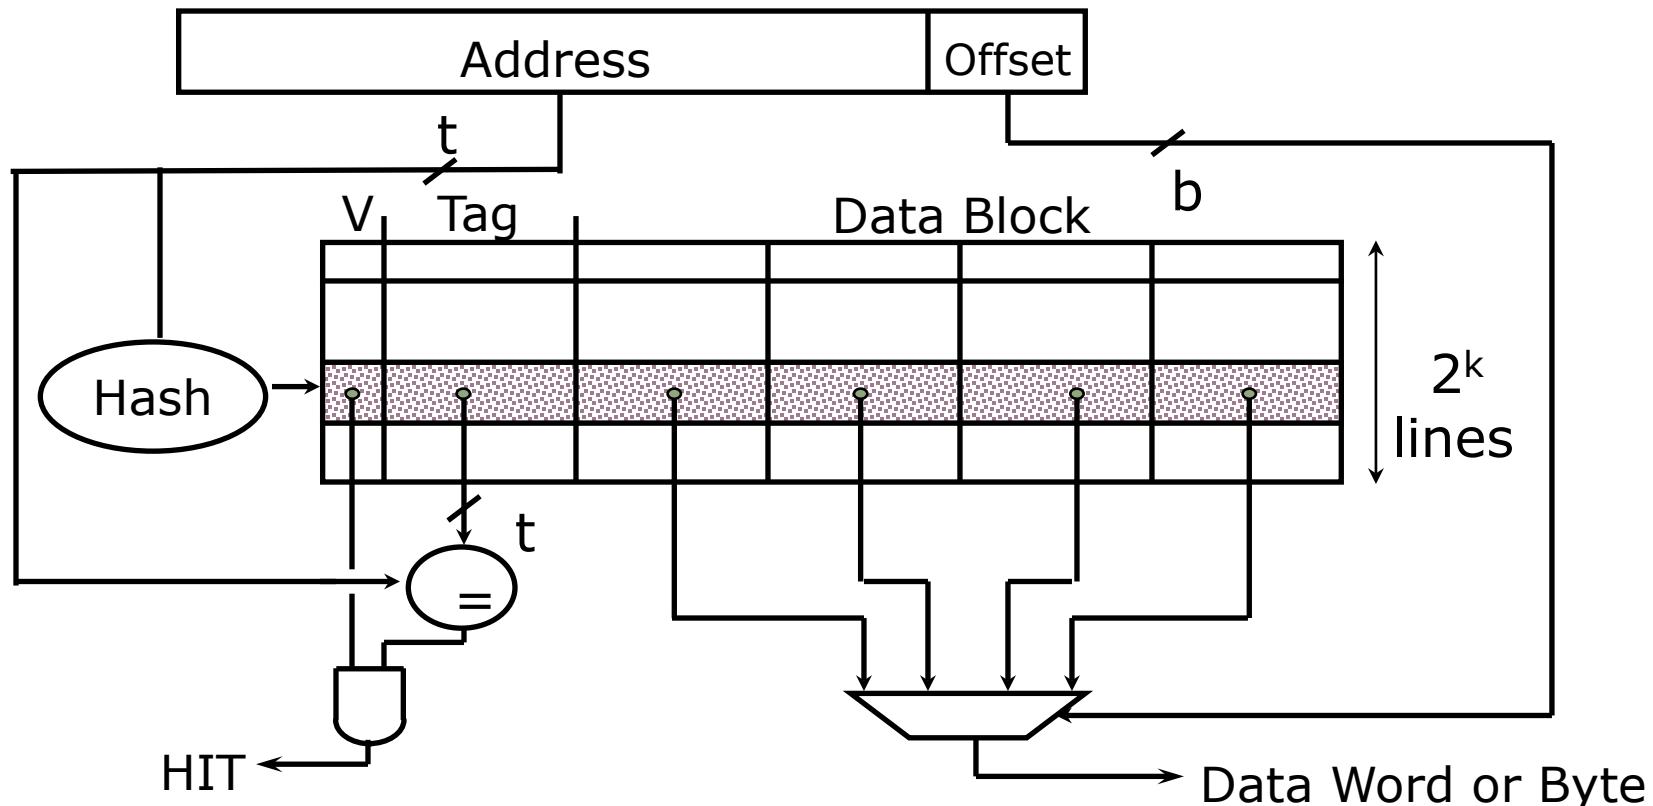
Direct Map Address Selection

higher-order vs. lower-order address bits



Q: Why might this be undesirable? Spatially local blocks conflict

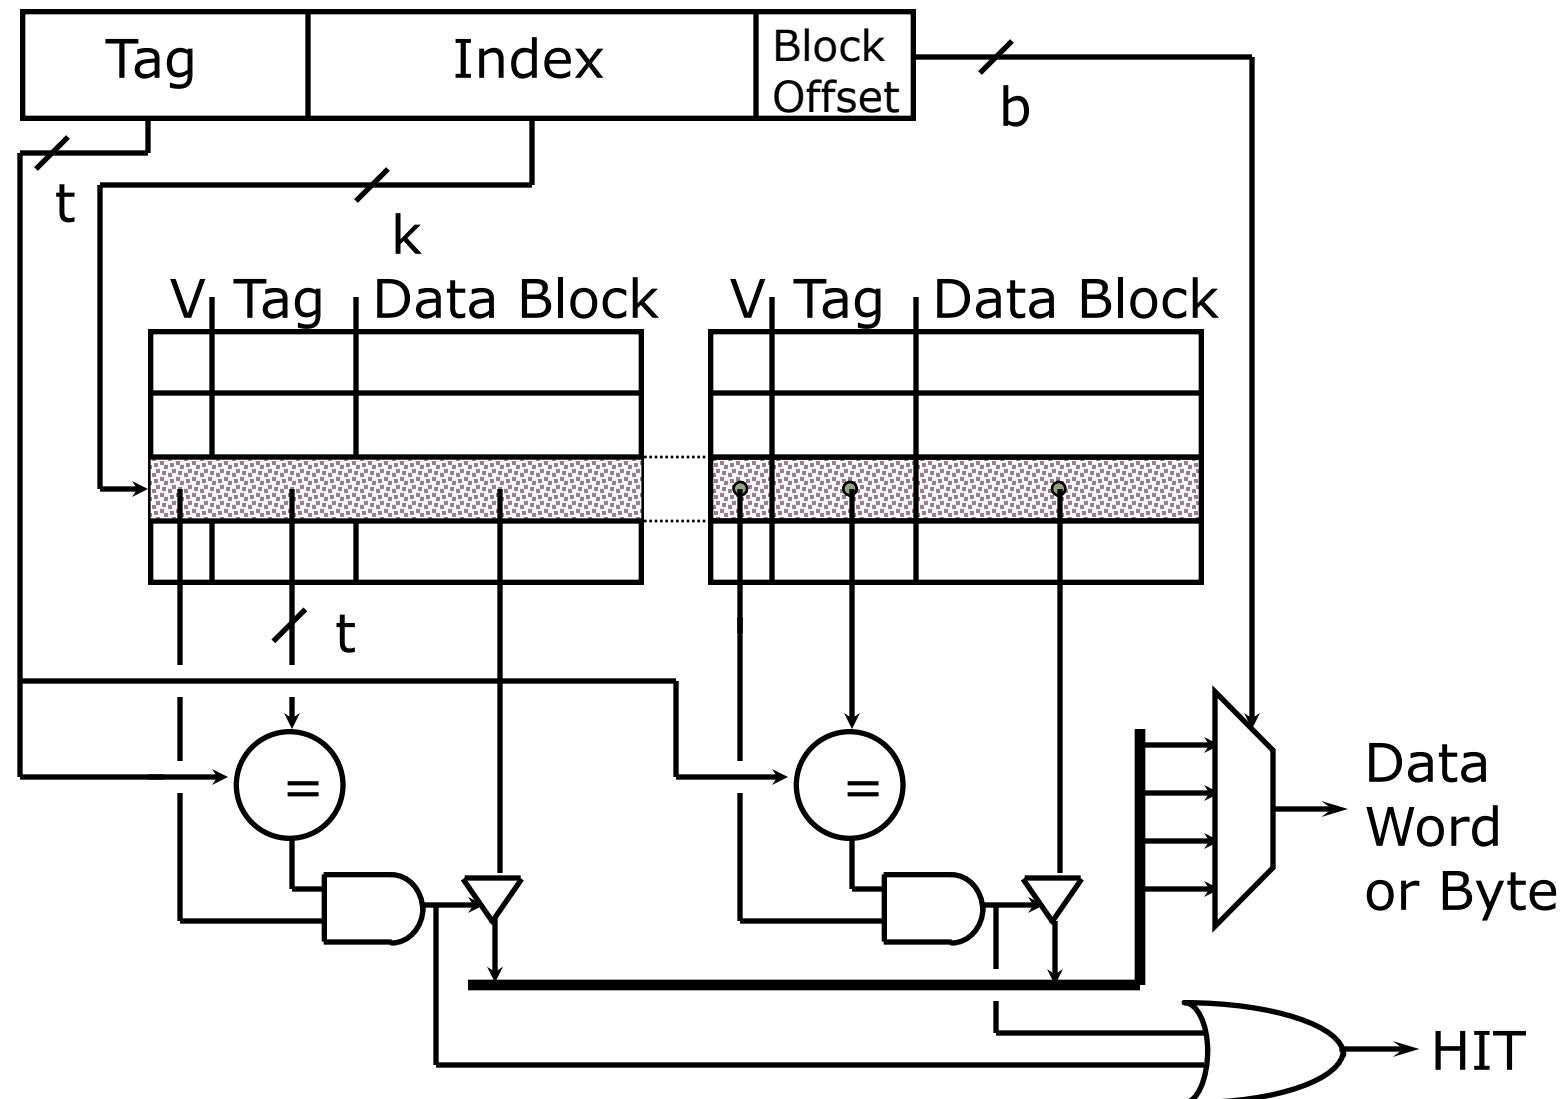
Hashed Address Mapping



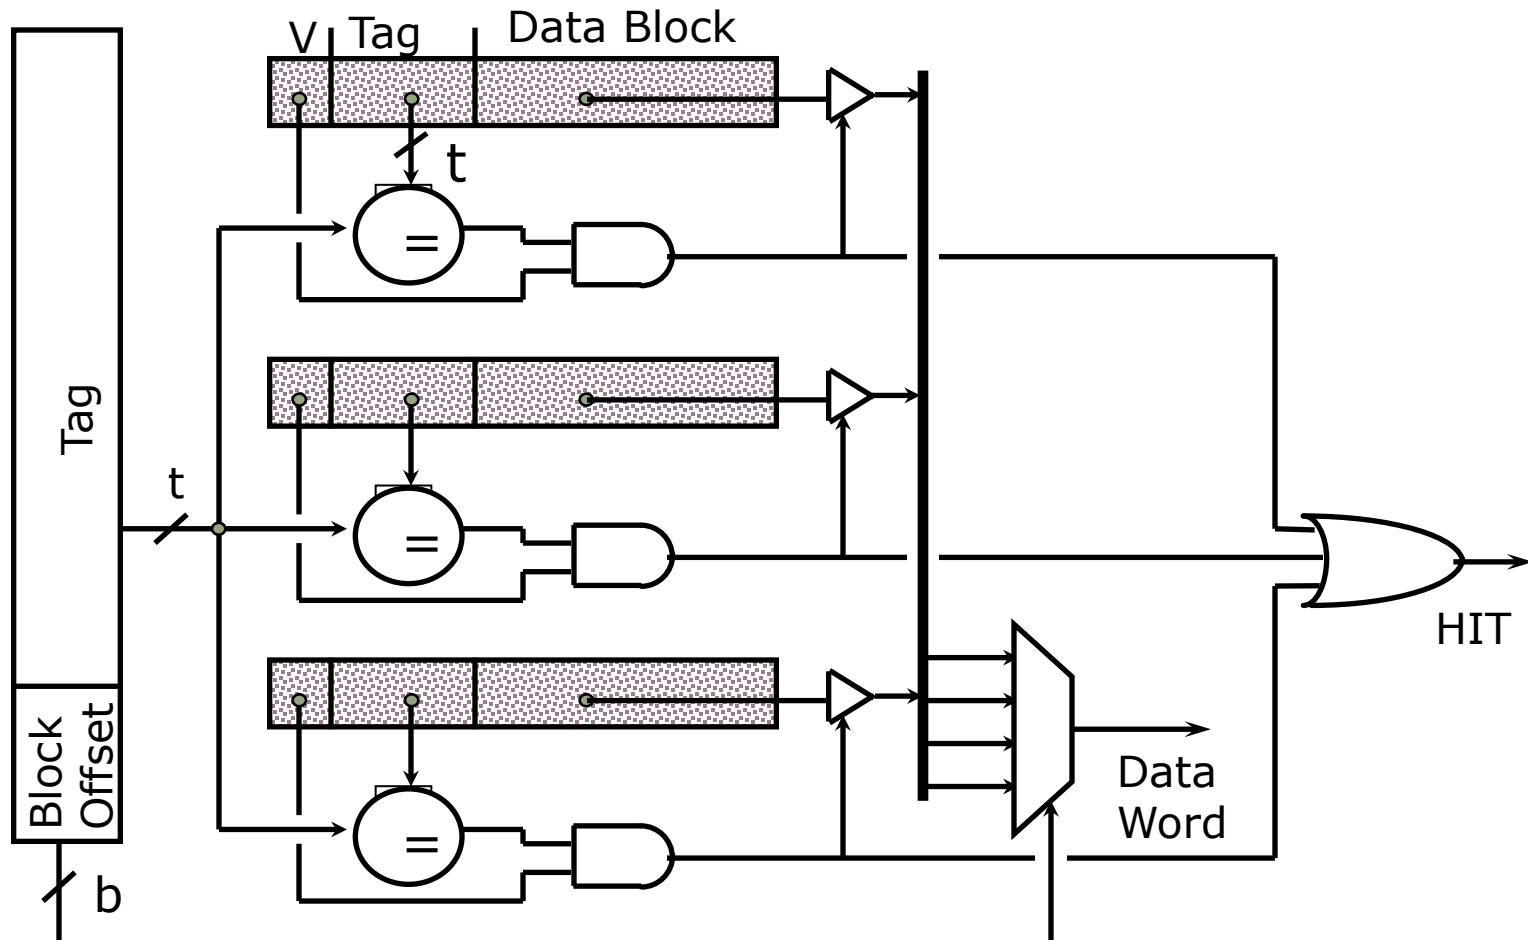
Q: What are the tradeoffs of hashing?

Good: Regular strides don't conflict
Bad: Hash adds latency
Tag is larger

2-Way Set-Associative Cache



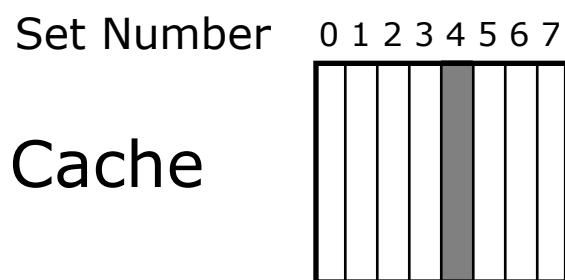
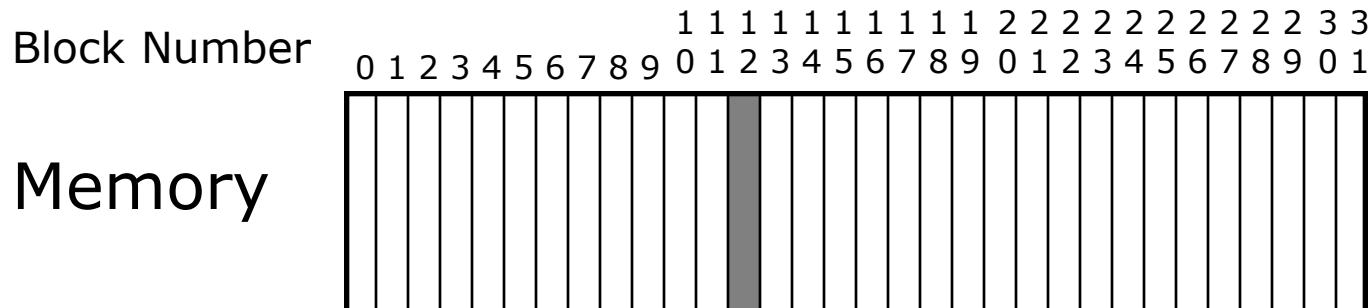
Fully Associative Cache



Q: Where are the index bits?

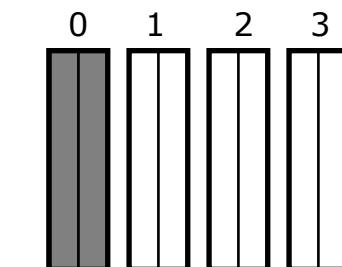
Not needed

Placement Policy

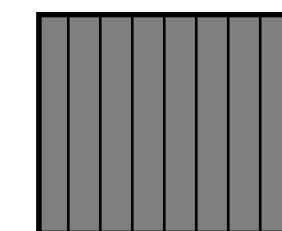


block 12
can be placed

Direct
Mapped
only into
block 4
(12 mod 8)



(2-way) Set
Associative
anywhere in
set 0
(12 mod 4)



Fully
Associative
anywhere

Cache Performance

Average memory access time (AMAT) =
Hit time + Miss rate x Miss penalty

To improve performance:

- reduce the hit time
- reduce the miss rate (e.g., larger, better policy)
- reduce the miss penalty (e.g., L2 cache)

What is the simplest design strategy?

*Biggest cache that doesn't increase hit time past 1-2 cycles
(approx. 16-64KB in modern technology)*

[design issues more complex with out-of-order superscalar processors]

Causes for Cache Misses

- *Compulsory:*
First reference to a block a.k.a. cold start misses
 - misses that would occur even with infinite cache
- *Capacity:*
cache is too small to hold all data the program needs
 - misses that would occur even under fully-associative placement & perfect replacement policy
- *Conflict:*
misses from collisions due to block-placement strategy
 - misses that would not occur with full associativity

Effect of Cache Parameters on Performance

| | Larger capacity cache | Higher associativity cache | Larger block size cache * |
|-------------------|-----------------------|----------------------------|---------------------------|
| Compulsory misses | = | = | ↓ |
| Capacity misses | ↓ | = | ↓ |
| Conflict misses | ↓ | ↓ | ? |
| Hit latency | ↑ | ↑ | = |
| Miss latency | = | = | ↑↓ |

* Assume substantial spatial locality

Thank you!

*Next lecture:
Memory Management*