



# Cache Memories

15-213/14-513/15-513: Introduction to Computer Systems  
10<sup>th</sup> Lecture, Sept 25, 2025

# Today


- Cache memory organization and operation CSAPP 6.4-6.5
- Performance impact of caches
  - The memory mountain CSAPP 6.6.1
  - Rearranging loops to improve spatial locality CSAPP 6.6.2
  - Using blocking to improve temporal locality CSAPP 6.6.3

# Recall: Locality

- **Principle of Locality:** Programs tend to use data and instructions with addresses near or equal to those they have used recently


- **Temporal locality:**

- Recently referenced items are likely to be referenced again in the near future




- **Spatial locality:**


- Items with nearby addresses tend to be referenced close together in time




# Recall: General Cache Concepts



# Recall: Cache Hit



# Recall: Cache Miss



*Data in block b is needed*

*Block b is not in cache:  
Miss!*

*Block b is fetched from  
memory*

*Block b is stored in cache*

- **Placement policy:**  
determines where b goes
- **Replacement policy:**  
determines which block gets evicted (victim)

# Recall: General Caching Concepts:

## 3 Types of Cache Misses

### ■ Cold (compulsory) miss

- Cold misses occur because this is the first reference to the block.  
(Misses with infinitely large cache with no placement restrictions)


### ■ Capacity miss

- Occurs when the set of active cache blocks is larger than the cache.  
(*Additional* misses from finite-sized cache with no placement restrictions)


### ■ Conflict miss

- Occurs when the cache is large enough, but too many data objects all map (by the **placement policy**) to the same limited set of blocks  
(*Additional* misses due to actual placement policy)

# General Cache Organization (S, E, B)




# Cache Read



# Example: Direct Mapped Cache ( $E = 1$ )

Direct mapped: One line per set


Assume: cache block size  $B=8$  bytes



# Example: Direct Mapped Cache ( $E = 1$ )

Direct mapped: One line per set


Assume: cache block size  $B=8$  bytes



# Example: Direct Mapped Cache ( $E = 1$ )

Direct mapped: One line per set

Assume: cache block size  $B=8$  bytes



If tag doesn't match (= miss): old line is evicted and replaced

# Direct-Mapped Cache Simulation

$t=1 \quad s=2 \quad b=1$

x	xx	x
---	----	---

4-bit addresses (address space size  $M=16$  bytes)  
 $S=4$  sets,  $E=1$  Blocks/set,  $B=2$  bytes/block

Address trace (reads, one byte per read):


0	$[0\underline{00}0_2]$ ,	miss
1	$[0\underline{00}1_2]$ ,	hit
7	$[0\underline{11}1_2]$ ,	miss
8	$[1\underline{00}0_2]$ ,	miss
0	$[0\underline{00}0_2]$	miss

	v	Tag	Block
<b>Set 0</b>	1	0	$M[0-1]$
<b>Set 1</b>	0		
<b>Set 2</b>	0		
<b>Set 3</b>	1	0	$M[6-7]$

# E-way Set Associative Cache (Here: E = 2)

$E = 2$ : Two lines per set


Assume: cache block size  $B=8$  bytes



# E-way Set Associative Cache (Here: E = 2)

E = 2: Two lines per set


Assume: cache block size B=8 bytes



# E-way Set Associative Cache (Here: E = 2)

$E = 2$ : Two lines per set

Assume: cache block size  $B=8$  bytes



## No match or not valid (= miss):

- One line in set is selected for eviction and replacement
- Replacement policies: random, least recently used (LRU), ...

# 2-Way Set Associative Cache Simulation

$t=2 \quad s=1 \quad b=1$



4-bit addresses ( $M=16$  bytes)

$S=2$  sets,  $E=2$  blocks/set,  $B=2$  bytes/block

Address trace (reads, one byte per read):

0	$[00\underline{00}_2]$ ,	miss
1	$[00\underline{01}_2]$ ,	hit
7	$[01\underline{11}_2]$ ,	miss
8	$[10\underline{00}_2]$ ,	miss
0	$[00\underline{00}_2]$	hit

	v	Tag	Block
Set 0	1	00	$M[0-1]$
	1	10	$M[8-9]$
Set 1	1	01	$M[6-7]$
	0		

# What about writes?

## ■ Multiple copies of data exist:

- L1, L2, L3, Main Memory

## ■ What to do on a write-hit?



- **Write-through** (write immediately to memory)
- **Write-back** (defer write to memory until replacement of line)
  - Each cache line needs a dirty bit (set if data has been written to)

## ■ What to do on a write-miss?

- **Write-allocate** (load into cache, update line in cache)
  - Good if more writes to the location will follow
- **No-write-allocate** (writes straight to memory, does not load into cache)

## ■ Typical


- Write-through + No-write-allocate
- **Write-back + Write-allocate**



# Practical Write-back Write-allocate

- A write to address X is issued

- If it is a hit
  - Update the contents of block
  - Set dirty bit to 1 (bit is sticky and only cleared on eviction)




- If it is a miss
  - Fetch block from memory (per a read miss)
  - Then perform the write operations (per a write hit)

- If a line is evicted and dirty bit is set to 1
  - The entire block of  $2^b$  bytes are written back to memory
  - Dirty bit is cleared (set to 0)
  - Line is replaced by new contents

# Intel Core i7 Cache Hierarchy

Processor package



**L1 i-cache and d-cache:**

32 KB, 8-way,  
Access: 4 cycles

**L2 unified cache:**

256 KB, 8-way,  
Access: 10 cycles

**L3 unified cache:**

8 MB, 16-way,  
Access: 40-75 cycles

**Block size:** 64 bytes for  
all caches.

# Cache Performance Metrics

## ■ Miss Rate

- Fraction of memory references not found in cache (misses / accesses)  
=  $1 - \text{hit rate}$
- Typical numbers (in percentages):
  - 3-10% for L1
  - can be quite small (e.g., < 1%) for L2, depending on size, etc.

## ■ Hit Time

- Time to deliver a line in the cache to the processor
  - includes time to determine whether the line is in the cache
- Typical numbers:
  - 4 clock cycle for L1
  - 10 clock cycles for L2

## ■ Miss Penalty

- Additional time required because of a miss
  - typically 50-200 cycles for main memory (Trend: increasing!)

# Let's think about those numbers

- Huge difference between a hit and a miss
  - Could be 100x, if just L1 and main memory
- Would you believe 99% hits is twice as good as 97%?
  - Consider this simplified example:  
cache hit time of 1 cycle  
miss penalty of 100 cycles
  - Average access time:  
97% hits: 1 cycle +  $0.03 \times 100$  cycles = **4 cycles**  
99% hits: 1 cycle +  $0.01 \times 100$  cycles = **2 cycles**
- This is why “miss rate” is used instead of “hit rate”

# Writing Cache Friendly Code

- **Make the common case go fast**
  - Focus on the inner loops of the core functions
- **Minimize the misses in the inner loops**
  - Repeated references to variables are good (**temporal locality**)
  - Stride-1 reference patterns are good (**spatial locality**)

**Key idea: Our qualitative notion of locality is quantified through our understanding of cache memories**

# Quiz

<https://canvas.cmu.edu/courses/49105/quizzes/150049>

# Today

- Cache organization and operation
- **Performance impact of caches**
  - The memory mountain
  - Rearranging loops to improve spatial locality
  - Using blocking to improve temporal locality

# The Memory Mountain

- **Read throughput (read bandwidth)**
  - Number of bytes read from memory per second (MB/s)
- **Memory mountain: Measured read throughput as a function of spatial and temporal locality.**
  - Compact way to characterize memory system performance.

# Memory Mountain Test Function

```

long data[MAXELEMS]; /* Global array to traverse */

/* test - Iterate over first "elems" elements of
 *         array "data" with stride of "stride",
 *         using 4x4 loop unrolling.
 */
int test(int elems, int stride) {
    long i, sx2=stride*2, sx3=stride*3, sx4=stride*4;
    long acc0 = 0, acc1 = 0, acc2 = 0, acc3 = 0;
    long length = elems, limit = length - sx4;

    /* Combine 4 elements at a time */
    for (i = 0; i < limit; i += sx4) {
        acc0 = acc0 + data[i];
        acc1 = acc1 + data[i+stride];
        acc2 = acc2 + data[i+sx2];
        acc3 = acc3 + data[i+sx3];
    }

    /* Finish any remaining elements */
    for (; i < length; i++) {
        acc0 = acc0 + data[i];
    }
    return ((acc0 + acc1) + (acc2 + acc3));
}

```

*mountain/mountain.c*


Call `test()` with many combinations of `elems` and `stride`.

For each `elems` and `stride`:

1. Call `test()` once to warm up the caches.
2. Call `test()` again and measure the read throughput (MB/s)

# The Memory Mountain

Aggressive prefetching




Core i7 Haswell  
2.1 GHz  
32 KB L1 d-cache  
256 KB L2 cache  
8 MB L3 cache  
64 B block size

# Today

- Cache organization and operation
- Performance impact of caches
  - The memory mountain
  - Rearranging loops to improve spatial locality
  - Using blocking to improve temporal locality

# Remember matrix multiplication



$\text{Out}[i, j] =$   
dot product( $A[i, ..], B[.., j]$ )  
= sum (  
     $a[i, 0] * b[0, j],$   
     $a[i, 1] * b[1, j]$   
)

# Matrix Multiplication Example

## ■ Description:

- Multiply  $N \times N$  matrices
- Matrix elements are doubles (8 bytes)
- $O(N^3)$  total operations
- $N$  reads per source element
- $N$  values summed per destination
  - but may be able to hold in register

```
/* ijk */
for (i=0; i<n; i++) t
    for (j=0; j<n; j++) {
        sum = 0.0; ←
        for (k=0; k<n; k++)
            sum += a[i][k] * b[k][j];
        c[i][j] = sum;
    }
}
```

*Variable sum held in register*

*matmult/mm.c*


# Miss Rate Analysis for Matrix Multiply

## ■ Assume:

- Block size =  $32B$  (big enough for four doubles)
- Matrix dimension ( $N$ ) is very large
  - Approximate  $1/N$  as 0.0
- Cache is not even big enough to hold multiple rows


## ■ Analysis Method:

- Look at access pattern of inner loop



# Layout of C Arrays in Memory (review)

## ■ C arrays allocated in row-major order



## ■ Stepping through columns in one row:

- `for (i = 0; i < N; i++)  
 sum += a[0][i]`
- if block size ( $B$ ) > `sizeof(aij)` bytes, exploit spatial locality
  - miss rate =  $\text{sizeof}(a_{ij}) / B$

## ■ Stepping through rows in one column:


- `for (i = 0; i < M; i++)  
 sum += a[i][0];`
- accesses distant elements: no spatial locality!
  - miss rate = 1 (i.e. 100%)

# Matrix Multiplication (ijk)

```
/* ijk */
for (i=0; i<n; i++) {
    for (j=0; j<n; j++) {
        sum = 0.0;
        for (k=0; k<n; k++)
            sum += a[i][k] * b[k][j];
        c[i][j] = sum;
    }
}
```

*matmult/mm.c*

Inner loop:



Miss rate for inner loop iterations:

A

B

C


**Block size = 32B (four doubles)**

# Matrix Multiplication (ijk)

```
/* ijk */
for (i=0; i<n; i++) {
    for (j=0; j<n; j++) {
        sum = 0.0;
        for (k=0; k<n; k++)
            sum += a[i][k] * b[k][j];
        c[i][j] = sum;
    }
}
```

*matmult/mm.c*

Inner loop:



Miss rate for inner loop iterations:

<u>A</u>	<u>B</u>	<u>C</u>
0.25	1.0	0.0


**Block size = 32B (four doubles)**

# Matrix Multiplication ( $kij$ )

```
/* kij */
for (k=0; k<n; k++) {
    for (i=0; i<n; i++) {
        r = a[i][k];
        for (j=0; j<n; j++)
            c[i][j] += r * b[k][j];
    }
}
```

*matmult/mm.c*

Inner loop:



Miss rate for inner loop iterations:

A

B

C


**Block size = 32B (four doubles)**

# Matrix Multiplication ( $kij$ )

```
/* kij */
for (k=0; k<n; k++) {
    for (i=0; i<n; i++) {
        r = a[i][k];
        for (j=0; j<n; j++)
            c[i][j] += r * b[k][j];
    }
}
```

*matmult/mm.c*

Inner loop:



Miss rate for inner loop iterations:

<u>A</u>	<u>B</u>	<u>C</u>
0.0	0.25	0.25


**Block size = 32B (four doubles)**

# Matrix Multiplication (jki)

```
/* jki */
for (j=0; j<n; j++) {
    for (k=0; k<n; k++) {
        r = b[k][j];
        for (i=0; i<n; i++)
            c[i][j] += a[i][k] * r;
    }
}
```

*matmult/mm.c*

Inner loop:



Miss rate for inner loop iterations:

A

B

C


**Block size = 32B (four doubles)**

# Matrix Multiplication (jki)

```
/* jki */
for (j=0; j<n; j++) {
    for (k=0; k<n; k++) {
        r = b[k][j];
        for (i=0; i<n; i++)
            c[i][j] += a[i][k] * r;
    }
}
```

*matmult/mm.c*

Inner loop:



Miss rate for inner loop iterations:

A  
1.0

B  
0.0

C  
1.0

**Block size = 32B (four doubles)**

# Summary of Matrix Multiplication

```

for (i=0; i<n; i++) {
    for (j=0; j<n; j++) {
        sum = 0.0;
        for (k=0; k<n; k++)
            sum += a[i][k] * b[k][j];
        c[i][j] = sum;
    }
}

```

```

for (k=0; k<n; k++) {
    for (i=0; i<n; i++) {
        r = a[i][k];
        for (j=0; j<n; j++)
            c[i][j] += r * b[k][j];
    }
}

```

```

for (j=0; j<n; j++) {
    for (k=0; k<n; k++) {
        r = b[k][j];
        for (i=0; i<n; i++)
            c[i][j] += a[i][k] * r;
    }
}

```

## **ijk (& jik):**

- 2 loads, 0 stores
- avg misses/iter = **1.25**

## **kij (& ikj):**

- 2 loads, 1 store
- avg misses/iter = **0.5**

## **jki (& kji):**

- 2 loads, 1 store
- avg misses/iter = **2.0**

# Core i7 Matrix Multiply Performance

Cycles per inner loop iteration

100

jki / kji (2.0)

- jki
- kji
- ijk
- jik
- kij
- ikj

10

ijk / jik (1.25)

1

kij / ikj (0.5)




# Today

- Cache organization and operation
- Performance impact of caches
  - The memory mountain
  - Rearranging loops to improve spatial locality
  - Using blocking to improve temporal locality

# Example: Matrix Multiplication

```
c = (double *) calloc(sizeof(double), n*n);  
  
/* Multiply n x n matrices a and b */  
void mmm(double *a, double *b, double *c, int n) {  
    int i, j, k;  
    for (i = 0; i < n; i++)  
        for (j = 0; j < n; j++)  
            for (k = 0; k < n; k++)  
                c[i*n + j] += a[i*n + k] * b[k*n + j];  
}
```




# Cache Miss Analysis

## ■ Assume:

- Matrix elements are doubles
- Cache block = 8 doubles
- Cache size  $C \ll n$  (much smaller than  $n$ )

## ■ First iteration:

- $n/8 + n = 9n/8$  misses



- Afterwards **in cache:**  
(schematic)


# Cache Miss Analysis

## ■ Assume:

- Matrix elements are doubles
- Cache block = 8 doubles
- Cache size  $C \ll n$  (much smaller than  $n$ )

## ■ Second iteration:

- Again:  
 $n/8 + n = 9n/8$  misses




## ■ Total misses:

- $9n/8 n^2 = (9/8) n^3$

# Blocked Matrix Multiplication

```
c = (double *) calloc(sizeof(double), n*n);

/* Multiply n x n matrices a and b */
void mmm(double *a, double *b, double *c, int n) {
    int i, j, k;
    for (i = 0; i < n; i+=B)
        for (j = 0; j < n; j+=B)
            for (k = 0; k < n; k+=B)
                /* B x B mini matrix multiplications */
                for (i1 = i; i1 < i+B; i1++)
                    for (j1 = j; j1 < j+B; j1++)
                        for (k1 = k; k1 < k+B; k1++)
                            c[i1*n+j1] += a[i1*n + k1]*b[k1*n + j1];
}
                                            matmult/bmm.c
```



Notation Note  
This “B” is not the cache block size B


# Cache Miss Analysis

## ■ Assume:


- Cache block = 8 doubles
- Cache size  $C \ll n$  (much smaller than  $n$ )
- Three blocks  fit into cache:  $3B^2 < C$

## ■ First (block) iteration:

- $B^2/8$  misses for each block
- $2n/B \times B^2/8 = nB/4$   
(omitting matrix  $c$ )



- Afterwards in cache  
(schematic)




# Cache Miss Analysis

## ■ Assume:

- Cache block = 8 doubles
- Cache size  $C \ll n$  (much smaller than  $n$ )
- Three blocks ■ fit into cache:  $3B^2 < C$

## ■ Second (block) iteration:

- Same as first iteration
- $2n/B \times B^2/8 = nB/4$



## ■ Total misses:

- $nB/4 * (n/B)^2 = n^3/(4B)$

# Blocking Summary


- No blocking:  $(9/8) n^3$  misses
- Blocking:  $(1/(4B)) n^3$  misses
- Use largest block size  $B$ , such that  $B$  satisfies  $3B^2 < C$ 
  - Fit three blocks in cache! Two input, one output.
- Reason for dramatic difference:
  - Matrix multiplication has inherent temporal locality:
    - Input data:  $3n^2$ , computation  $2n^3$
    - Every array elements used  $O(n)$  times!
  - But program has to be written properly

# Cache Summary


- Cache memories can have significant performance impact
- You can write your programs to exploit this!
  - Focus on the inner loops, where bulk of computations and memory accesses occur.
  - Try to maximize spatial locality by reading data objects sequentially with stride 1.
  - Try to maximize temporal locality by using a data object as often as possible once it's read from memory.

# Supplemental slides

# Recall: Modern CPU Design




# What it Really Looks Like




Nehalem


AMD FX 8150




CPU chip



Core i7-3960X



# What it Really Looks Like (Cont.)




Intel Sandy Bridge  
Processor Die

- L1: 32KB Instruction + 32KB Data
- L2: 256KB
- L3: 3–20MB

# Why Index Using Middle Bits?


Direct mapped: One line per set

Assume: cache block size 8 bytes



**Standard Method:  
Middle bit indexing**

Address of int:



find set

**Alternative Method:  
High bit indexing**


Address of int:



find set

# Illustration of Indexing Approaches

- **64-byte memory**
  - 6-bit addresses
- **16 byte, direct-mapped cache**
- **Block size = 4. (Thus, 4 sets; why?)**
- **2 bits tag, 2 bits index, 2 bits offset**




				0000xx
				0001xx
				0010xx
				0011xx
				0100xx
				0101xx
				0110xx
				0111xx
				1000xx
				1001xx
				1010xx
				1011xx
				1100xx
				1101xx
				1110xx
				1111xx

# Middle Bit Indexing

- Addresses of form **TTSSBB**

- **TT** Tag bits
- **SS** Set index bits
- **BB** Offset bits

- Makes good use of spatial locality





0000xx	Green	Green	Green	Green
0001xx	Yellow	Yellow	Yellow	Yellow
0010xx	Red	Red	Red	Red
0011xx	Blue	Blue	Blue	Blue
0100xx	Green	Green	Green	Green
0101xx	Yellow	Yellow	Yellow	Yellow
0110xx	Red	Red	Red	Red
0111xx	Blue	Blue	Blue	Blue
1000xx	Green	Green	Green	Green
1001xx	Yellow	Yellow	Yellow	Yellow
1010xx	Red	Red	Red	Red
1011xx	Blue	Blue	Blue	Blue
1100xx	Green	Green	Green	Green
1101xx	Yellow	Yellow	Yellow	Yellow
1110xx	Red	Red	Red	Red
1111xx	Blue	Blue	Blue	Blue

# High Bit Indexing

## ■ Addresses of form **SSTTBB**

- **SS** Set index bits
- **TT** Tag bits
- **BB** Offset bits

## ■ Program with high spatial locality would generate lots of conflicts



# Example: Core i7 L1 Data Cache

**32 kB 8-way set associative**

**64 bytes/block**


**47 bit address range**

**B =**

**S = , s =**


**E = , e =**

**C =**



Hex	Decimal	Binary
0	0	0000
1	1	0001
2	2	0010
3	3	0011
4	4	0100
5	5	0101
6	6	0110
7	7	0111
8	8	1000
9	9	1001
A	10	1010
B	11	1011
C	12	1100
D	13	1101
E	14	1110
F	15	1111

**Address of word:**



**Block offset:** . bits

**Set index:** . bits

**Tag:** . bits

**Stack Address:**  
**0x00007f7262a1e010**

**Block offset:** **0x??**  
**Set index:** **0x??**  
**Tag:** **0x??**

# Example: Core i7 L1 Data Cache

**32 kB 8-way set associative**

**64 bytes/block**


**47 bit address range**

$$B = 64$$

$$S = 64, s = 6$$


$$E = 8, e = 3$$

$$C = 64 \times 64 \times 8 = 32,768$$



Hex	Decimal	Binary
0	0	0000
1	1	0001
2	2	0010
3	3	0011
4	4	0100
5	5	0101
6	6	0110
7	7	0111
8	8	1000
9	9	1001
A	10	1010
B	11	1011
C	12	1100
D	13	1101
E	14	1110
F	15	1111

**Address of word:**



**Block offset:** 6 bits

**Set index:** 6 bits

**Tag:** 35 bits

**Stack Address:**

**0x00007f7262a1e010**

0000 0001 0000

**Block offset:**

**0x10**

**Set index:**

**0x0**


**Tag:**

**0x7f7262a1e**

# The Memory Mountain


Aggressive  
prefetching

Core i5 Haswell  
3.1 GHz  
32 KB L1 d-cache  
256 KB L2 cache  
8 MB L3 cache  
64 B block size



# Cache Capacity Effects from Memory Mountain

Core i7 Haswell  
3.1 GHz  
32 KB L1 d-cache  
256 KB L2 cache  
8 MB L3 cache  
64 B block size




Slice through  
memory  
mountain with  
stride=8

# Cache Block Size Effects from Memory Mountain

Core i7 Haswell  
2.26 GHz  
32 KB L1 d-cache  
256 KB L2 cache  
8 MB L3 cache  
64 B block size


Throughput for size = 128K




# Modeling Block Size Effects from Memory Mountain

Core i7 Haswell  
2.26 GHz  
32 KB L1 d-cache  
256 KB L2 cache  
8 MB L3 cache  
64 B block size

Throughput for size = 128K



# 2008 Memory Mountain




# Matrix Multiplication (jik)

```
/* jik */
for (j=0; j<n; j++) {
    for (i=0; i<n; i++) {
        sum = 0.0;
        for (k=0; k<n; k++)
            sum += a[i][k] * b[k][j];
        c[i][j] = sum
    }
}
```

*matmult/mm.c*

Inner loop:



Misses per inner loop iteration:

<u>A</u>	<u>B</u>	<u>C</u>
0.25	1.0	0.0


**Block size = 32B (four doubles)**

# Matrix Multiplication (ikj)

```
/* ikj */
for (i=0; i<n; i++) {
    for (k=0; k<n; k++) {
        r = a[i][k];
        for (j=0; j<n; j++)
            c[i][j] += r * b[k][j];
    }
}
```

*matmult/mm.c*

Inner loop:



Misses per inner loop iteration:


<u>A</u>	<u>B</u>	<u>C</u>
0.0	0.25	0.25

**Block size = 32B (four doubles)**

# Matrix Multiplication (kji)

```
/* kji */
for (k=0; k<n; k++) {
    for (j=0; j<n; j++) {
        r = b[k][j];
        for (i=0; i<n; i++)
            c[i][j] += a[i][k] * r;
    }
}
matmult/mm.c
```

Inner loop:



Misses per inner loop iteration:

A	B	C
1.0	0.0	1.0

**Block size = 32B (four doubles)**