


# RAID and the Warehouse

# Magnetic Disk – common I/O device


- A kind of computer memory
  - Information stored by magnetizing ferrite material on surface of rotating disk
    - Similar to tape recorder except digital rather than analog data
- A type of *non-volatile* storage
  - Retains its value without applying power to disk.
- Two Types of Magnetic Disk
  - Hard Disk Drives (HDD) – faster, more dense, non-removable.
  - Floppy disks – slower, less dense, removable  
(now replaced by USB “flash drive”, only roll these days is as the “Save Icon”).
- Purpose in computer systems (Hard Drive):
  - Working file system + long-term backup for files
  - Secondary “backing store” for main-memory. Large, inexpensive, slow level in the memory hierarchy (virtual memory)



# Photo of Disk Head, Arm, Actuator





# Disk Device Terminology




- Several platters, with information recorded magnetically on both surfaces (usually)
- Bits recorded in tracks, which in turn divided into sectors (e.g., 512 Bytes)
- Actuator moves head (end of arm) over track (“seek”), wait for sector rotate under head, then read or write

# Hard Drives are Sealed

- The closer the head to the disk, the smaller the “spot size” and thus the denser the recording.
  - Modern drives can store up to 20 TB of data
- Disks are sealed to keep the dust out.
  - Heads are designed to “fly” at around 3-20nm above the surface of the disk.
  - 99.999% of the head/arm weight is supported by the air bearing force (air cushion) developed between the disk and the head
- Some drives are even sealed with Helium
  - Lower drag than air




# Disk Device Performance




- **Disk Access Time = Seek Time + Rotation Time + Transfer Time + Controller Overhead**
  - Seek Time = time to position the head assembly at the proper cylinder
  - Rotation Time = time for the disk to rotate to the point where the first sectors of the block to access reach the head
  - Transfer Time = time taken by the sectors of the block and any gaps between them to rotate past the head

# Flash Memory / SSD Technology




Multi-level NAND flash memory stores 3 bits of data in each transistor.

- NMOS transistor with an additional conductor between gate and source/drain which “traps” electrons. The presence/absence is a 1 or 0
- Memory cells can withstand a limited number of program-erase cycles. Controllers use a technique called *wear leveling* to distribute writes as evenly as possible across all the flash blocks in the SSD.
- Erase or write a block, no way to change a block







In the basic functional block used in multilevel NAND flash memories, 32 rows of bit lines and 32 control-gate lines form a building block that's repeated many times to form the memory array. The select gate lines are used with the control gate lines to control access to the array.

# Flash Memory Key to Success of Smart Phones



# Flash Memory in Laptops – Solid State Drive (SSD)




# Flash and Latency...

- Flash bandwidth is similar to spinning disk
  - And spinning disk is still a better storage/\$ and storage/cm<sup>3</sup>
- But Flash's big advantage: ***no seek time!***
  - No additional latency for random access vs sequential access of a block
- This is huge:
  - HDD access times are measured in milliseconds, SSD times are measured in microseconds
- ***NEVER*** put your main OS on a spinning disk!

# RAID: Redundancy for Disks

- Spinning disk is still a critical technology
  - Although worse latency than SSD...
- Disk has equal or greater bandwidth and an order of magnitude better storage density ( $\text{bits/cm}^3$ ) and cost density ( $\text{bits}/\$$ )
- So when you need to store a petabyte or three...
  - You need to use disk, not SSDs
- Oh, and SSDs can fail too

# Evolution of the Disk Drive



IBM RAMAC 305, 1956

First commercial computer that used a moving-head hard disk drive (magnetic disk storage) for secondary storage.



up to 22.7 billion bytes (gigabytes) of storage

IBM 3390K, 1989




Apple SCSI, 1986

# Arrays of Small Disks

Can smaller disks be used to close gap in performance between disks and CPUs?


Conventional:  
4 disk designs



Low End → High End

Disk Array:  
1 disk design

3.5" →



# Replace Small Number of Large Disks in 1988

	IBM 3390K	IBM 3.5" 0061	x70
Capacity	20 GBytes	320 MBytes	23 GBytes
Volume	97 cu. ft.	0.1 cu. ft.	11 cu. ft. <span style="color:red">9X</span>
Power	3 KW	11 W	1 KW <span style="color:red">3X</span>
Data Rate	15 MB/s	1.5 MB/s	105 MB/s <span style="color:red">7X</span>
I/O Rate	600 I/Os/s	55 I/Os/s	3900 IOs/s <span style="color:red">6X</span>
MTTF	250 K Hrs	50 K Hrs	??? Hrs
Cost	\$250K	\$2K	\$150K

Disk Arrays have potential for large data and I/O rates, high MB per cu. ft., high MB per KW, **but what about reliability?**

# But MTTF goes through the roof...

- If 1 disk has MTTF of 50k hours...
  - 70 disks will have a MTTF of ~700 hours!!!
  - This is assuming failures are independent...
- But fortunately we know when failures occur!
  - Disks use a lot of CRC coding, so we don't have corrupted data, just no data
- We can have both “Soft” and “Hard” failures
  - Soft failure just the read is incorrect/failed, the disk is still good
  - Hard failures kill the disk, necessitating replacement
    - Most RAID setups are “Hot swap”: Unplug the disk and put in a replacement while things are still going
    - Most modern RAID arrays also have “hot spares”: An already installed disk that is used automatically if another disk fails.


# RAID: Redundant Arrays of (Inexpensive) Disks

- Files are "striped" across multiple disks, ex:
- Redundancy yields high data availability
  - Availability: service still provided to user, even if some components failed
- Disks will still fail
- Contents reconstructed from data redundantly stored in the array
  - Capacity penalty to store redundant info
  - Bandwidth penalty to update redundant info on writes
- 6 Raid *Levels*, 0, 1, 5, 6 most common today

# RAID 0: Striping

- "RAID 0" is not actually RAID (no redundancy)
  - It is simply spreading the data across multiple disks
- So, e.g., for 4 disks, stripe-unit address 0 is on disk 0, address 1 is on disk 1, address 2 is on disk 2, address 4 on disk 0...
- Improves bandwidth linearly
  - With 4 disks you have 4x the disk bandwidth
- Doesn't really help latency
  - Still have the individual disks seek and rotation time
- Failures will happen...


# RAID 1: Disk Mirroring/Shadowing (online sparing)



- Each disk is fully duplicated onto its “mirror”  
Very high availability can be achieved
- Writes go to disk and mirror - limited by single-disk speed
- Reads from original disk, unless failure

Most expensive solution: 100% (2x) capacity overhead


# RAID 2: Hamming Code for Error Correction



***Not actually used***, don't worry about remembering it!


# RAID 3: Single Parity Disk

- Disk drives themselves code data and detect failures
- Reconstruction of data can be done with single parity disk **if we know which disk failed**
- Writes change data disk and P disk




# RAID 4: High I/O Rate Parity

- Interleave data at the sector level (data block) rather than bit level - permits more parallelism (independent small reads, parallel large reads)
- Reconstruction of data can be done with single parity disk
- Reading (w/o) fault involve only data disk
- Writing involves data disk + parity disk




# Inspiration for RAID 5

- RAID 4 works well for reads, but
- Parity Disk is the bottleneck for writes: Write to D0, D5 both also write to P disk



# RAID 5: High I/O Rate Interleaved Parity



# RAID 6

- RAID 5 is no longer the “gold standard”
- Can experience 1 disk failure and continue operation
  - RAID array is in a “degraded” state
- But disk failures are not actually independent!
  - When one disk has failed, there’s a decent chance another will fail soon
- RAID 6: Add another parity block per stripe
  - Now 2 blocks per stripe rather than 1
  - Sacrifice capacity for increased redundancy
  - Now the array can tolerate **2** disk failures and continue operating

# Berkeley's Role in Definition of RAID (December 1987)

## A Case for Redundant Arrays of Inexpensive Disks (RAID)

The screenshot shows a Google Scholar search results page. The search query "Case for Raid" is entered in the search bar. The results section is titled "Scholar" and shows "About 138,000 results (0.08 sec)". The top result is a book entry: "[BOOK] A case for redundant arrays of inexpensive disks (RAID) DA Patterson, G Gibson, RH Katz - 1988 - dl.acm.org". The abstract discusses the need for RAID to match CPU and memory performance improvements. It mentions that while SLED capacity has grown rapidly, its performance improvement has been modest. The paper introduces five levels of RAID and compares them to an IBM 3380 and a Fujitsu Super Eagle.

Scholar About 138,000 results (0.08 sec)

Articles [BOOK] [A case for redundant arrays of inexpensive disks \(RAID\)](#)  
DA Patterson, G Gibson, RH Katz - 1988 - dl.acm.org

Legal documents Abstract Increasing performance of CPUs and memories will be squandered if not matched by a similar performance increase in I/O. While the capacity of Single Large Expensive Disks (SLED) has grown rapidly, the performance improvement of SLED has been modest. ...

Any time Cited by 2814 Related articles All 239 versions Cite More▼

*Increasing performance of CPUs and memories will be squandered if not matched by a similar performance increase in I/O. While the capacity of Single Large Expensive Disk (SLED) has grown rapidly, the performance improvement of SLED has been modest. Redundant Arrays of Inexpensive Disks (RAID), based on the magnetic disk technology developed for personal computers, offers an attractive alternative to SLED, promising improvements of an order of magnitude in performance, reliability, power consumption, and scalability.*

*This paper introduces five levels of RAID, giving their relative cost/performance, and compares RAIDs to an IBM 3380 and a Fujitsu Super Eagle.*

# RAID Version 1

Computer Science 61C Spring 2022

and Weaver

- RAID-I (1989)
  - Consisted of a Sun 4/280 workstation with 128 MB of DRAM, four dual-string SCSI controllers, 28 5.25-inch SCSI disks and specialized disk striping software



# RAID Version 2

- 1990-1993
- Early Network Attached Storage (NAS) System running a Log Structured File System (LFS)
- Impact:
  - \$25 Billion/year in 2002
  - Over \$150 Billion in RAID device sold since 1990-2002
  - 200+ RAID companies (at the peak)
  - Software RAID a standard component of modern OSs



# RAID Is Not Enough By Itself

- You don't just have one disk die...
  - You can have more die in a short period of time
  - Thank both the "bathtub curve" and common environmental conditions
- If you care about your data, RAID isn't sufficient
  - You need to also consider a separate backup solution
- A good practice in clusters/warehouse scale computers:
  - RAID-6 in each cluster node with auto-failover and a hot spare
  - Distributed filesystem on top
    - Replicates amongst the cluster nodes so that nodes can fail
  - And then distribute to a different WSC...

# In Conclusion ...

- We have methods to mitigate faults in electronic systems:
  - Design bugs, Manufacturing defects, and Runtime Faults
- Dependability Measures let us quantify
- Dealing with Runtime Faults requires redundancy
  - either more hardware (cost) or more time (performance)
- Redundancy most commonly used in memory systems (DRAM, SRAM, Disks, SSD), also for communications

# Warehouse Scale Cat-puting



# Agenda

- Warehouse-Scale Computing
- Cloud Computing
- Request-Level Parallelism (RLP)
- Map-Reduce Data Parallelism



# Google's WSCs



Ex: In Oregon

# WSC Architecture




## 1U Server:

64 cores, 64 GiB DRAM,  
4x8 TB disk



## Rack:


40-80 servers,  
Local Ethernet (1-10Gbps) switch  
(30\$/1Gbps/server)



## Array (aka cluster):

16-32 racks  
Expensive switch  
(10X bandwidth → 100x cost)

# WSC Storage Hierarchy



1U Server:

DRAM: 64GB, 100ns

Disk: 10TB, 10ms

Rack (80 servers):

DRAM: 5TB, 300µs


Disk: 800TB, 11ms

Array (30 racks):

DRAM: 150TB, 500µs

Disk: 24PB, 12ms

# Early Google Server Internals



# Power Usage Effectiveness


- Energy efficiency
  - Primary concern in the design of WSC
  - Important component of the total cost of ownership
- Power Usage Effectiveness (PUE):

**Total Building Power**

- 
- Power efficiency measure for WSC
  - Not considering efficiency of servers, networking
  - Perfection = 1.0
  - Google WSC's PUE = 1.2

**IT equipment Power**

# Power Usage Effectiveness



$$\text{PUE} = \text{Total Power}/\text{IT Power}$$
$$\text{PUE} = 2.5$$

# Cheating on Cooling

- Normally cooling the air requires big air-conditioning units
  - These suck a lot of power and still consume a lot of water
    - Evaporation of water to dissipate the energy
- Cheat #1: Heat-exchange to a water source
  - Locate your data center on a river or the ocean
  - Or even just put it in a sealed container dropped onto the sea bottom
- Cheat #2: Just have things open to the air!
  - Ups the failure rate, but if the power savings exceed the costs incurred by additional machines dying, it becomes worth it




# Cloud Distinguished by ...

- Shared platform with illusion of isolation
  - Collocation with other tenants
  - Exploits technology of VMs and hypervisors
  - At best “fair” allocation of resources, but not true isolation
- Attraction of low-cost cycles
  - Economies of scale driving move to consolidation
  - Statistical multiplexing to achieve high utilization/efficiency of resources
- Elastic service
  - Pay for what you need, get more when you need it
  - But no performance guarantees: assumes uncorrelated demand for resources

# Cloud Services

Computer Science 61C Spring 2022

- **SaaS:** deliver apps over Internet, eliminating need to install/run on customer's computers, simplifying maintenance and support
  - E.g., Google Docs, Win Apps in the Cloud
- **PaaS:** Deliver computing “stack” as a service, using cloud infrastructure to implement apps. Deploy apps without cost/complexity of buying and managing underlying layers
  - E.g., Hadoop on EC2, Apache Spark on GCP
- **IaaS:** Rather than purchasing servers, software, data center space or net equipment, clients buy resources as an outsourced service. Billed on utility basis.  
Amount of resources consumed/cost reflect level of activity
  - E.g., Amazon Elastic Compute Cloud, Google Compute Platform



# Request-Level Parallelism (RLP)

- Hundreds of thousands of requests per second
  - Popular Internet services like web search, social networking, ...
  - Such requests are largely independent
    - Often involve read-mostly databases
    - Rarely involve read-write sharing or synchronization across requests
- Computation easily partitioned across different requests and even within a request
  - Can often "load balance" just at the DNS level:  
Just tell different people to use a different computer

# Scaled Communities, Processing, and Data

Google cal stanford big game

All Videos News Shopping Maps More Tools

About 7,740,000 results (0.84 seconds)

Top stories




The Mercury News  
The Big Game: Cal claims The Axe after steamrolling Stanford  
14 hours ago



ESPN  
California vs. Stanford - Game Recap - November 20, 2021  
14 hours ago



San Francisco Chronicle  
Cal sets Big Game record for yards, overwhelms Stanford 41-11  
13 hours ago




More images

Big Game

American football

# Google Query-Serving Architecture



# Web Search Result

Google uc berkeley About 429,000,000 results (0.65 seconds)

All News Images Maps Videos More Tools

<https://www.berkeley.edu> ::

**University of California, Berkeley: Home**

@UCBerkeley. Berkeley Talks transcript: Scholars on new book, 'Atmospheres of Violence'  
[news.berkeley.edu/2021/11/19/ ...](https://news.berkeley.edu/2021/11/19/)

Results from berkeley.edu

**Admissions**  
Freshmen Requirements - Contact Us - Dates & Deadlines - Cost

**Graduate Division**  
The Graduate Division oversees graduate admissions ...

**Academic departments**  
Campus Life - Admissions - Skip to

**Academics**  
From 10 faculty members, 40 students and three fields of ...

**Schools & colleges**  
Engineering - Letters & Science - Haas School of Business - ...

**News**  
All news - Campus news - About

 See photos 

**University of California, Berkeley**

Directions Save  
Land-grant university in Berkeley, California

The University of California, Berkeley is a public land-grant research university in Berkeley, California.

# Anatomy of a Web Search (1/3)

- Google “UC Berkeley”
  1. Direct request to “closest” Google Warehouse-Scale Computer
  2. Front-end load balancer directs request to one of many clusters of servers within WSC
  3. Within cluster, select one of many Google Web Servers (GWS) to handle the request and compose the response pages
  4. GWS communicates with Index Servers to find documents that contain the search words, “UC”, “Berkeley”, uses location of search as well as user information
  5. Send information about this search to the node in charge of tracking you
  6. Return document list with associated relevance score

# Anatomy of a Web Search (2/3)

- In parallel,
  - Ad system: if anyone has bothered to advertise for you
  - Customization based on your account
- Use docids (document IDs) to access indexed documents to get snippets of stuff
- Compose the page
  - Result document extracts (with keyword in context) ordered by relevance score
  - A bunch of advertisements (along the top and side)
    - Initially they were easy to see...  
But now they are almost indistinguishable from the desired content

# Anatomy of a Web Search (3/3)

- Implementation strategy
  - Randomly distribute the entries
  - Make many copies of data (aka “replicas”)
  - Load balance requests across replicas
- ***Redundant copies*** of indices and documents
  - Breaks up hot spots — especially popular queries
  - Increases opportunities for ***request-level parallelism***
  - Makes the system more ***tolerant of failures***

# Data-Level Parallelism (DLP)

- SIMD
  - Supports data-level parallelism in a single machine
  - Additional instructions & hardware (e.g., AVX)  
e.g., Matrix multiplication in memory
- DLP on WSC
  - Supports data-level parallelism across ***multiple machines***
  - MapReduce & scalable file systems

# Problem Statement

- How process large amounts of raw data (crawled documents, request logs, ...) every day to compute derived data (inverted indices, page popularity, ...)
  - Each computation is relatively simple but the input data is huge (petabytes) and distributed across 100s or 1000s of servers
- Challenge: Parallelize computation, distribute data, tolerate faults without obscuring simple computation with complex code to deal with issues

# Solution: MapReduce

- Simple data-parallel ***programming model*** and ***implementation*** for processing large datasets
- Users specify the computation in terms of
  - a ***map*** function, and
  - a ***reduce*** function
- Underlying runtime system
  - Automatically ***parallelize*** the computation across large scale clusters of machines
  - ***Handles*** machine ***failure***
  - ***Schedule*** inter-machine communication to make efficient use of the networks


# Inspiration: Map & Reduce Functions, ex: Python

Calculate :

$$\sum_{n=1}^4 n^2$$

```
A = [1, 2, 3, 4]
```

```
def square(x):  
    return x * x  
  
def sum(x, y):  
    return x + y  
  
reduce(sum, map(square,  
A))
```



## Divide and Conquer!

# MapReduce Programming Model

- **Map:**  $(in\_key, in\_value) \rightarrow list(interm\_key, interm\_val)$

```
map(in_key, in_val):  
    // DO WORK HERE  
    emit(interm_key, interm_val)
```

- Slice data into “shards” or “splits” and distribute to workers
- Compute set of intermediate key/value pairs

- **Reduce:**  $(interm\_key, list(interm\_value)) \rightarrow list(out\_value)$

```
reduce(interm_key, list(interm_val)):  
    // DO WORK HERE  
    emit(out_key, out_val)
```


- Combines all intermediate values for a particular key
- Produces a set of merged output values (usually just one)

# MapReduce Execution

Fine granularity tasks: many more map tasks than machines

Bucket sort to get same keys together

2000 servers =>  
≈ 200,000 Map Tasks, ≈ 5,000 Reduce tasks



# MapReduce Word Count Example

## Distribute

that that is is that that is not is not is that it it is			
<b>Map 1</b> that, that, that, 1	<b>Map 2</b> is 1, that 1, that 1	<b>Map 3</b> is 1, not 1, not 1, not 1	<b>Map 4</b> is 1, that, 1t, it, 1t, it, that 1

## Local Sort

## Shuffle



is 6; it 2; not 2; that 5

## Collect

# MapReduce Word Count Example

User-written **Map** function reads the document data and parses the words. For each word, it writes the (key, value) pair of (word, 1). The word is treated as the intermediate key and the associated value of 1 means that we saw the word once.

**Map** phase: (doc name, doc contents) → list(word, count)


```
// "I do I learn" → [ ("I", 1), ("do", 1), ("I", 1), ("learn", 1) ]  
map(key, value):  
    for each word w in value:  
        emit(w, 1)
```

# MapReduce Word Count Example


Intermediate data is then sorted by MapReduce by keys and the user's **Reduce** function is called for each unique key. In this case, Reduce is called with a list of a "1" for each occurrence of the word that was parsed from the document. The function adds them up to generate a total word count for that word.

**Reduce** phase: (word, list(counts)) → (word, count\_sum)

```
// ("I", [1,1]) → ("I", 2)
reduce(key, values):
    result = 0
    for each v in values:
        result += v
    emit(key, result)
```




# MapReduce Processing



1. MR 1st splits the input files into  $M$  “splits” then starts many copies of program on servers


# MapReduce Processing



2. One copy—the master—is special. The rest are workers. The master picks idle workers and assigns each 1 of M map tasks or 1 of R reduce tasks.


# MapReduce Processing

3. A map worker reads the input split. It parses key/value pairs of the input data and passes each pair to the user-defined map function.



# MapReduce Processing


4. Periodically, the buffered pairs are written to local disk, partitioned into  $R$  regions by the partitioning function.



# MapReduce Processing


5. When a reduce worker has read all intermediate data for its partition, it bucket sorts using intermediate keys so that occurrences of same keys are grouped together

(The sorting is needed because typically many different keys map to the same reduce task )




# MapReduce Processing

6. Reduce worker iterates over sorted intermediate data and for each unique intermediate key, it passes key and corresponding set of values to the user's reduce function.



# MapReduce Processing

7. When all map tasks and reduce tasks have been completed, the master wakes up the user program. The MapReduce call in user program returns



Output of MR is in  $R$  output files (1 per reduce task, with file names specified by user); often passed into another MR job so don't concatenate into a single file

# Big Data Frameworks: Hadoop & Spark

- Apache Hadoop
  - Open-source MapReduce Framework
  - Hadoop Distributed File System (HDFS)
  - MapReduce Java APIs
- Apache Spark
  - Fast and general engine for large-scale data processing
  - Originally developed in the AMP lab at UC Berkeley
  - Running on top of HDFS
  - Provides Java, Scala, Python APIs for
    - Database
    - Machine learning
    - Graph algorithms



# Apache Spark


- Resilient Distributed Data Set (RDD): A collection of items partitioned across the members of a cluster
  - Can program against it just like an ordinary list, but operations are carried out in parallel on different machines
- Uses the same file system/infrastructure as Hadoop
  - Reuse existing systems, make it easier for users to transition
- Users can think about writing “ordinary” code to operate against RDDs rather than an explicit map/reduce structure
- Keep intermediate results in memory where possible
  - Issue with Hadoop: Write to disk after each map/reduce cycle, slow and inefficient when we want to compose many operations together (e.g., iterative method)

# Word Count in Spark's Python API

```
file = sc.textFile("hdfs://...")  
// Two kinds of operations:  
// Actions: RDD → Value  
// Transformations: RDD → RDD  
// e.g. flatMap, Map, reduceByKey  
file.flatMap(lambda line: line.split())  
    .map(lambda word: (word, 1))  
    .reduceByKey(lambda a, b: a + b)
```

See <http://spark.apache.org/examples.html>


# What about a *real* application of Spark?



- 50K Blocks in 15 min. -> 8 million Blocks in 15 min.
- How: Spatial partition of the problem

# MapReduce Processing Time Line

- Master assigns map + reduce tasks to “worker” servers
- As soon as a map task finishes, worker server can be assigned a new map or reduce task
- Data shuffle begins as soon as a given Map finishes
- Reduce task begins as soon as all data shuffles finish
- To tolerate faults, reassign task if a worker server “dies”



# A 2003 example...


- ~41 minutes total
  - ~29 minutes for Map tasks & Shuffle tasks
  - ~12 minutes for Reduce tasks
  - 1707 worker servers used
- Map (Green) tasks                    read 0.8 TB, write 0.5 TB
- Shuffle (Red) tasks                read 0.5 TB, write 0.5 TB
- Reduce (Blue) tasks                read 0.5 TB, write 0.5 TB

# MapReduce status: MR\_Indexer-beta6-large-2003\_10\_28\_00\_03

Started: Fri Nov 7 09:51:07 2003 -- up 0 hr 00 min 18 sec

323 workers; 0 deaths

Type	Shards	Done	Active	Input(MB)	Done(MB)	Output(MB)
<a href="#">Map</a>	13853	0	323	878934.6	1314.4	717.0
Shuffle	500	0	323	717.0	0.0	0.0
<a href="#">Reduce</a>	500	0	0	0.0	0.0	0.0



Counters


Variable	
Mapped (MB/s)	72.5
Shuffle (MB/s)	0.0
Output (MB/s)	0.0
doc-index-hits	145825686
docs-indexed	506631
dups-in-index-merge	0
mr-operator-calls	508192
mr-operator-counters	506631

# MapReduce status: MR\_Indexer-beta6-large-2003\_10\_28\_00\_03

Started: Fri Nov 7 09:51:07 2003 -- up 0 hr 05 min 07 sec

1707 workers; 1 deaths

Type	Shards	Done	Active	Input(MB)	Done(MB)	Output(MB)
<a href="#">Map</a>	13853	1857	1707	878934.6	191995.8	113936.6
Shuffle	500	0	500	113936.6	57113.7	57113.7
<a href="#">Reduce</a>	500	0	0	57113.7	0.0	0.0



## Counters


Variable	
Mapped (MB/s)	699.1
Shuffle (MB/s)	349.5
Output (MB/s)	0.0
doc-index-hits	5004411944
docs-indexed	17290135
dups-in-index-merge	0
mr-operator-calls	17331371
mr-operator-outputs	17290135

# MapReduce status: MR\_Indexer-beta6-large-2003\_10\_28\_00\_03

Started: Fri Nov 7 09:51:07 2003 -- up 0 hr 10 min 18 sec

1707 workers; 1 deaths

Type	Shards	Done	Active	Input(MB)	Done(MB)	Output(MB)
<a href="#">Map</a>	13853	5354	1707	878934.6	406020.1	241058.2
Shuffle	500	0	500	241058.2	196362.5	196362.5
<a href="#">Reduce</a>	500	0	0	196362.5	0.0	0.0



## Counters


Variable	
Mapped (MB/s)	704.4
Shuffle (MB/s)	371.9
Output (MB/s)	0.0
doc-index-hits	5000364228
docs-indexed	17300709
dups-in-index-merge	0
mr-operator-calls	17342493
mr-operator-outputs	17300709

# MapReduce status: MR\_Indexer-beta6-large-2003\_10\_28\_00\_03

Started: Fri Nov 7 09:51:07 2003 -- up 0 hr 15 min 31 sec

1707 workers; 1 deaths

Type	Shards	Done	Active	Input(MB)	Done(MB)	Output(MB)
<a href="#">Map</a>	13853	8841	1707	878934.6	621608.5	369459.8
Shuffle	500	0	500	369459.8	326986.8	326986.8
<a href="#">Reduce</a>	500	0	0	326986.8	0.0	0.0



## Counters


Variable	
Mapped (MB/s)	706.5
Shuffle (MB/s)	419.2
Output (MB/s)	0.0
doc-index-hits	4982870667
docs-indexed	17229926
dups-in-index-merge	0
mr-operator-calls	17272056
mr-operator-outputs	17229926

# MapReduce status: MR\_Indexer-beta6-large-2003\_10\_28\_00\_03

Started: Fri Nov 7 09:51:07 2003 -- up 0 hr 29 min 45 sec

1707 workers; 1 deaths

Type	Shards	Done	Active	Input(MB)	Done(MB)	Output(MB)
<a href="#">Map</a>	13853	13853	0	878934.6	878934.6	523499.2
Shuffle	500	195	305	523499.2	523389.6	523389.6
<a href="#">Reduce</a>	500	0	195	523389.6	2685.2	2742.6




Counters		
Variable		
Mapped (MB/s)	0.3	
Shuffle (MB/s)	0.5	
Output (MB/s)	45.7	
doc-index-hits	2313178	$10^6$
docs-indexed	7936	
dups-in-index-merge	0	
mr-merge-calls	1954105	
mr-merge-outputs	1954105	

MapReduce status: MR\_Indexer-beta6-large-2003\_10\_28\_00\_03

Started: Fri Nov 7 09:51:07 2003 -- up 0 hr 31 min 34 sec

1707 workers; 1 deaths

Type	Shards	Done	Active	Input(MB)	Done(MB)	Output(MB)
<a href="#">Map</a>	13853	13853	0	878934.6	878934.6	523499.2
Shuffle	500	500	0	523499.2	523499.5	523499.5
<a href="#">Reduce</a>	500	0	500	523499.5	133837.8	136929.6



## Counters


Variable		
Mapped (MB/s)		0.0
Shuffle (MB/s)		0.1
Output (MB/s)		1238.8
doc- index-hits		0 10
docs- indexed		0
dups-in- index- merge		0
mr- merge- calls	51738599	
mr- merge- outputs	51738599	

# MapReduce status: MR\_Indexer-beta6-large-2003\_10\_28\_00\_03

Started: Fri Nov 7 09:51:07 2003 -- up 0 hr 33 min 22 sec

1707 workers; 1 deaths

Type	Shards	Done	Active	Input(MB)	Done(MB)	Output(MB)
<a href="#">Map</a>	13853	13853	0	878934.6	878934.6	523499.2
Shuffle	500	500	0	523499.2	523499.5	523499.5
<a href="#">Reduce</a>	500	0	500	523499.5	263283.3	269351.2



## Counters


Variable	
Mapped (MB/s)	0.0
Shuffle (MB/s)	0.0
Output (MB/s)	1225.1
doc-index-hits	0 100
docs-indexed	0
dups-in-index-merge	0
mr-merge-calls	51842100
mr-merge-outputs	51842100

# MapReduce status: MR\_Indexer-beta6-large-2003\_10\_28\_00\_03

Started: Fri Nov 7 09:51:07 2003 -- up 0 hr 35 min 08 sec

1707 workers; 1 deaths

Type	Shards	Done	Active	Input(MB)	Done(MB)	Output(MB)
<a href="#">Map</a>	13853	13853	0	878934.6	878934.6	523499.2
Shuffle	500	500	0	523499.2	523499.5	523499.5
<a href="#">Reduce</a>	500	0	500	523499.5	390447.6	399457.2



## Counters


Variable		
Mapped (MB/s)	0.0	
Shuffle (MB/s)	0.0	
Output (MB/s)	1222.0	
doc-index-hits	0	100
docs-indexed	0	
dups-in-index-merge	0	
mr-merge-calls	51640600	
mr-merge-outputs	51640600	

# MapReduce status: MR\_Indexer-beta6-large-2003\_10\_28\_00\_03

Started: Fri Nov 7 09:51:07 2003 -- up 0 hr 37 min 01 sec

1707 workers; 1 deaths

Type	Shards	Done	Active	Input(MB)	Done(MB)	Output(MB)
<a href="#">Map</a>	13853	13853	0	878934.6	878934.6	523499.2
Shuffle	500	500	0	523499.2	520468.6	520468.6
<a href="#">Reduce</a>	500	406	94	520468.6	512265.2	514373.3




Counters	
Variable	
Mapped (MB/s)	0.0
Shuffle (MB/s)	0.0
Output (MB/s)	849.5
doc-index-hits	0 100
docs-indexed	0
dups-in-index-merge	0
mr-merge-calls	35083350
mr-merge-outputs	35083350

# MapReduce status: MR\_Indexer-beta6-large-2003\_10\_28\_00\_03

Started: Fri Nov 7 09:51:07 2003 -- up 0 hr 38 min 56 sec

1707 workers; 1 deaths

Type	Shards	Done	Active	Input(MB)	Done(MB)	Output(MB)
<a href="#">Map</a>	13853	13853	0	878934.6	878934.6	523499.2
Shuffle	500	500	0	523499.2	519781.8	519781.8
<a href="#">Reduce</a>	500	498	2	519781.8	519394.7	519440.7



## Counters


Variable	Value	Unit
Mapped (MB/s)	0.0	
Shuffle (MB/s)	0.0	
Output (MB/s)	9.4	
doc-index-hits	0	1056
docs-indexed	0	3
dups-in-index-merge	0	
mr-merge-calls	394792	
mr-merge-outputs	394792	

# MapReduce status: MR\_Indexer-beta6-large-2003\_10\_28\_00\_03

Started: Fri Nov 7 09:51:07 2003 -- up 0 hr 40 min 43 sec

1707 workers; 1 deaths

Type	Shards	Done	Active	Input(MB)	Done(MB)	Output(MB)
<a href="#">Map</a>	13853	13853	0	878934.6	878934.6	523499.2
Shuffle	500	500	0	523499.2	519774.3	519774.3
<a href="#">Reduce</a>	500	499	1	519774.3	519735.2	519764.0



Counters		
Variable	Value	Unit
Mapped (MB/s)	0.0	
Shuffle (MB/s)	0.0	
Output (MB/s)	1.9	
doc-index-hits	1051	
docs-indexed	0	
dups-in-index-merge	0	
mr-merge-calls	73442	
mr-merge-outputs	73442	

# Important Limitations

- This model only works for certain classes of problems
  - Need parallel compute over data and parallel reduction steps
  - **Critically:** Can divide a problem into many independent subproblems, minimal need for communication among workers when performing their computations
    - “Embarrassingly Parallel”
- Significant Overhead
  - Hadoop Distributed File System: 3x+ redundant storage
  - Lots of startup and control overhead:  
So unless you have many GiB/TiB of data, don't bother!
- For many cases, you are still better served sticking with a traditional database approach with big hardware behind it

# Summary

- Warehouse-Scale Computers (WSCs)
  - New class of computers
  - Scalability, energy efficiency, high failure rate
- Cloud Computing
  - Benefits of WSC computing for third parties
  - “Elastic” pay as you go resource allocation
- Request-Level Parallelism
  - High request volume, each largely independent of other
  - Use replication for better request throughput, availability
- MapReduce Data Parallelism
  - **Map**: Divide large data set into pieces for independent parallel processing
  - **Reduce**: Combine and process intermediate results to obtain final result
  - Hadoop, Spark