

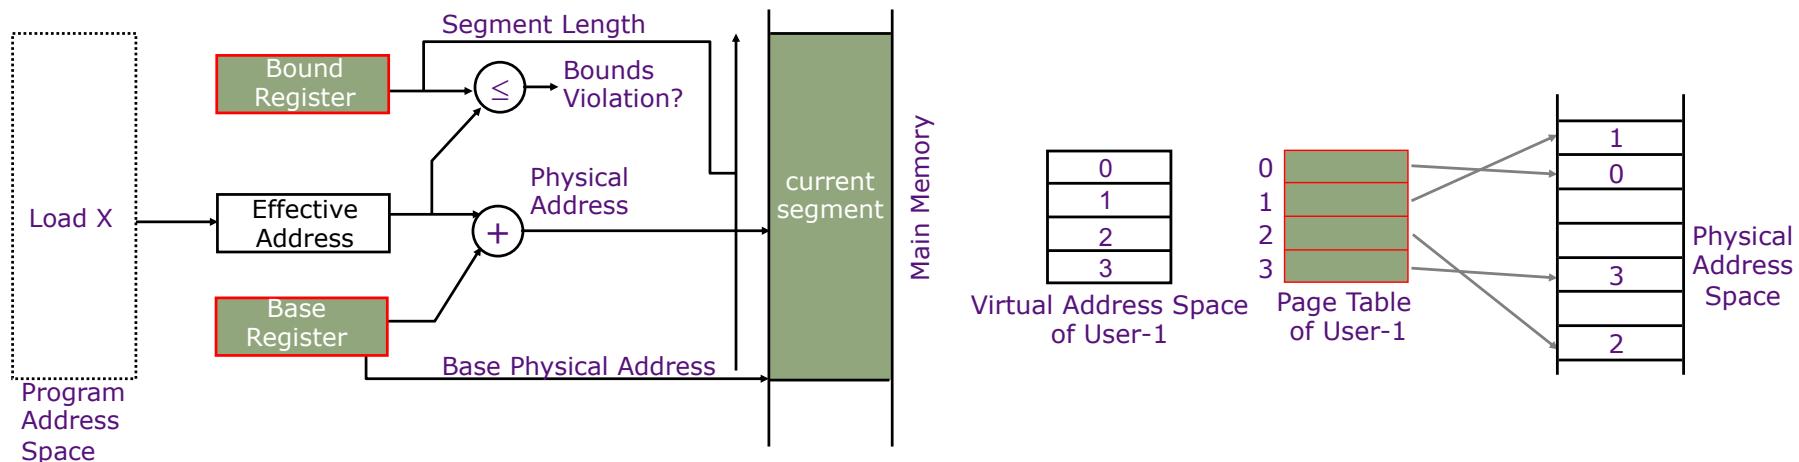
# Modern Virtual Memory Systems

*Daniel Sanchez*

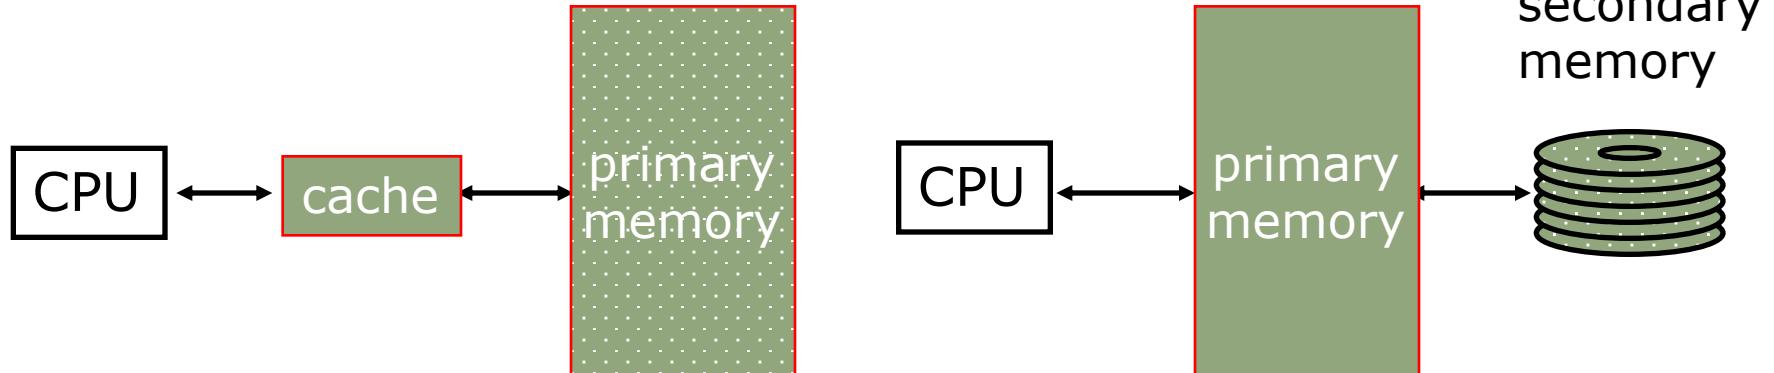
Computer Science and Artificial Intelligence Laboratory  
M.I.T.

# Recap: Evolution of Virtual Memory

- Initial need for multiprogramming
  - Segmentation (Base and bound translation)
- Problem: Fragmentation
  - Paged memory systems
- Problem: Program cannot fit in primary memory
  - Demand paging



# Caching vs. Demand Paging



## *Caching*

- cache entry
- cache block (~32 bytes)
- cache miss rate (1% to 20%)
- cache hit (~1 cycle)
- cache miss (~100 cycles)
- a miss is handled  
*in hardware*

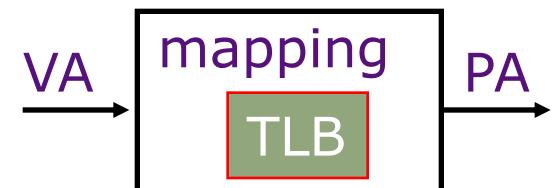
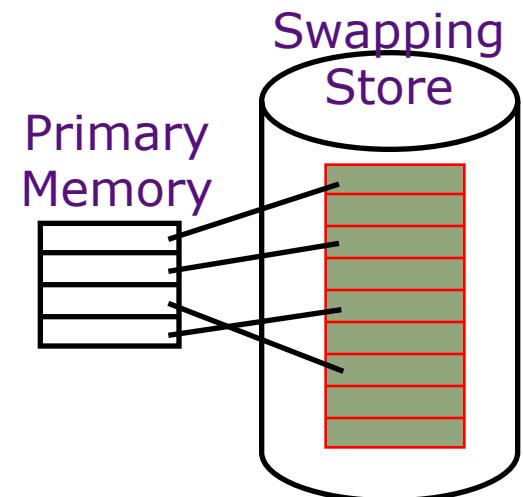
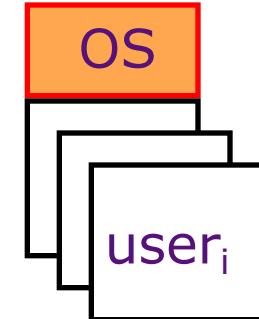
## *Demand paging*

- page frame
- page (~4K bytes)
- page miss rate (<0.001%)
- page hit (~100 cycles)
- page miss (~5M cycles)
- a miss is handled  
*mostly in software*

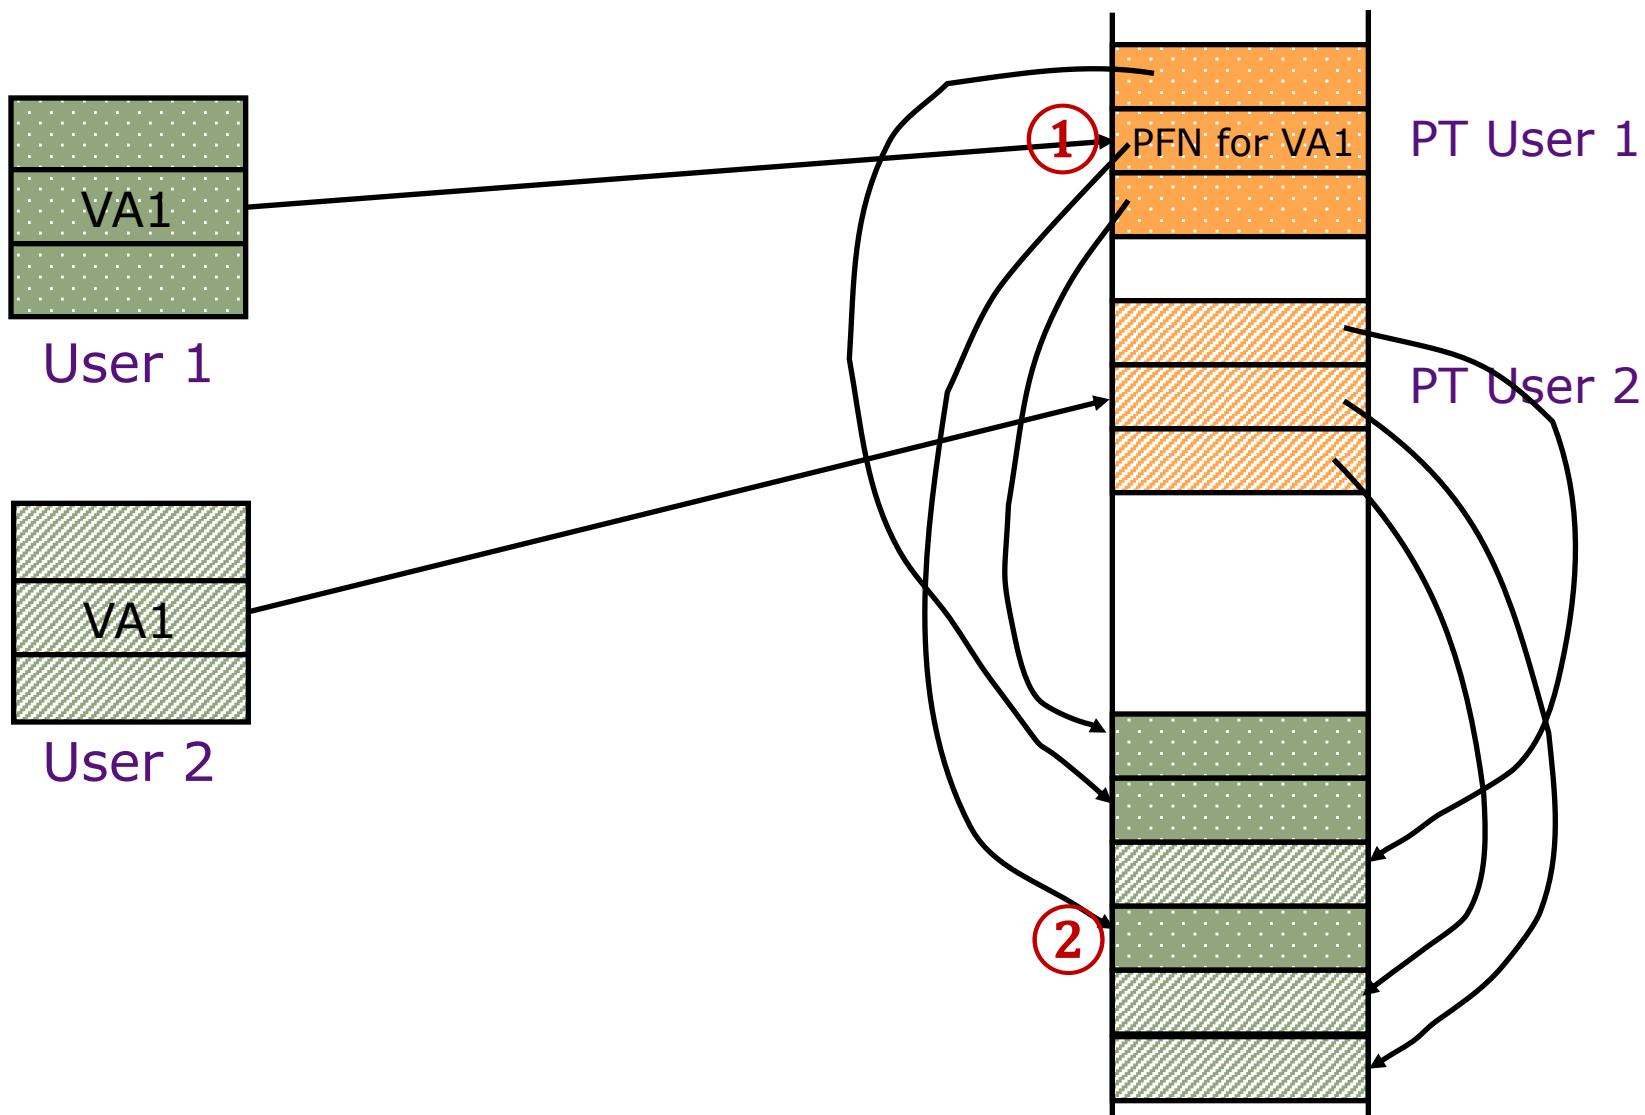
# Modern Virtual Memory Systems

*Illusion of a large, private, uniform store*

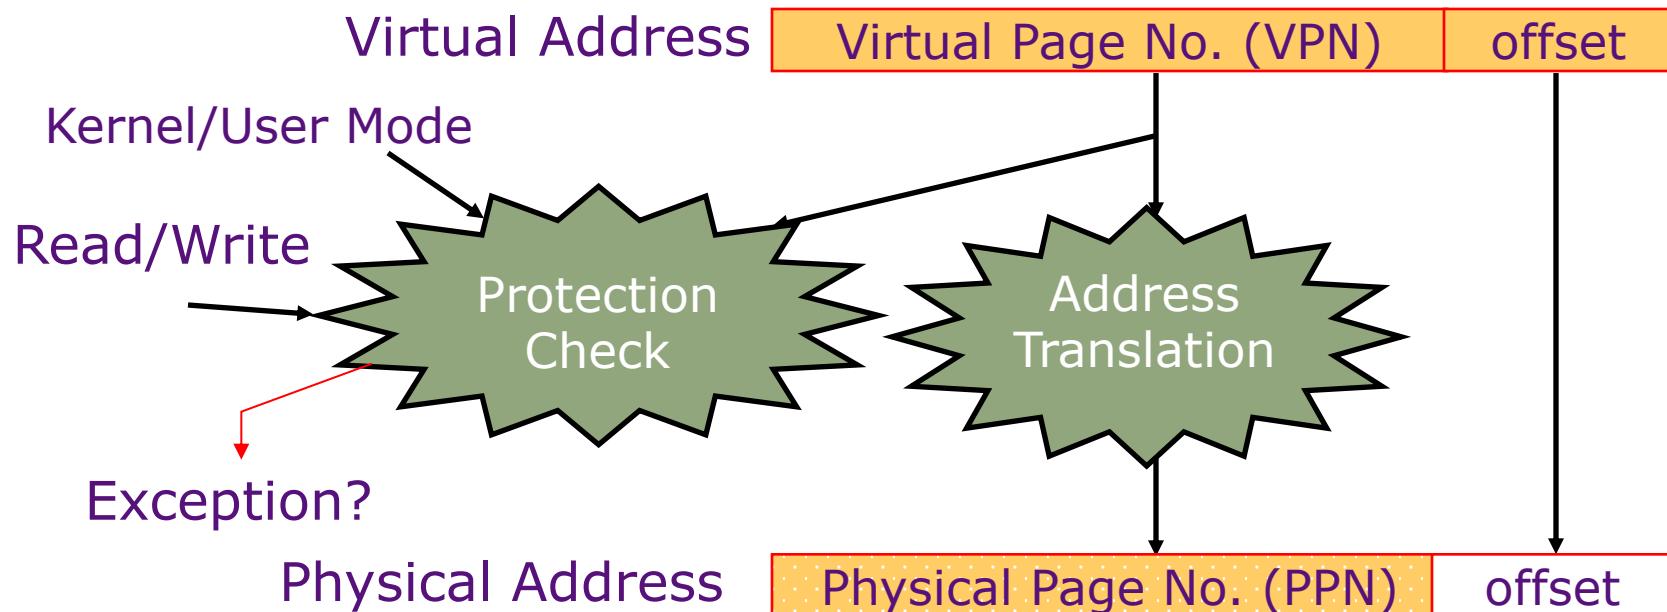
- Protection & Privacy
  - several users, each with their private address space and one or more shared address spaces
  - page table  $\equiv$  memory view  $\equiv$  name space
- Demand Paging
  - Provides the ability to run programs larger than the primary memory
  - Hides differences in machine configurations
- *The price is address translation on each memory reference*



# Reminder: Page Tables are Stored in Physical Memory



# Address Translation & Protection

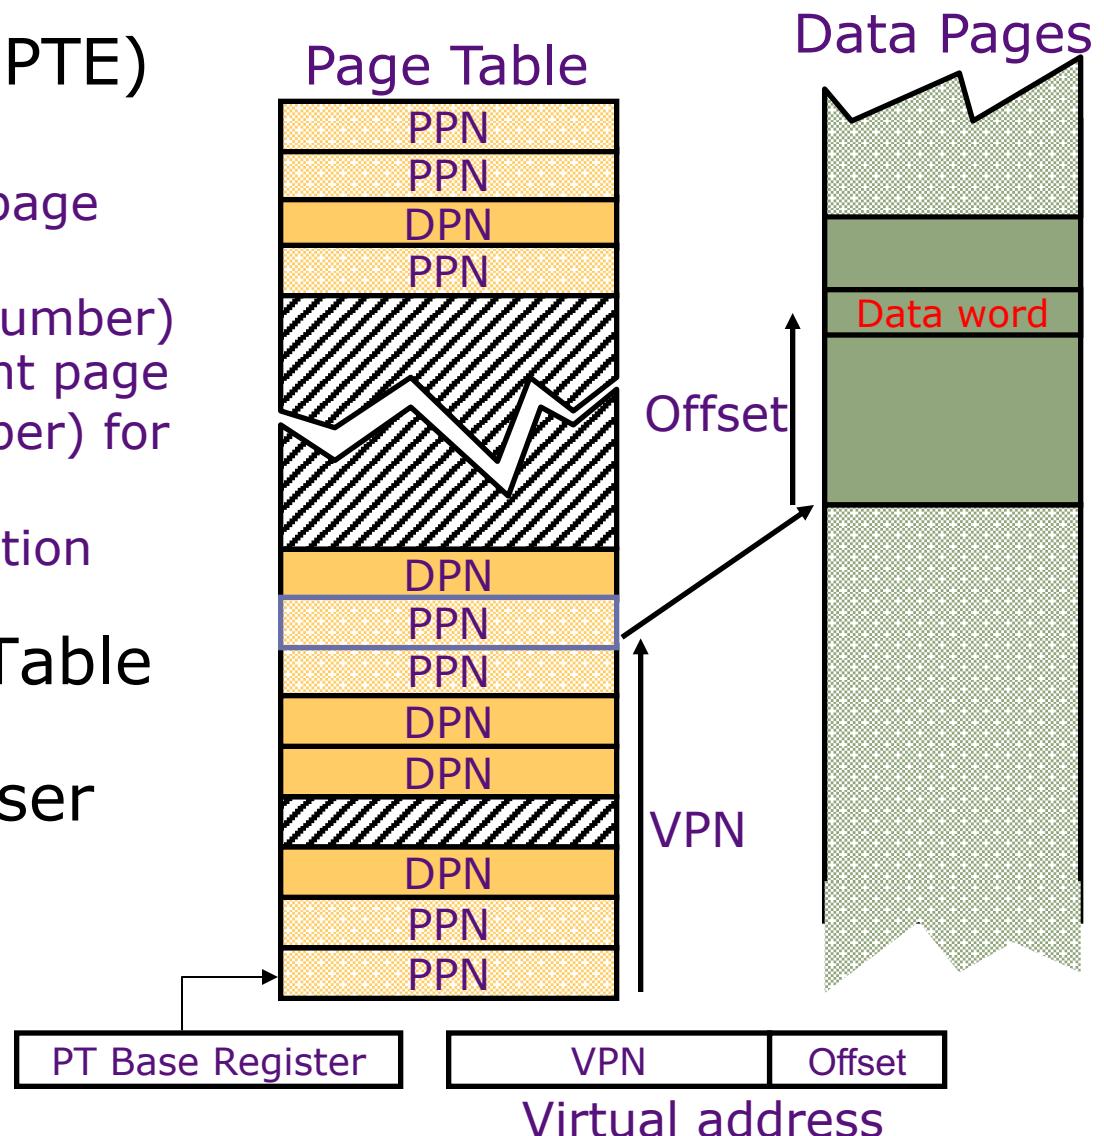


- Every instruction and data access needs address translation and protection checks

*A good Virtual Memory design needs to be space-efficient and fast (~one cycle)*

# Linear Page Table

- Page Table Entry (PTE) contains:
  - A bit to indicate if a page exists
  - PPN (physical page number) for a memory-resident page
  - DPN (disk page number) for a page on the disk
  - Status bits for protection and usage
- OS sets the Page Table Base Register whenever active user process changes



# Size of Linear Page Table

---

With 32-bit addresses, 4 KB pages & 4-byte PTEs:

- ⇒  $2^{20}$  PTEs, i.e., 4 MB page table per user
- ⇒ 4 GB of swap space needed to back up the full virtual address space

Larger pages?

- Internal fragmentation (Not all memory in a page is used)
- Larger page fault penalty (more time to read from disk)

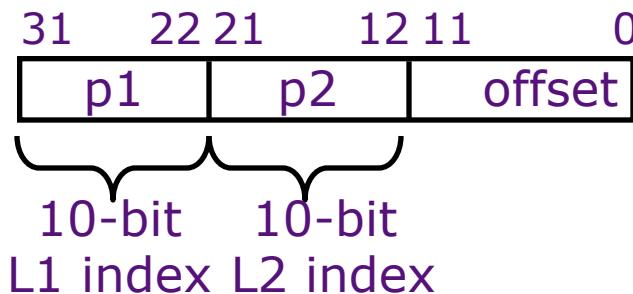
What about 64-bit virtual address space???

- Even 1MB pages would require  $2^{44}$  8-byte PTEs (35 TB!)

*What is the “saving grace”?*

# Hierarchical Page Table

Virtual Address

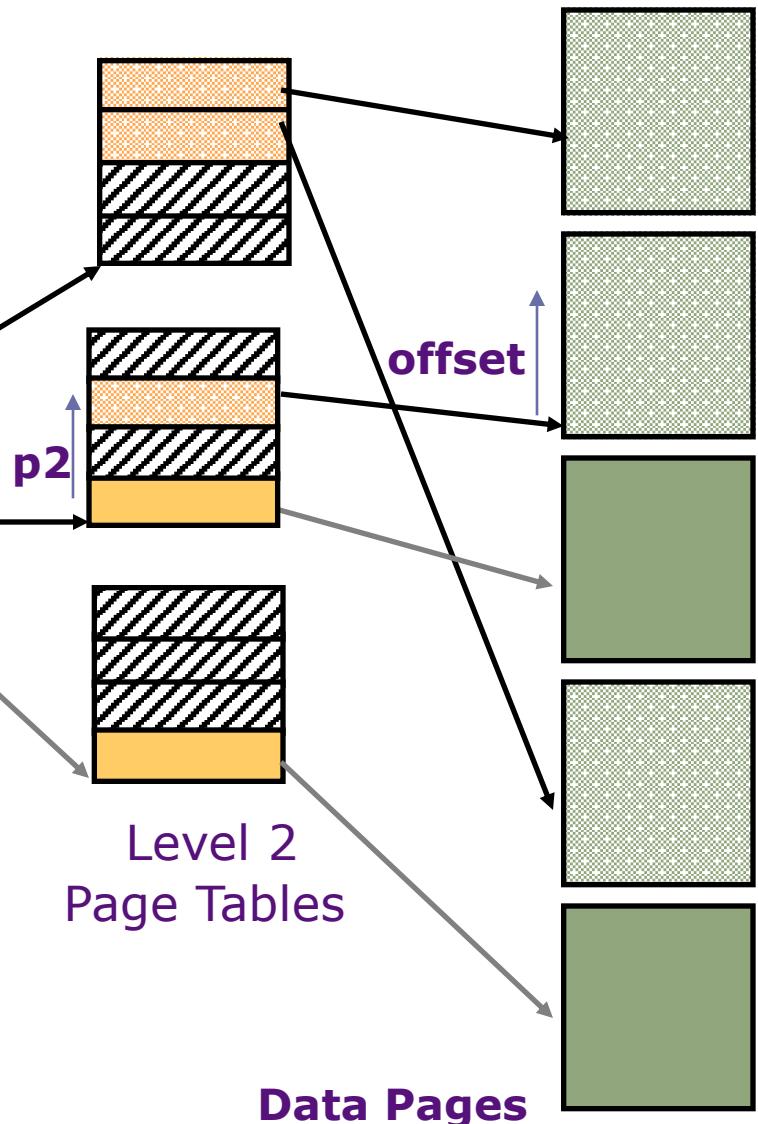


Root of the Current  
Page Table



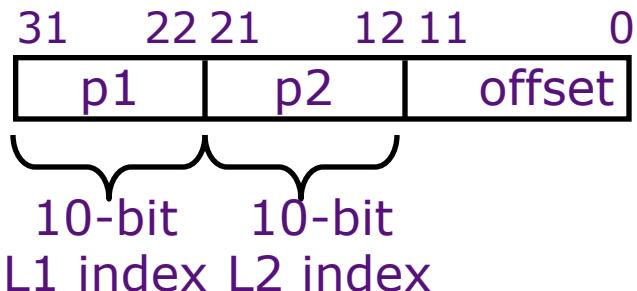
(Processor  
Register)

Level 1  
Page Table



# Supporting Multiple Page Sizes

Virtual Address

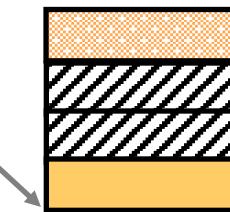
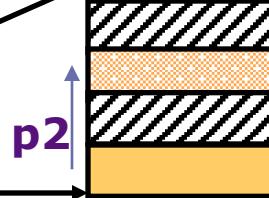


Root of the Current  
Page Table



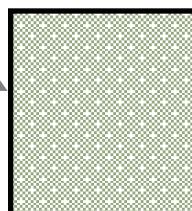
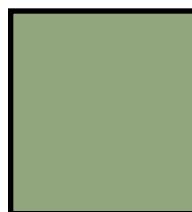
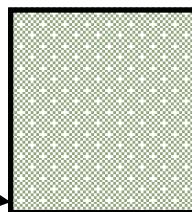
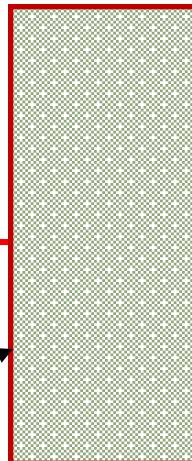
(Processor  
Register)

Level 1  
Page Table



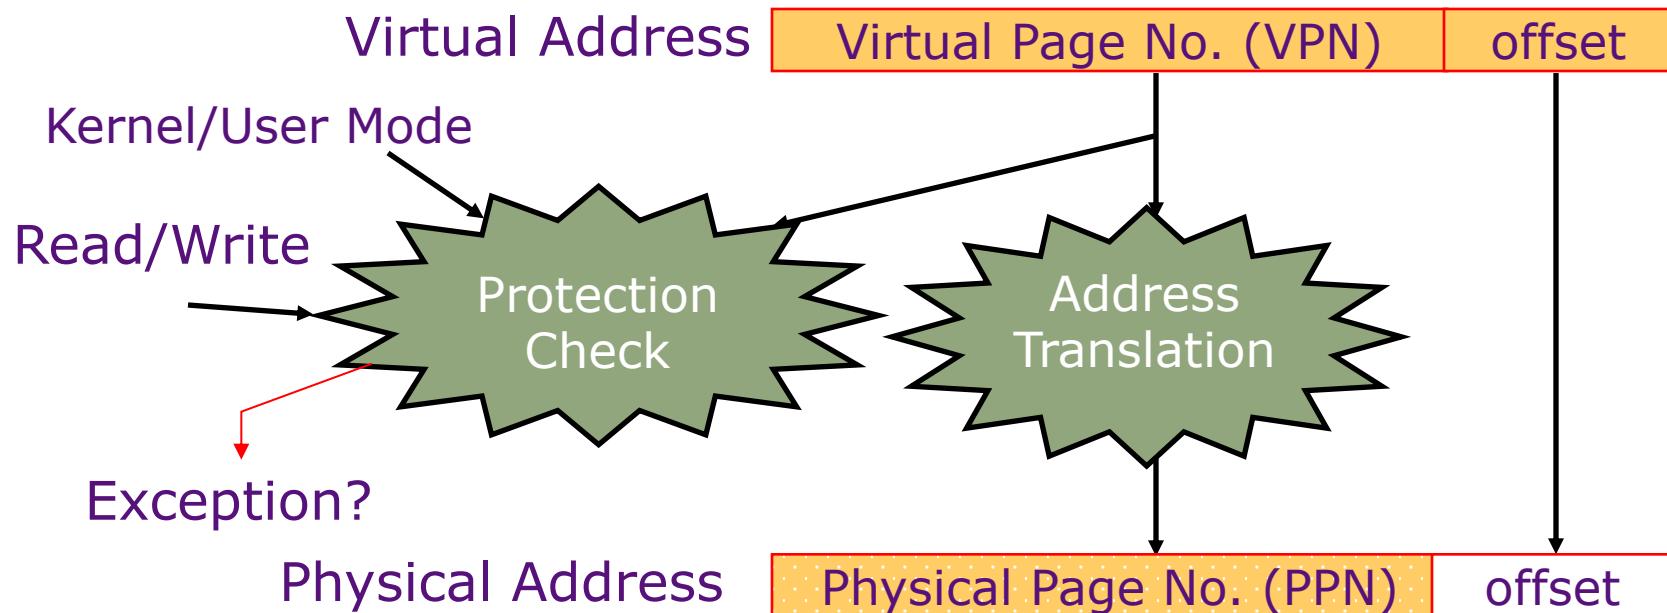
Level 2  
Page Tables

Data Pages



- page in primary memory
- large page in primary memory
- page in secondary memory
- PTE of a nonexistent page

# Address Translation & Protection



- Every instruction and data access needs address translation and protection checks

*A good Virtual Memory design needs to be space-efficient and fast (~one cycle)*

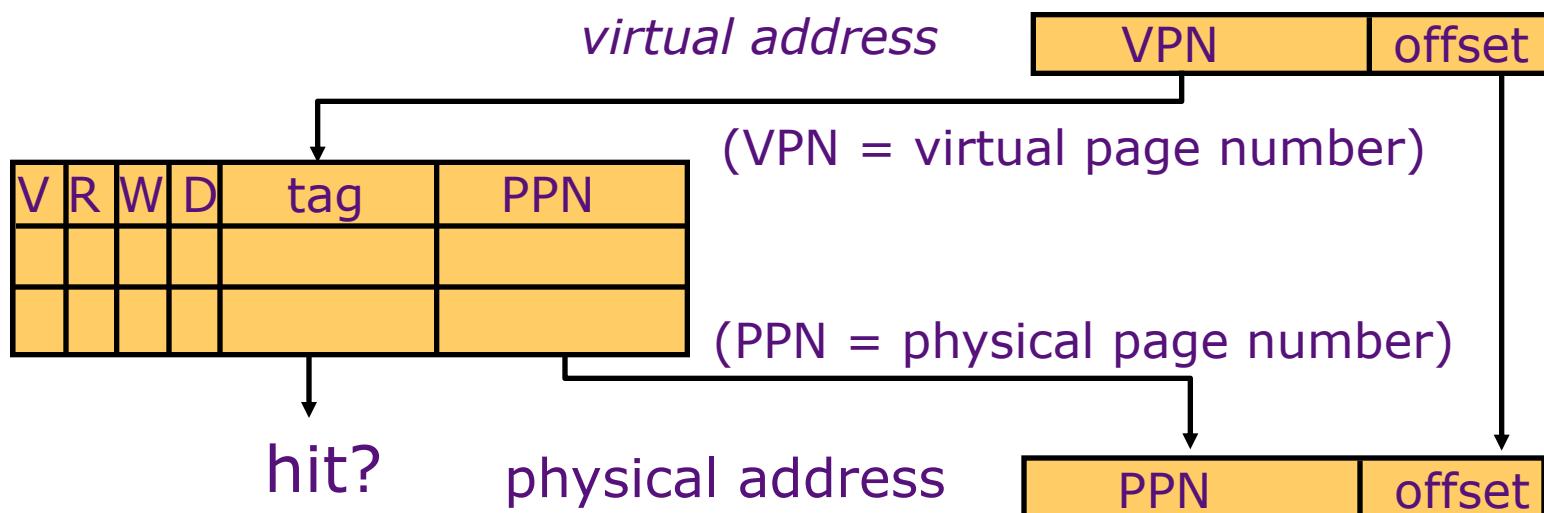
# Translation Lookaside Buffers

Address translation is very expensive!

In a hierarchical page table, each reference becomes several memory accesses

Solution: *Cache translations in TLB*

TLB hit       $\Rightarrow$  Single-cycle Translation  
TLB miss       $\Rightarrow$  Page Table Walk to refill



# TLB Designs

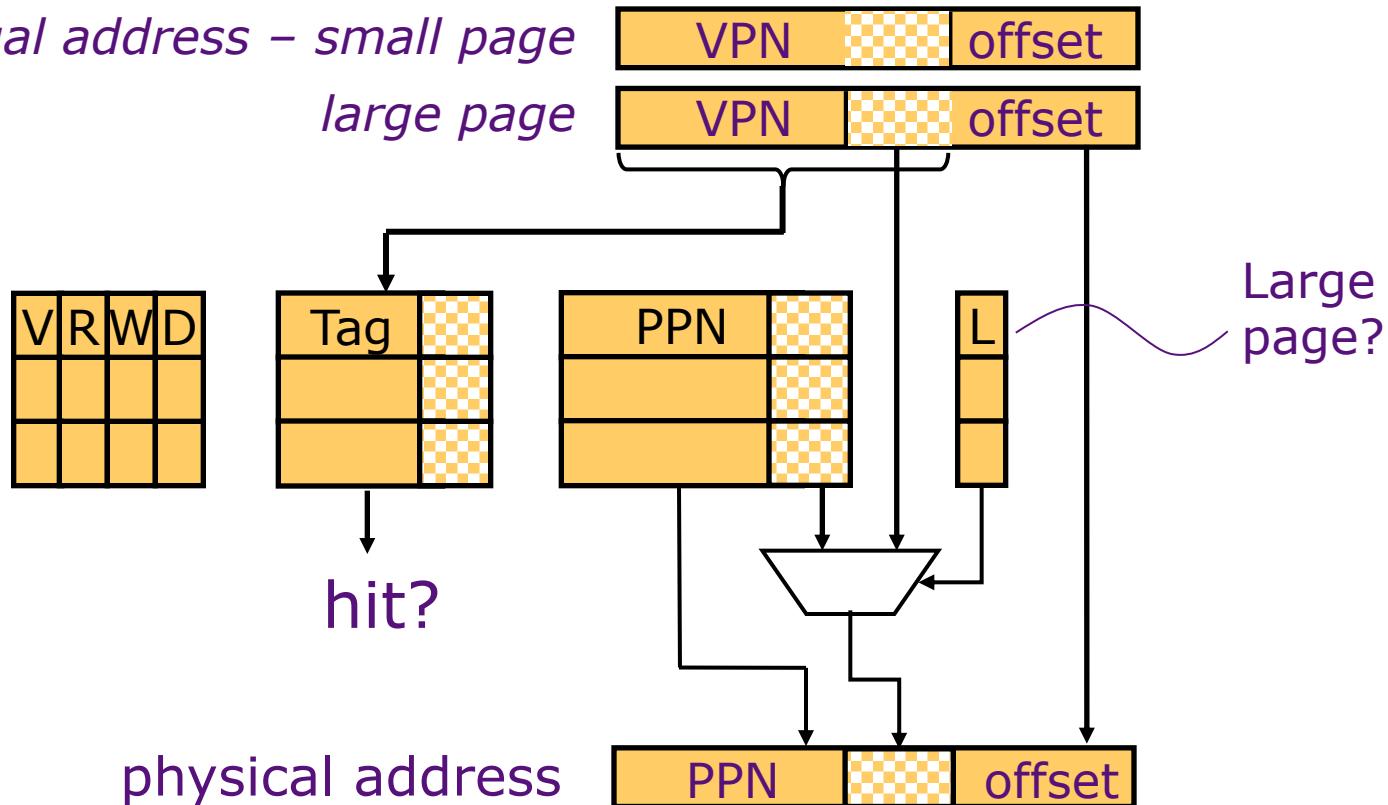
---

- Keep process information in TLB?
  - No process id → Must flush on context switch
  - Tag each entry with process id → No flush, but costlier
- Size and Associativity
  - Typically 32-128 entries, usually highly associative
- TLB Reach: Size of largest virtual address space that can be simultaneously mapped by TLB
  - Example: 64 TLB entries, 4KB pages, one page per entry
  - TLB Reach =  $64 \text{ entries} * 4 \text{ KB} = 256 \text{ KB (if contiguous)}$  ?
- Ways to increase TLB reach
  - Multi-level TLBs (e.g., Intel Skylake: 64-entry L1 data TLB, 128-entry L1 instruction TLB, 1.5K-entry L2 TLB)
  - Multiple page sizes, e.g., x86-64: 4KB, 2MB, 1GB

# TLB Supporting Multiple Page Sizes

*virtual address – small page*

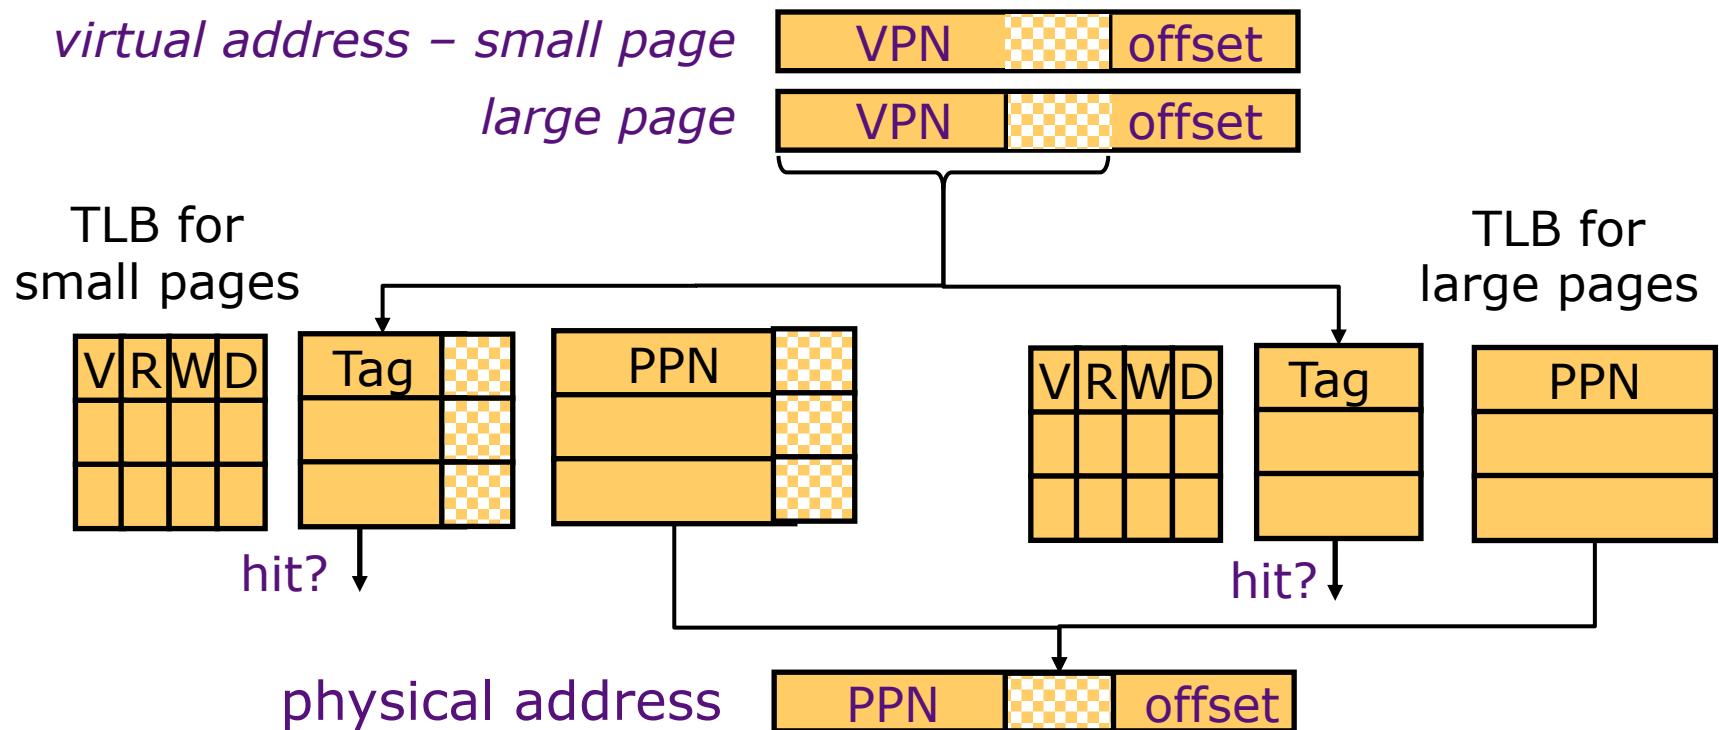
*large page*



Step 1: Assume 4KB page size, calculate index and probe

Step 2: If miss, assume 2MB page, re-calculate index and probe

# TLB Supporting Multiple Page Sizes



Example: Intel Skylake

	<b>4KB</b>	<b>2MB</b>	<b>1GB</b>
<b>L1-D TLB</b>	64	32	4
<b>L1-I TLB</b>	128	8	/
<b>L2 STLB</b>	1536		16

Alternatively, have a separate TLB for each page size

Pros/cons compared to unified TLB?

# Handling a TLB Miss

---

## Software (MIPS, Alpha)

TLB miss causes an exception and the operating system walks the page tables and reloads TLB. A *privileged “untranslated” addressing mode used for walk*

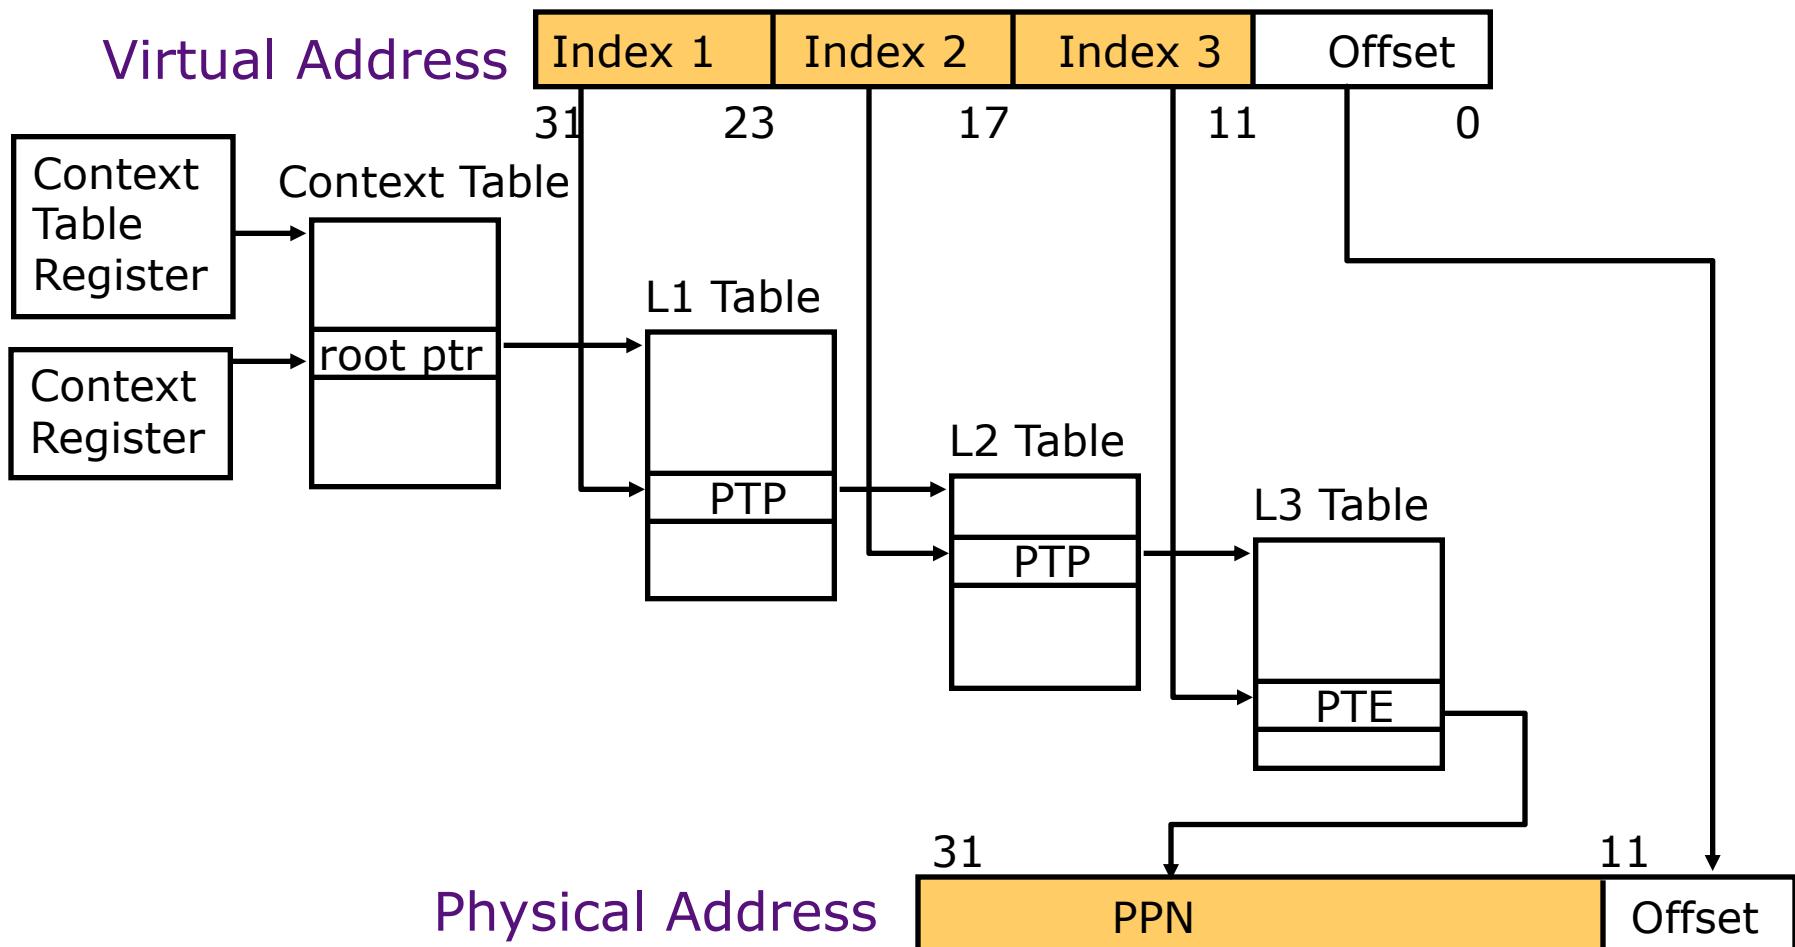
## Hardware (SPARC v8, x86, PowerPC)

A memory management unit (MMU) walks the page tables and reloads the TLB

If a missing (data or PT) page is encountered during the TLB reloading, MMU gives up and signals a Page-Fault exception for the original instruction

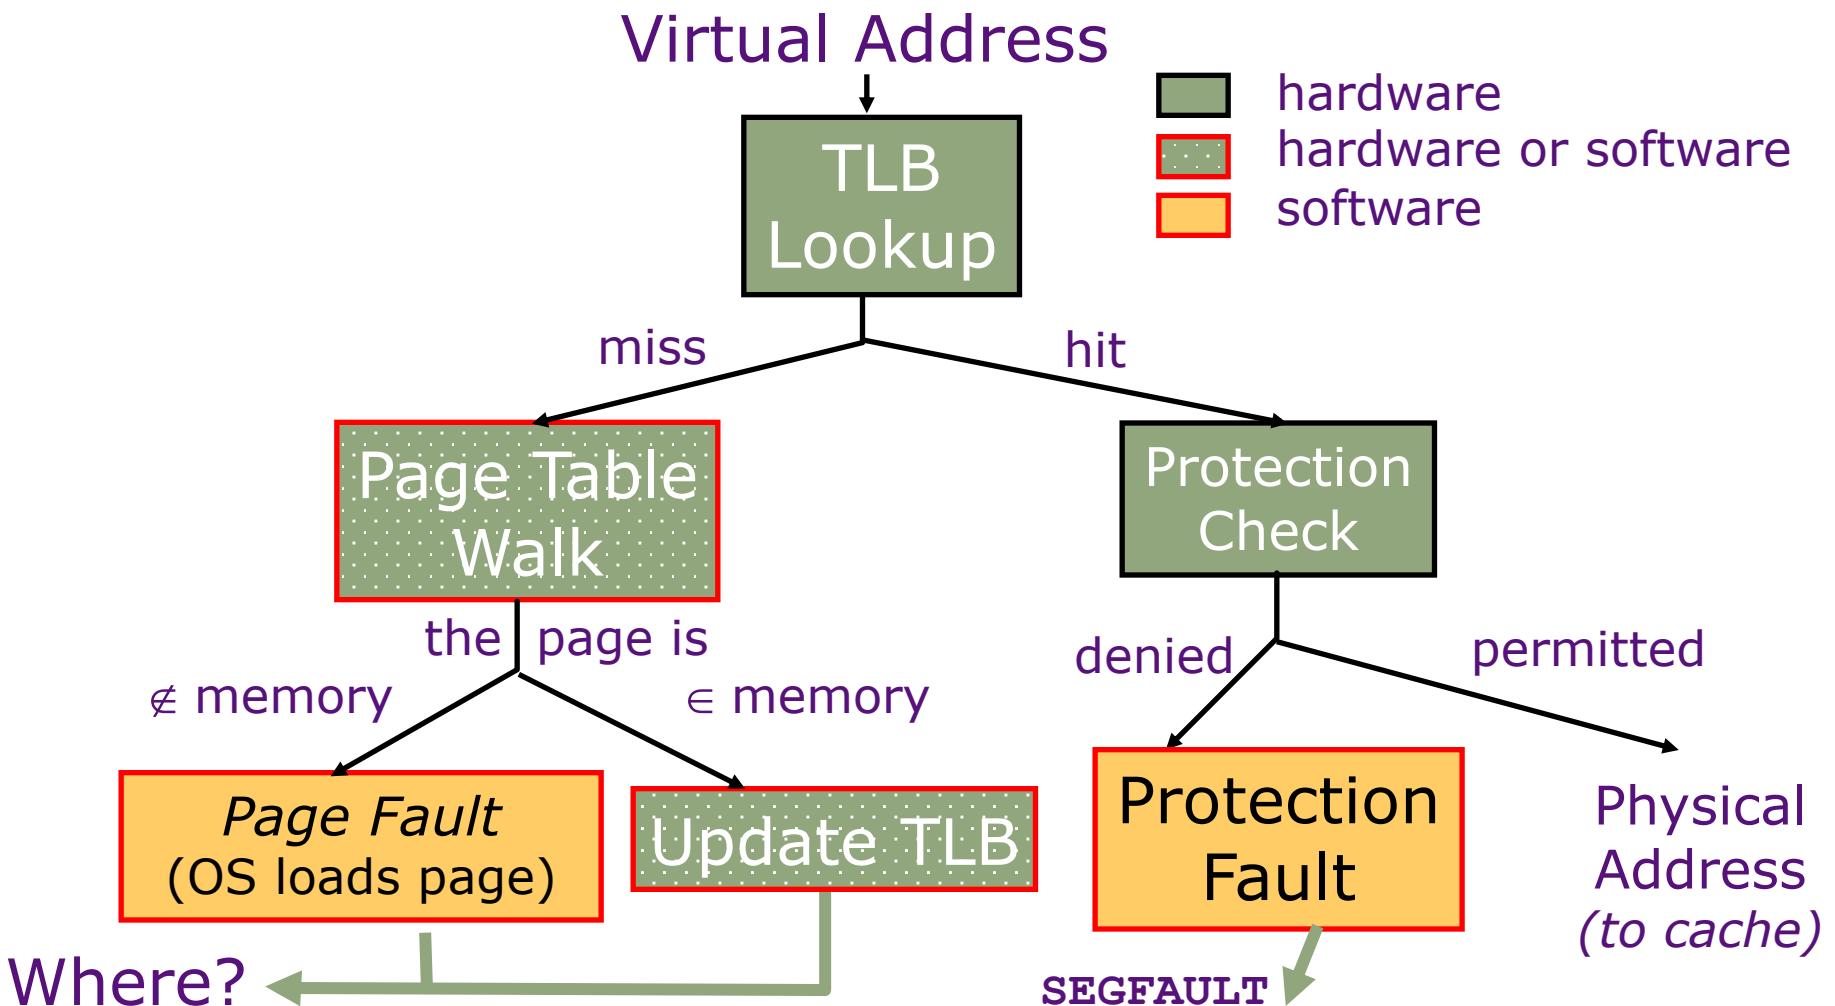
*Pros and cons of each?*

# Hierarchical Page Table Walk: SPARC v8



MMU does this table walk in hardware on a TLB miss

# Address Translation: *putting it all together*

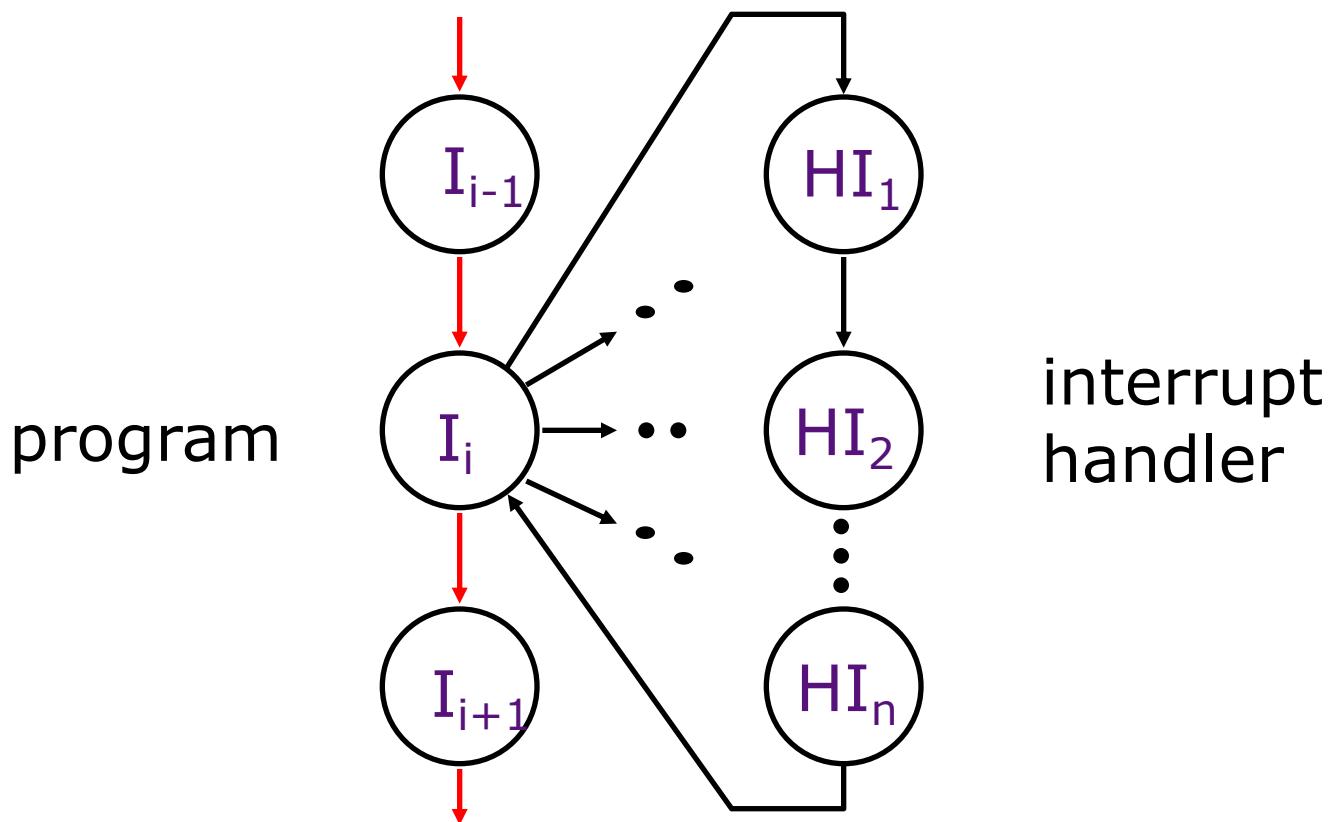


# Topics

---

- Modern VM implementation
  - Page tables: Linear, hierarchical, multiple page sizes
  - TLBs
- Interrupts
- TLB & Cache organization
- Modern usage

# Interrupts: altering the normal flow of control



An *external or internal event* that needs to be processed by another (system) program. The event is usually unexpected or rare from program's point of view.

# Causes of Interrupts

---

Interrupt: an *event* that requests the attention of the processor

- Asynchronous: an *external event*
  - input/output device service-request
  - timer expiration
  - power disruptions, hardware failure
- Synchronous: an *internal event (a.k.a. exception)*
  - undefined opcode, privileged instruction
  - arithmetic overflow, FPU exception
  - misaligned memory access
  - *virtual memory exceptions*: page faults, TLB misses, protection violations
  - *traps*: system calls, e.g., jumps into kernel

# Asynchronous Interrupts

## *Invoking the interrupt handler*

---

- An I/O device requests attention by asserting one of the *prioritized interrupt request lines*
- When the processor decides to process interrupt
  - It stops the current program at instruction  $I_i$ , completing all the instructions up to  $I_{i-1}$  (*precise interrupt*)
  - It saves the PC of instruction  $I_i$  in a special register (EPC)
  - It disables interrupts and transfers control to a designated interrupt handler running in kernel mode

# Interrupt Handler

---

- Saves EPC before enabling interrupts to allow nested interrupts ⇒
  - need an instruction to move EPC into GPRs
  - need a way to mask further interrupts at least until EPC can be saved
- Needs to read a *status register* that indicates the cause of the interrupt
- Uses a special indirect jump instruction `mret` (*return-from-exception*) that
  - enables interrupts
  - restores the processor to the user mode
  - restores hardware status and control state

# Synchronous Interrupts

---

- A synchronous interrupt (exception) is caused by a *particular instruction*
- In general, the instruction cannot be completed and needs to be *restarted* after the exception has been handled
  - With pipelining, requires undoing the effect of one or more partially executed instructions
- In case of a trap (system call), the instruction is considered to have been completed
  - A special jump instruction involving a change to privileged kernel mode

# Page Fault Handler

---

- When the referenced page is not in DRAM:
  - The missing page is located (or created)
  - It is brought in from disk, and page table is updated

*Another job may be run on the CPU while the first job waits for the requested page to be read from disk*
  - If no free pages are left, a page is swapped out

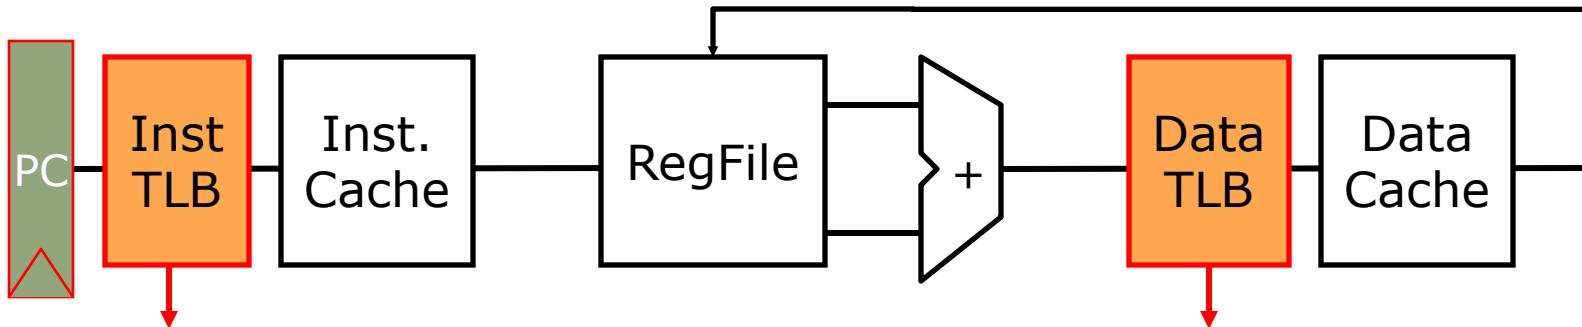
*Pseudo-LRU replacement policy*
- Since it takes a long time to transfer a page (msecs), page faults are handled completely in software by the OS
  - Untranslated addressing mode is essential to allow kernel to access page tables

# Topics

---

- Modern VM implementation
- Interrupts
- TLB & Cache organization
- Modern usage

# Address Translation in CPU

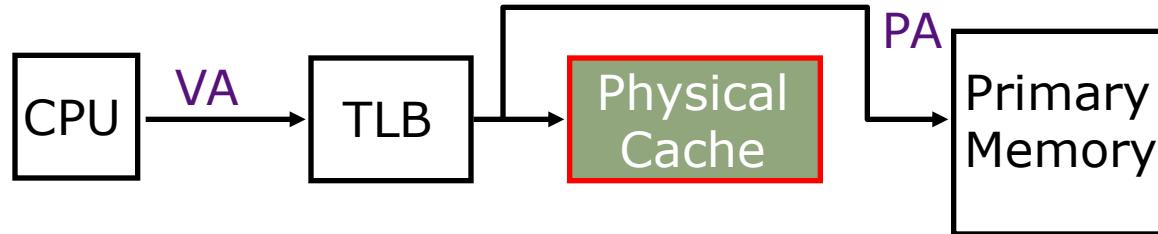


*TLB miss? Page Fault?  
Protection violation?*

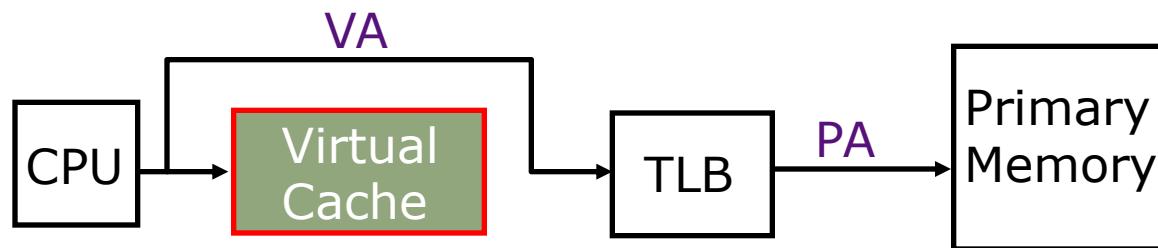
*TLB miss? Page Fault?  
Protection violation?*

- Software handlers need a *restartable* exception on page fault or protection violation
- Handling a TLB miss needs a *hardware* or *software* mechanism to refill TLB
- Need mechanisms to cope with the additional latency of TLB:
  - slow down the clock
  - pipeline the TLB and cache access
  - virtual-address caches
  - parallel TLB/cache access

# Virtual-Address Caches

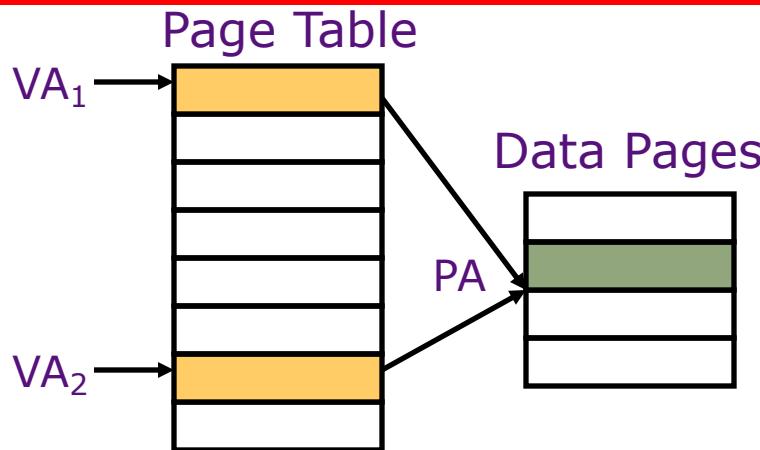


*Alternative: place the cache before the TLB*



- one-step process in case of a hit (+)
- cache needs to be flushed on a context switch unless address space identifiers (ASIDs) included in tags (-)
- *aliasing problems* due to the sharing of pages (-)

# Aliasing in Virtual-Address Caches



Two virtual pages share  
one physical page

Tag	Data
$VA_1$	1st Copy of Data at PA
$VA_2$	2nd Copy of Data at PA

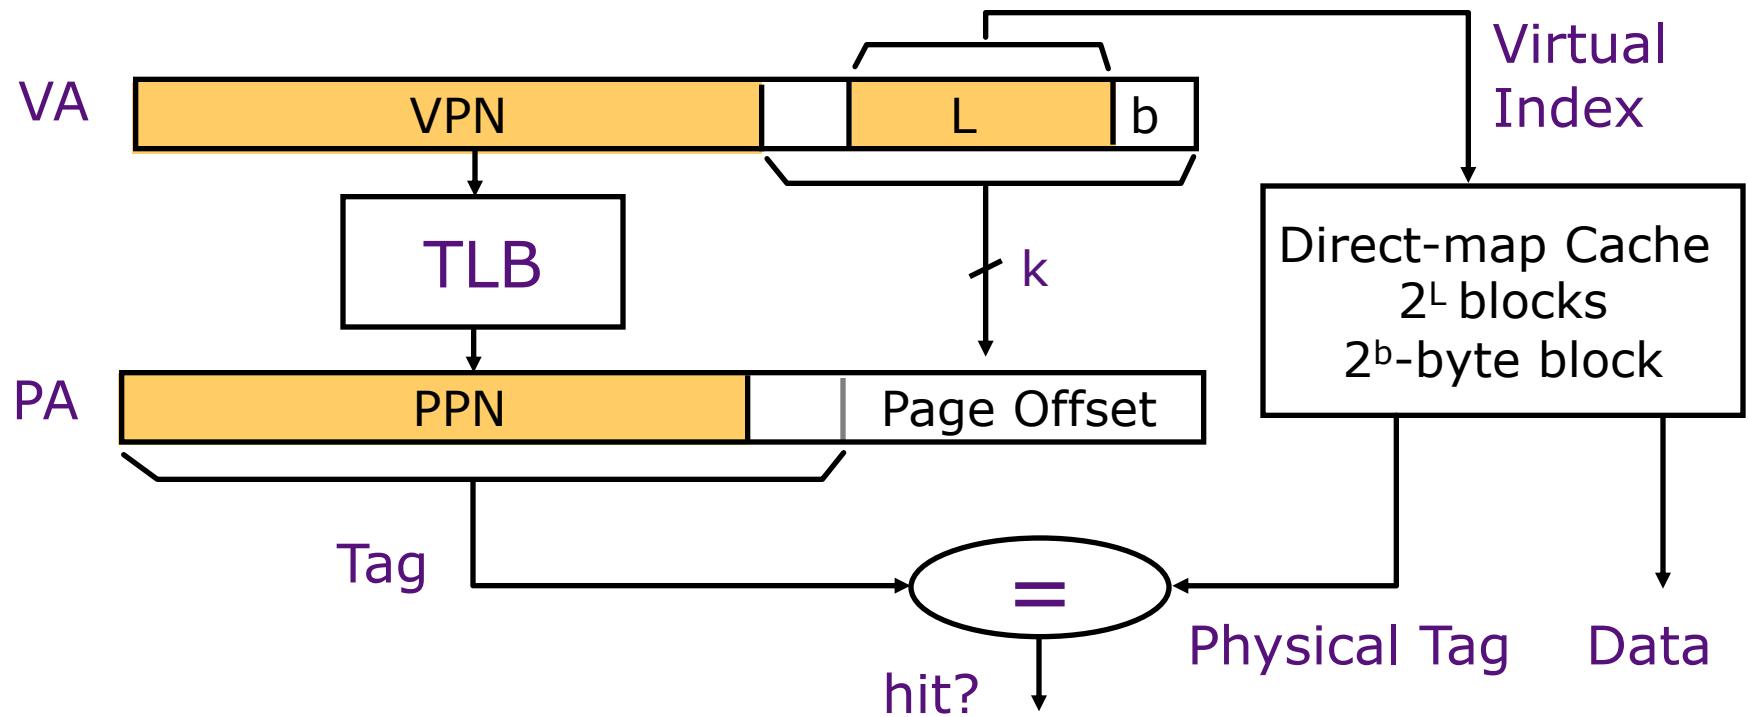
Virtual cache can have two copies of same physical data.  
Writes to one copy not visible to reads of other!

General Solution: *Disallow aliases to coexist in cache*

Software (i.e., OS) solution for direct-mapped cache

VAs of shared pages must agree in cache index bits; this ensures all VAs accessing same PA will conflict in direct-mapped cache (early SPARCcs)

# Concurrent Access to TLB & Cache



Index L is available without consulting the TLB

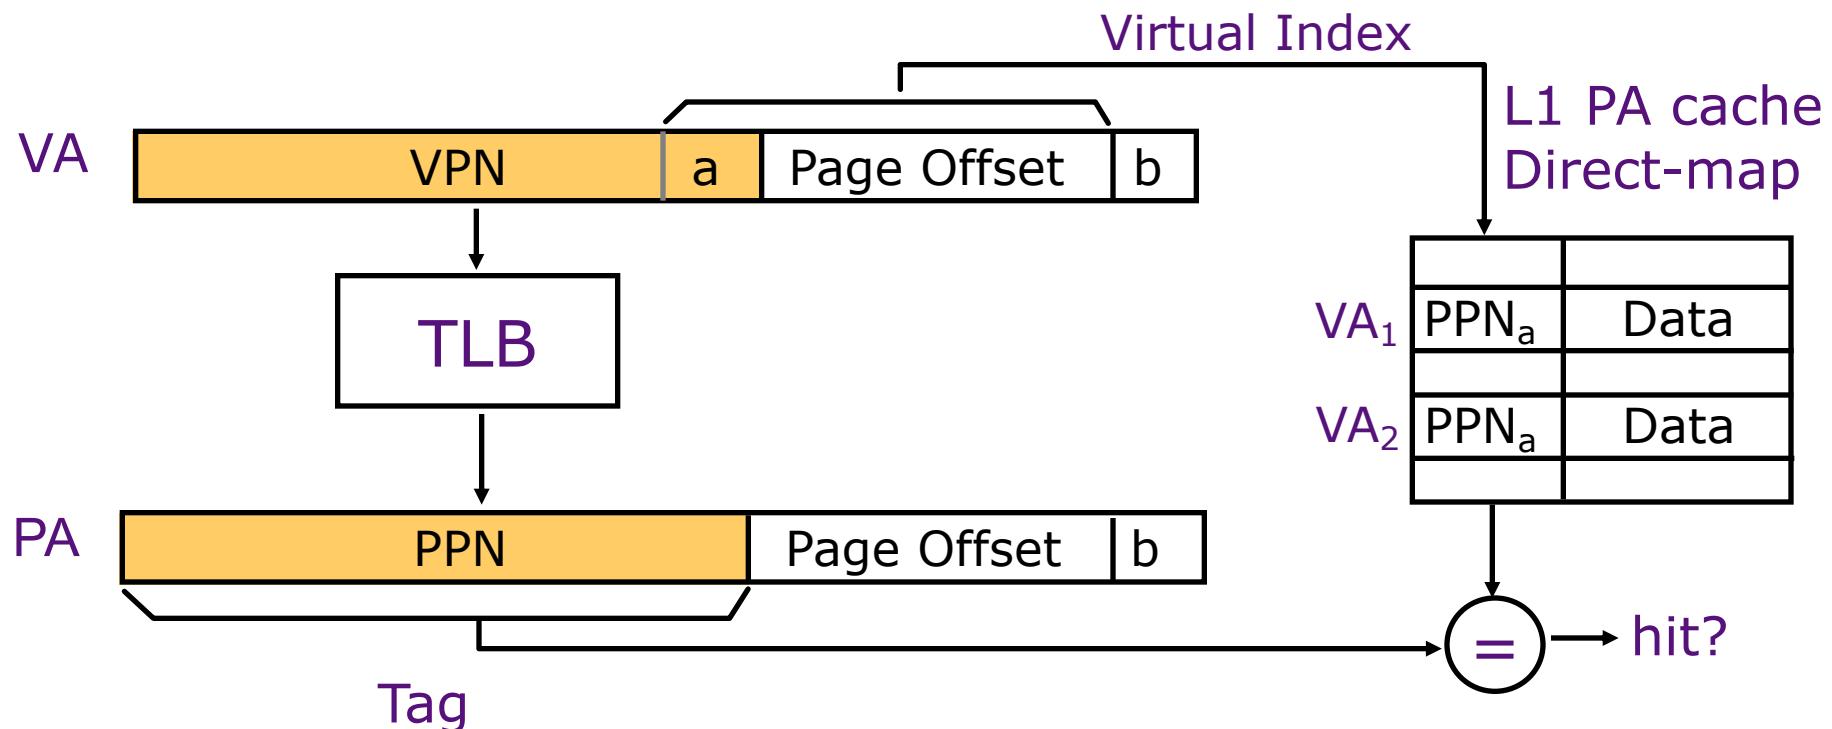
$\Rightarrow$  *cache and TLB accesses can begin simultaneously*

Tag comparison is made after both accesses are completed

*When does this work?*  $L + b < k \checkmark \quad L + b = k \checkmark \quad L + b > k \text{ } \textcolor{red}{X}$

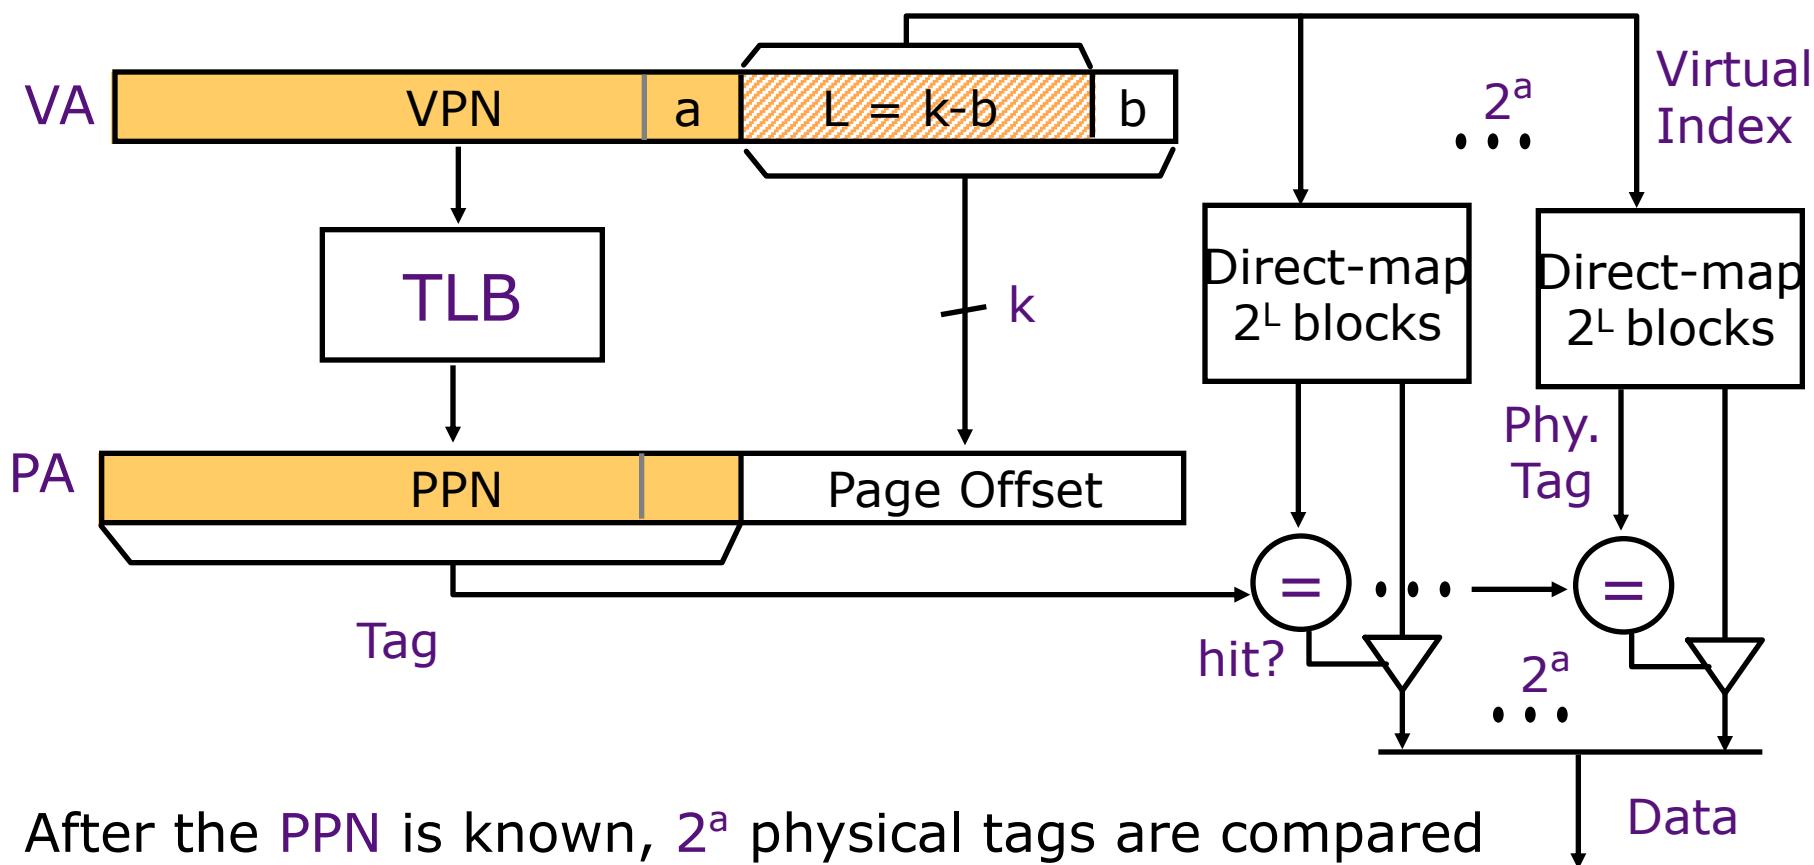
# Concurrent Access to TLB & Large L1

The problem with  $L1 > \text{Page size}$



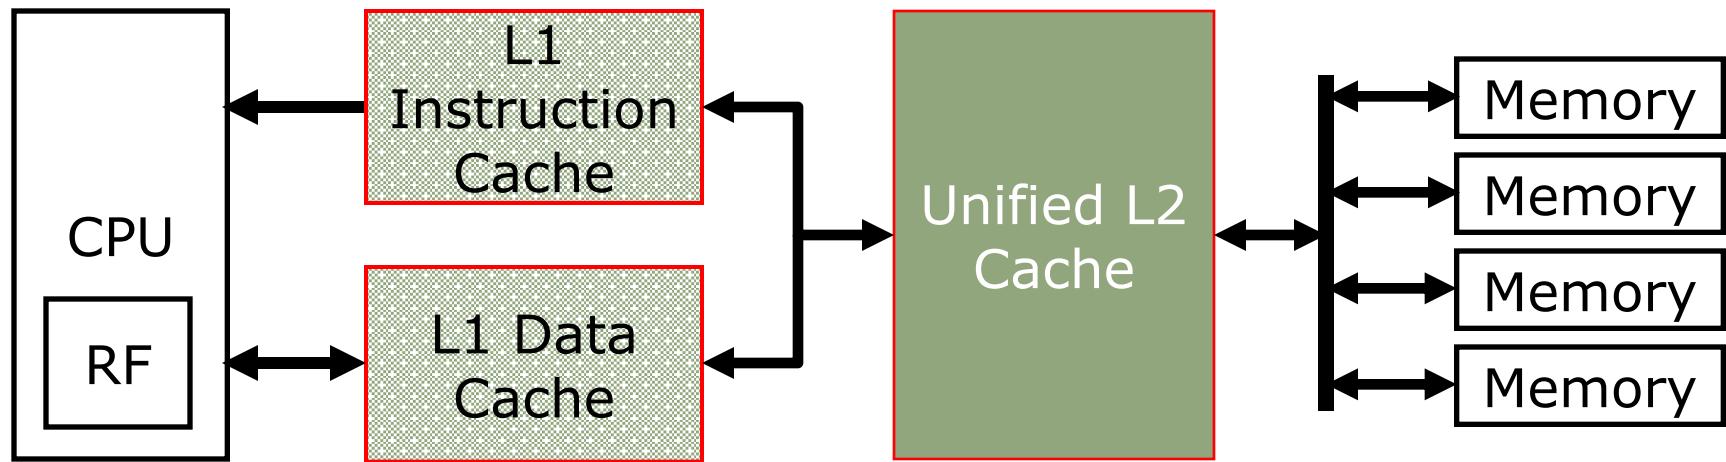
*Can VA<sub>1</sub> and VA<sub>2</sub> both map to PA? Yes*

# Virtual-Index Physical-Tag Caches: Associative Organization



*Is this scheme realistic?*

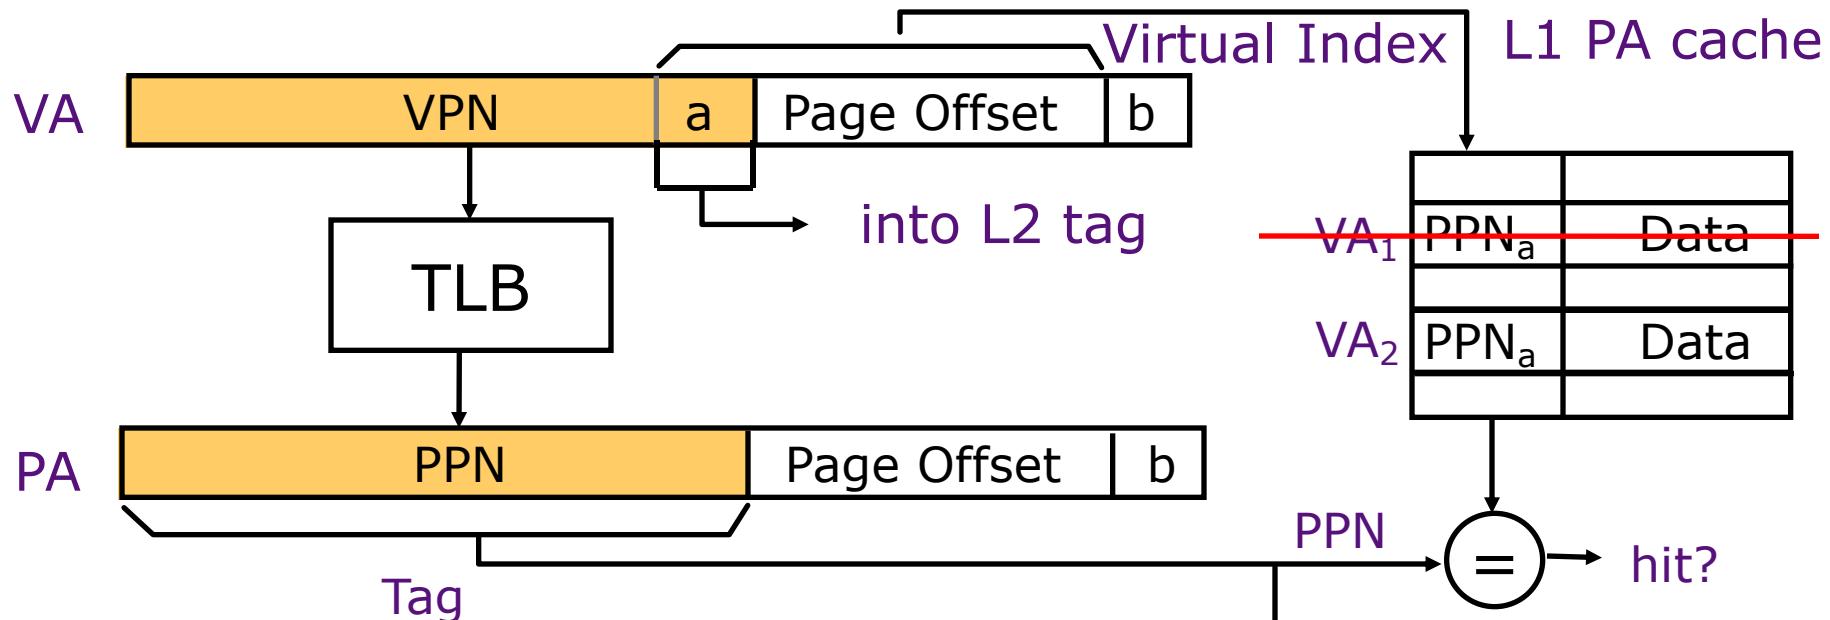
# A solution via Second-Level Cache



Usually a common L2 cache backs up both  
Instruction and Data L1 caches

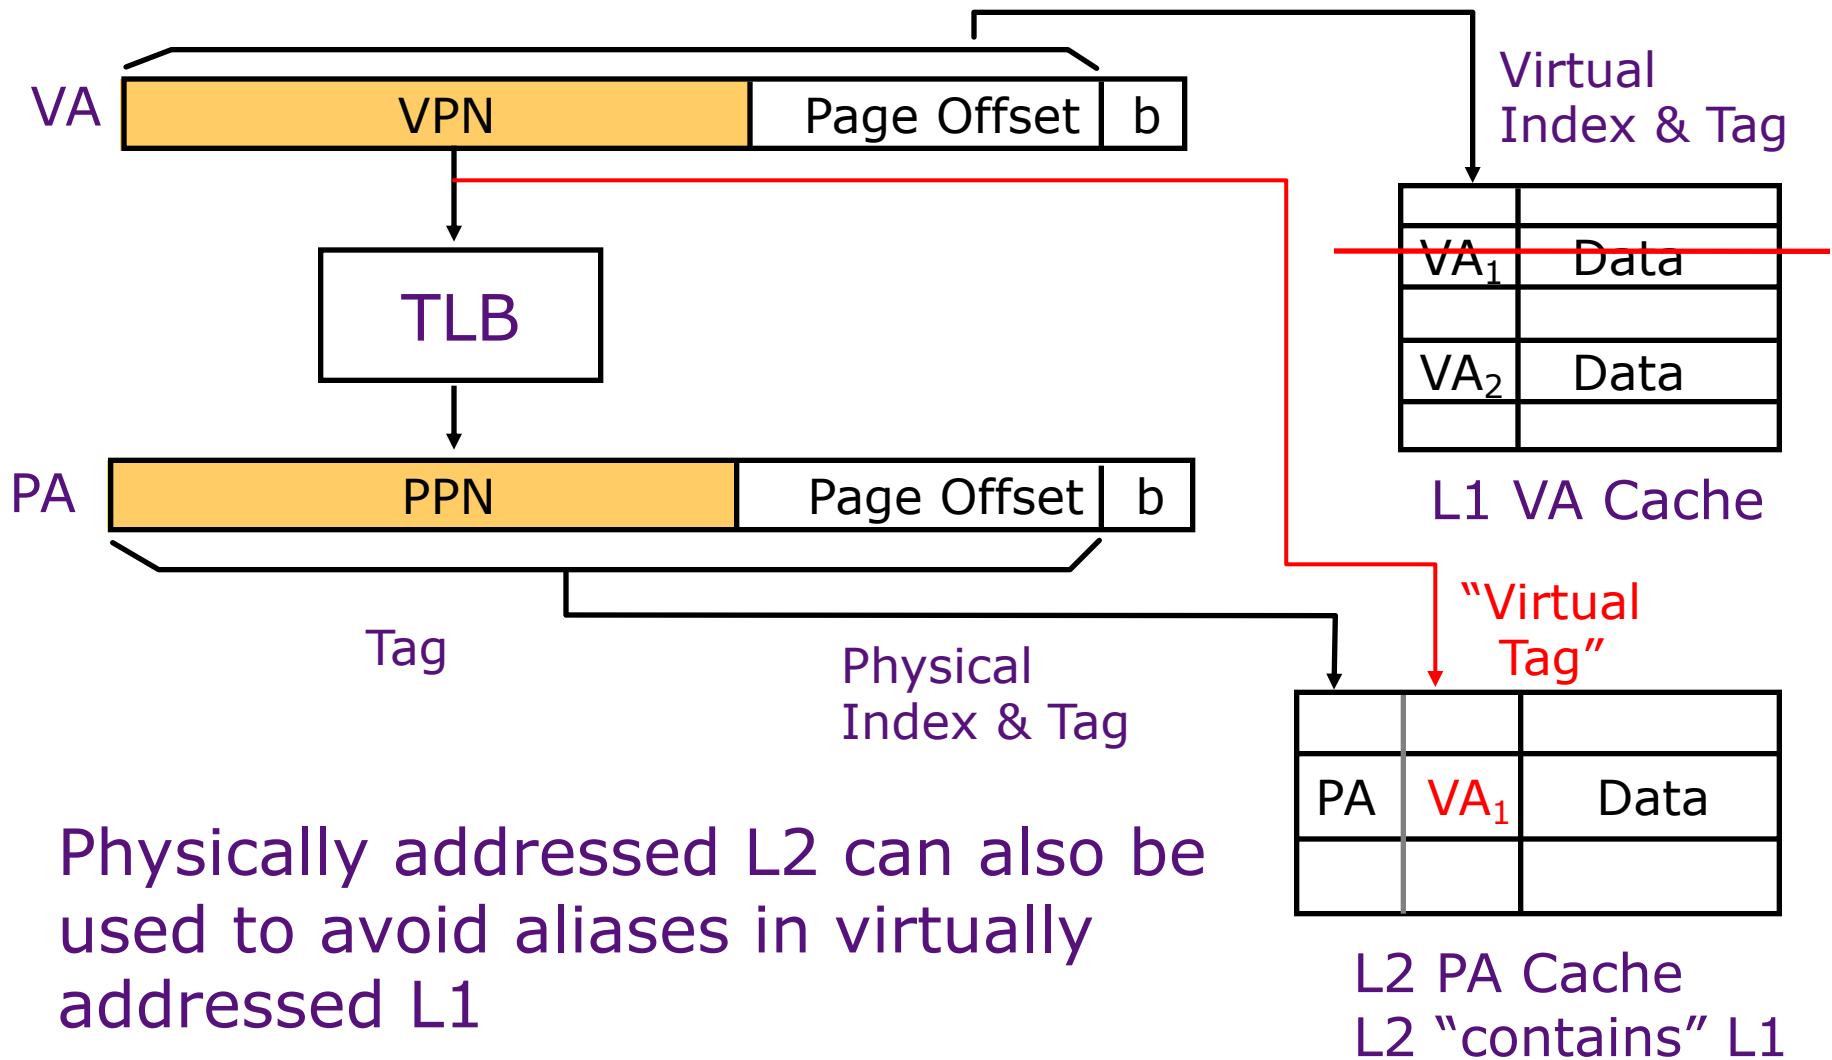
L2 is “inclusive” of both Instruction and Data caches

# Anti-Aliasing Using L2: MIPS R10000



- Suppose VA1 and VA2 both map to PA and VA1 is already in L1, L2 ( $VA_1 \neq VA_2$ )
- After VA2 is resolved to PA, collision is detected in L2. Collision  $\rightarrow$  **Field a is different.**
- VA1 will be purged from L1, and VA2 will be loaded  $\Rightarrow$  *no aliasing!*

# Virtually Addressed L1: Anti-Aliasing using L2



# Topics

---

- Modern VM implementation
- Interrupts
- TLB & Cache organization
- **Modern usage**

# Virtual Memory Use Today - 1

---

- Desktop/server/cellphone processors have full demand-paged virtual memory
  - Portability between machines with different memory sizes
  - Protection between multiple users or multiple tasks
  - Share small physical memory among active tasks
  - Simplifies implementation of some OS features
- Vector supercomputers and GPUs have translation and protection but not demand paging  
(Older Crays: base&bound, Japanese & Cray X1: pages)
  - Don't waste expensive processor time thrashing to disk (make jobs fit in memory)
  - Mostly run in batch mode (run set of jobs that fits in memory)
  - Difficult to implement restartable vector instructions

# Virtual Memory Use Today - 2

---

- Most embedded processors and DSPs provide physical addressing only
  - Can't afford area/speed/power budget for virtual memory support
  - Often there is no secondary storage to swap to!
  - Programs custom-written for particular memory configuration in product
  - Difficult to implement restartable instructions for exposed architectures

*Next lecture:*  
Pipelining!